



Machine Learning and Generative AI for Large-Scale Traveling Salesman Problems

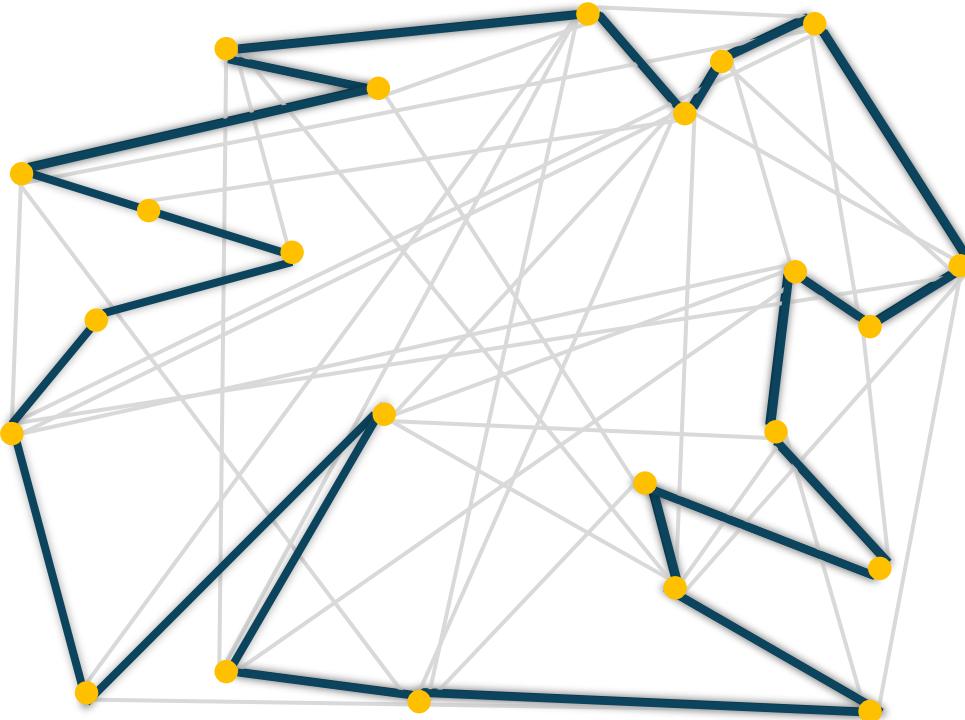
Bocheng Lin
Supervisors: Yi Mei, Mengjie Zhang
Victoria University of Wellington

2024.12

- 
- 1. Introduction**
 - 2. Motivations**
 - 3. Research Objectives**
 - 4. Preliminary Works**
 - 5. Proposed Contributions and Research Plan**

- 1. Introduction**
- 2. Motivations**
- 3. Research Objectives**
- 4. Preliminary Works**
- 5. Proposed Contributions and Research Plan**

Traveling Salesman Problem



a complete graph $G = (V, E)$

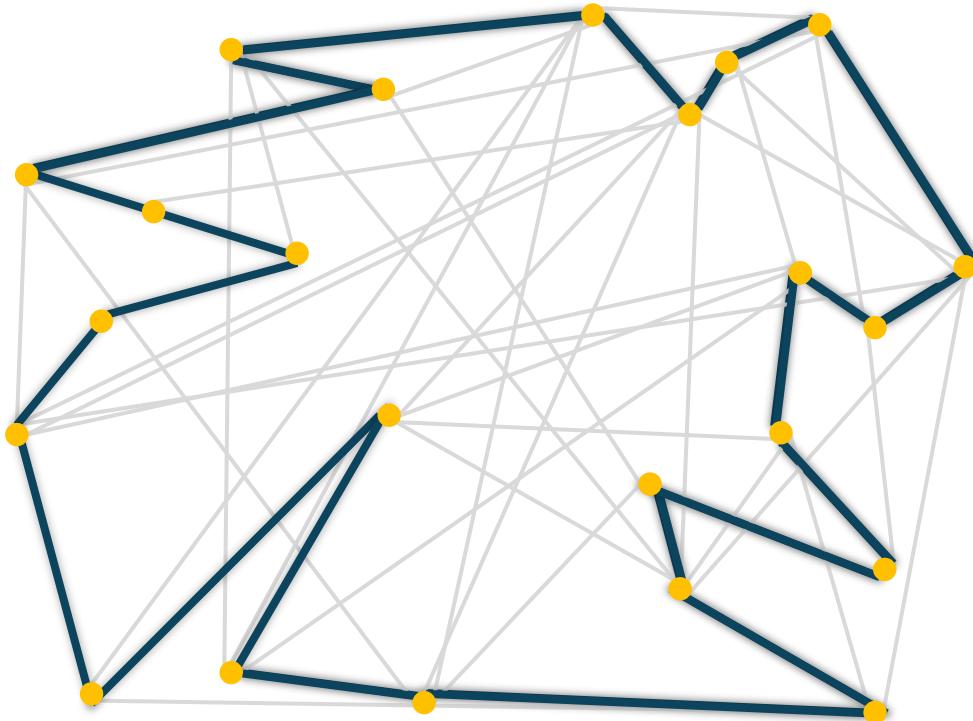
vertices V representing cities

edges E representing the paths between cities

$w(i, j)$ reflects the cost or distance between cities i and j

The objective is to find the shortest possible route that visits each city exactly once and returns to the starting point

Traveling Salesman Problem



NP-Hard



Large-Scale Traveling Salesman Problem

Exact Methods

find **provably optimal solutions** through exhaustive or branch-and-bound searches, albeit often with prohibitive computational costs.

Algorithms: MILP, Constraint Programming, cut-and-price, Concorde, etc.

Heuristic Methods

and (Hyper-Heuristics) Evolutionary
Machine Learning

produce **near-optimal solutions efficiently** by exploiting problem-specific insights, though without guaranteeing global optimality

Algorithms: Greedy Search, A*, Tabu Search, LNS, LKH-3, ACO, GPHH, GP-ACO, etc.

Neural Combinatorial Optimization (NCO)

Generative AI (Diffusion Models)

leverages **neural networks**, often paired with **reinforcement learning**, to learn and solution patterns from data

Algorithms: Ptr-Net, AM, POMO, GCN, GNN, Equity-Transformer, Sym-NCO, etc.

data-driven **generative models** to directly produce high-quality solutions, integrating global structural information for more holistic route designs

Algorithms: DIFUSCO, IDEQ, TSPDiffuser, DISCO, Fast t2t, SymmetricDiffusers, etc.



Large-Scale Traveling Salesman Problem

Exact Methods

Heuristic Methods
and (Hyper-Heuristics) Evolutionary
Machine Learning

Neural Combinatorial Optimization (NCO)

Generative AI (Diffusion Models)

In theory, optimal solution

High quality solution in reasonable time
Good Generalization ability

High potential to capture global
information with GNNs
High speed Inference with GPU

High quality solutions
High speed Inference with GPU

Unacceptable Computation Time

Hard to capture global information
Trade off between quality and efficiency
Low speed (improvement heuristics)

Low Generalization ability
High Computation resources

Huge Computation resources
Low Generalization ability
Huge data

Machine Learning is to automatically find a **function** from data.

Set scope

Define the set of candidate functions

**Deep Learning (CNN, Transformer...),
Decision Tree, etc.**

Establish criteria

Define the criteria for evaluating the quality of functions

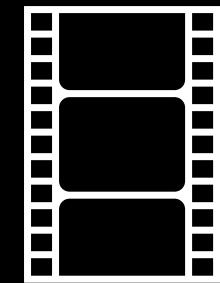
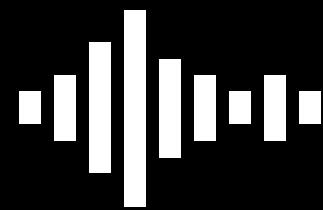
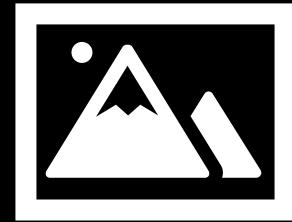
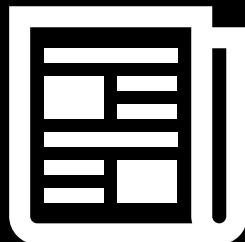
**Supervised Learning, Semi-supervised
Learning, Reinforcement Learning, etc.**

Achieve the goal

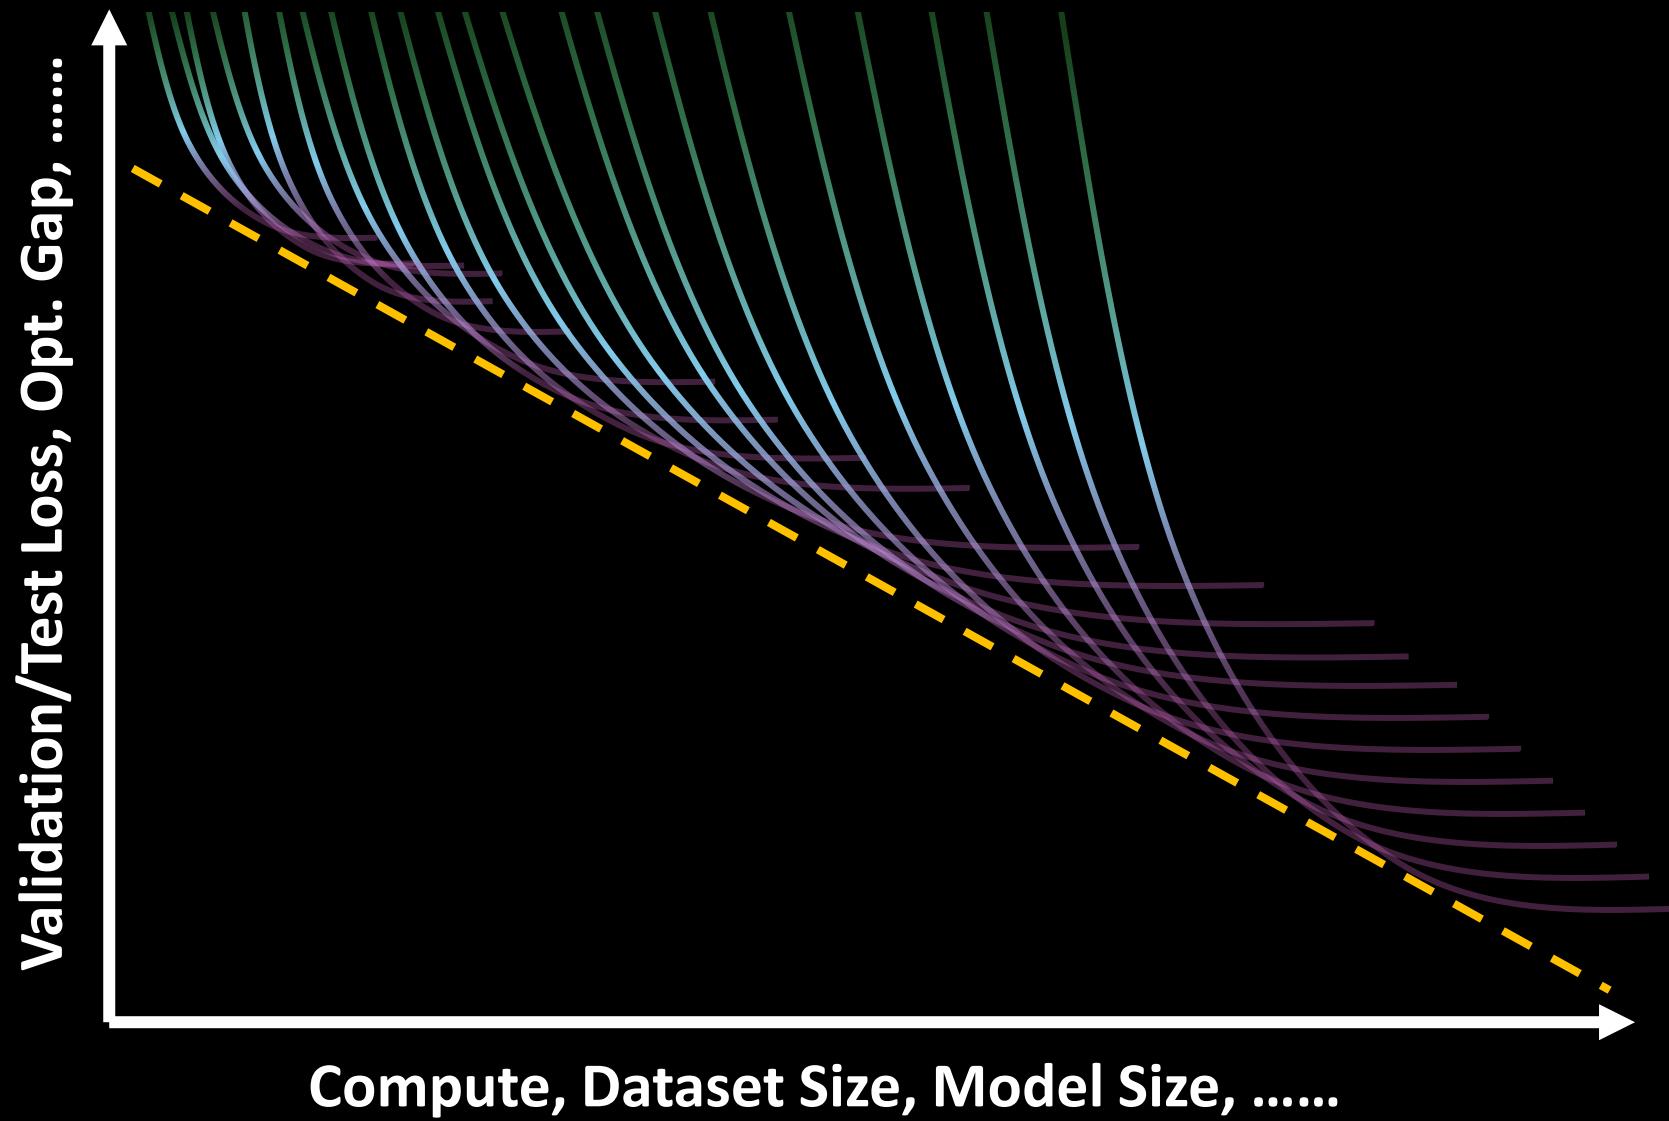
Find the best function → Optimization

Gradient Descent, Genetic Algorithm, etc.

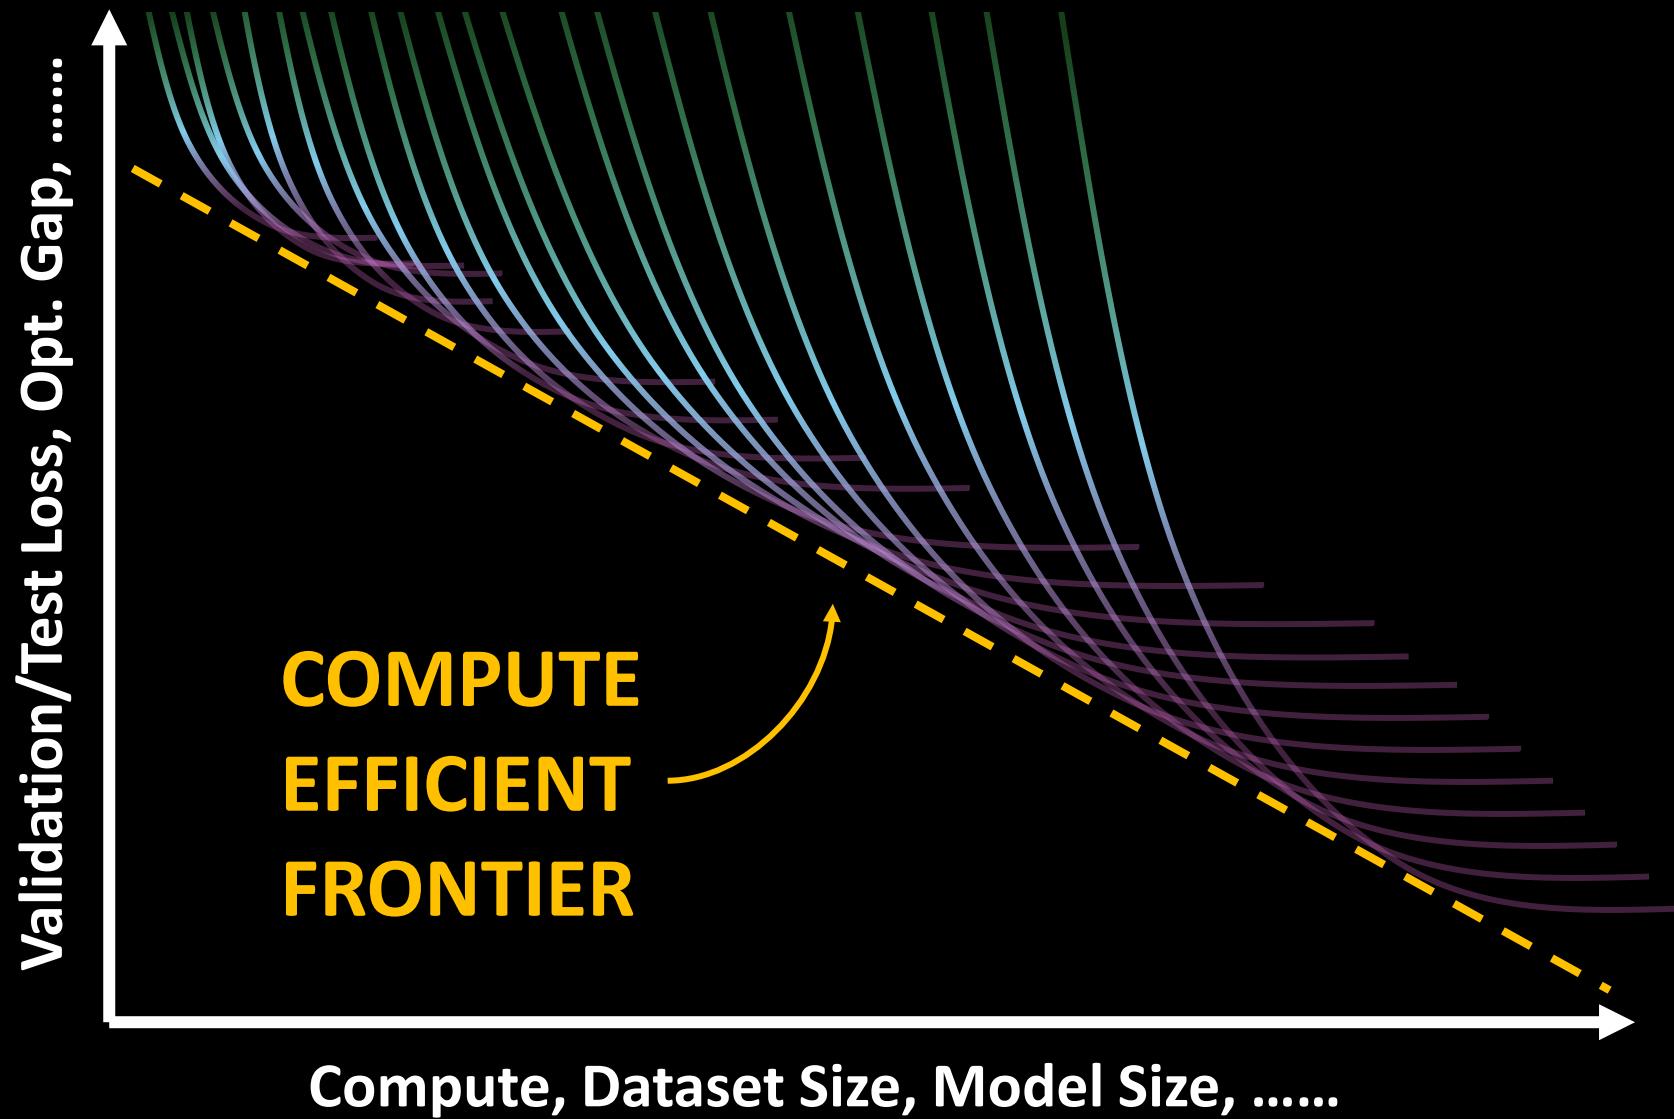
Generative AI
is to enable machines to **generate**
complex and structured things.



Scaling Laws



Scaling Laws





Large-Scale Traveling Salesman Problem

Exact Methods

Heuristic Methods
and (Hyper-Heuristics)
Evolutionary Machine Learning

Neural Combinatorial Optimization (NCO)

Generative AI (Diffusion Models)

In theory, optimal solution

High quality solution in reasonable time
Good Generalization ability

High potential to capture global
information with GNNs
High speed Inference with GPU

High quality solutions
High speed Inference with GPU

Unacceptable Computation Time

Hard to capture global information
Trade off between quality and efficiency
Low speed (improvement heuristics)

Low Generalization ability
High Computation resources

Huge Computation resources
Low Generalization ability
Huge data

1. Introduction
2. Motivations
3. Research Objectives
4. Preliminary Works
5. Proposed Contributions and Research Plan

Current Limitations

1 Insufficient Understanding of Scaling Laws

2 NCOs' Generalization Ability and Heuristics' Global Information Processing

3 Insufficient Adaptability and Customization in Generative AI and Heuristic Methods

4 Insufficient Research on Transfer Learning



Current Limitations

1 Insufficient Understanding of Scaling Laws

With the problem size increases, Heuristics and NCOs need more resources
Computing (Evaluation Times), Training data, Model size ...

But, how much?

Compute Efficient Frontier?

How to allocation?



Current Limitations

2 NCOs' Generalization Ability and Heuristics' Global Information

Exact Methods

Heuristic Methods
and (Hyper-Heuristics) Evolutionary
Machine Learning

Neural Combinatorial Optimization (NCO)

Generative AI (Diffusion Models)

In theory, optimal solution

High quality solution in reasonable time
Good Generalization ability

High potential to capture global
information with GNNs
High speed Inference with GPU

High quality solutions
High speed Inference with GPU

Unacceptable Computation Time

Hard to capture global information
Trade off between quality and efficiency
Low speed (improvement heuristics)

Low Generalization ability
High Computation resources

Huge Computation resources
Low Generalization ability
Huge data



Current Limitations

3 Insufficient Customization in Generative AI and Heuristics

Generative AI (Diffusion Models) *end-to-end*
cannot adjust model according to specific instance
high computation cost of inference

Heuristics *fixed strategy*
cannot adjust strategy according to specific instance
low computation cost (constructive heuristics)



Current Limitations

4 Insufficient Research on Transfer Learning

When facing new variant of TSPs

NCOs

Re-collect training data, re-train

Heuristics

Re-design heuristics manually

1. Introduction
2. Motivations
- 3. Research Objectives**
4. Preliminary Works
5. Proposed Contributions and Research Plan

Overall Goal

to develop **scalable** and **efficient** solutions for large-scale TSP
by leveraging machine learning and generative AI techniques

Objective 1

Exploring **Scaling Laws** in Solving Large-Scale TSP



Current Limitations

1 Insufficient Understanding of Scaling Laws

With the problem size increases, Heuristics and NCOs need more resources
Computing (Evaluation Times), Training data, Model size ...

But, how much? **Prediction**

Compute Efficient Frontier?

How to allocation?

Objective 1

Scaling Laws

1.1 Analysis the Scaling Behavior of **Heuristic Methods**

1.2 Analysis the Scaling Behavior of **NCO Methods**

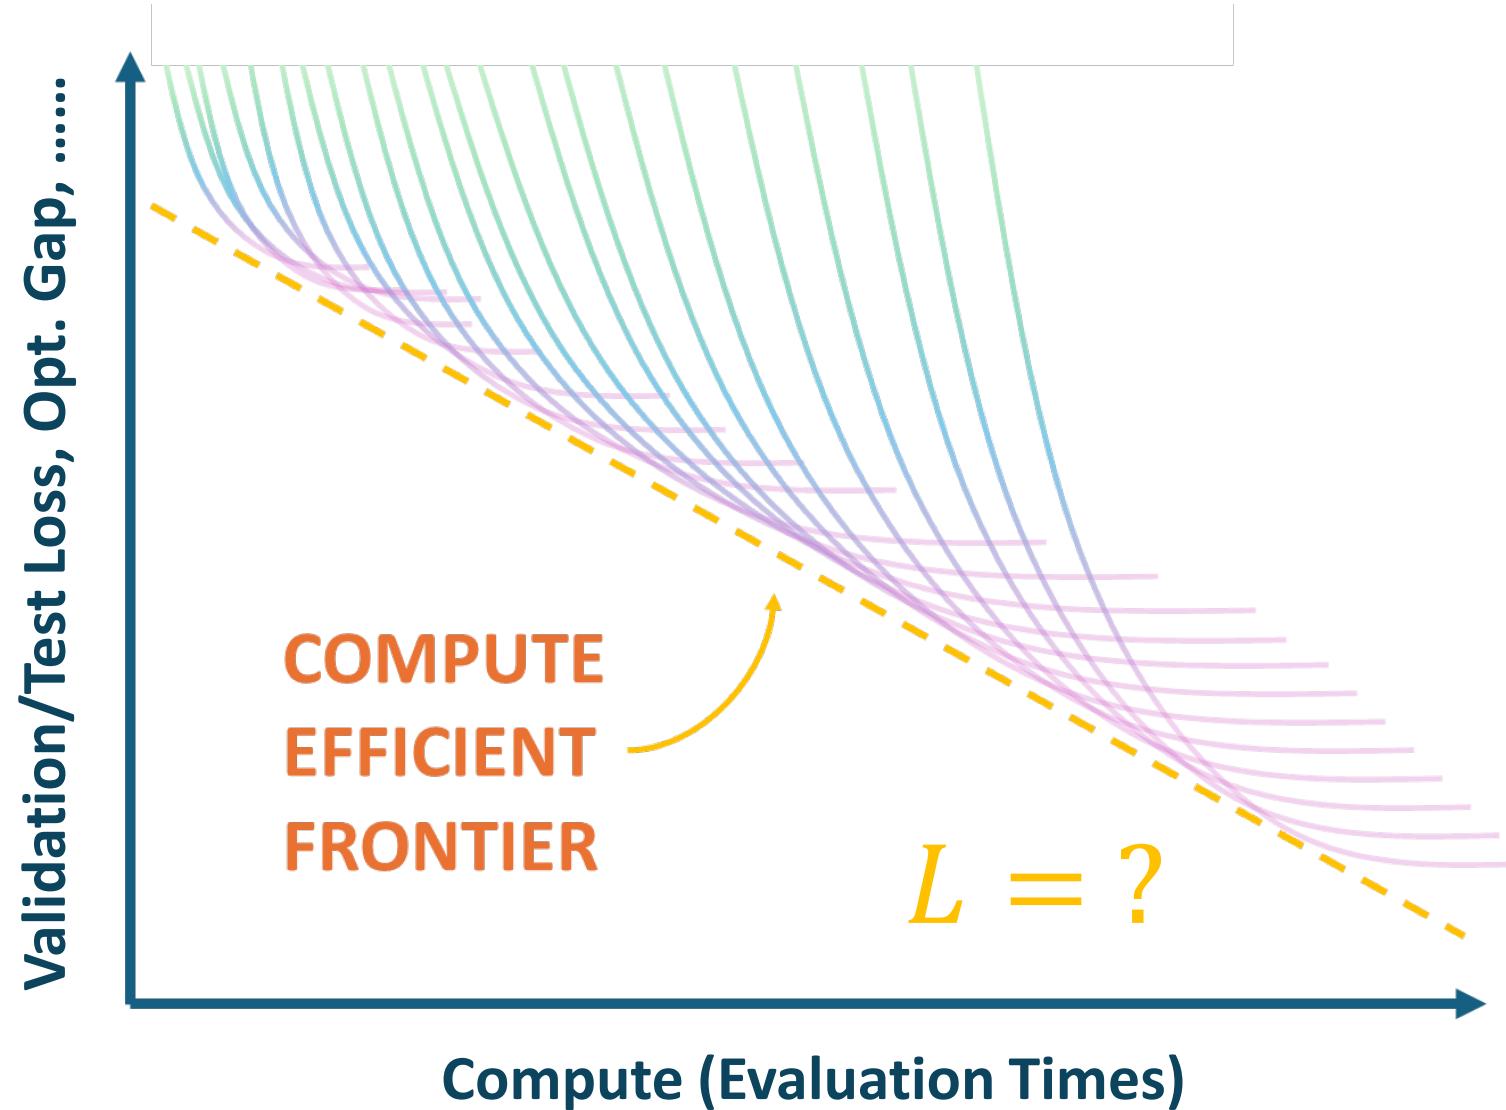
1.3 Formulate Scaling Laws to Guide **Resource Allocation**

Optional

1.4 Develop **Theoretical Models** to Explain Observed Scaling Behaviors

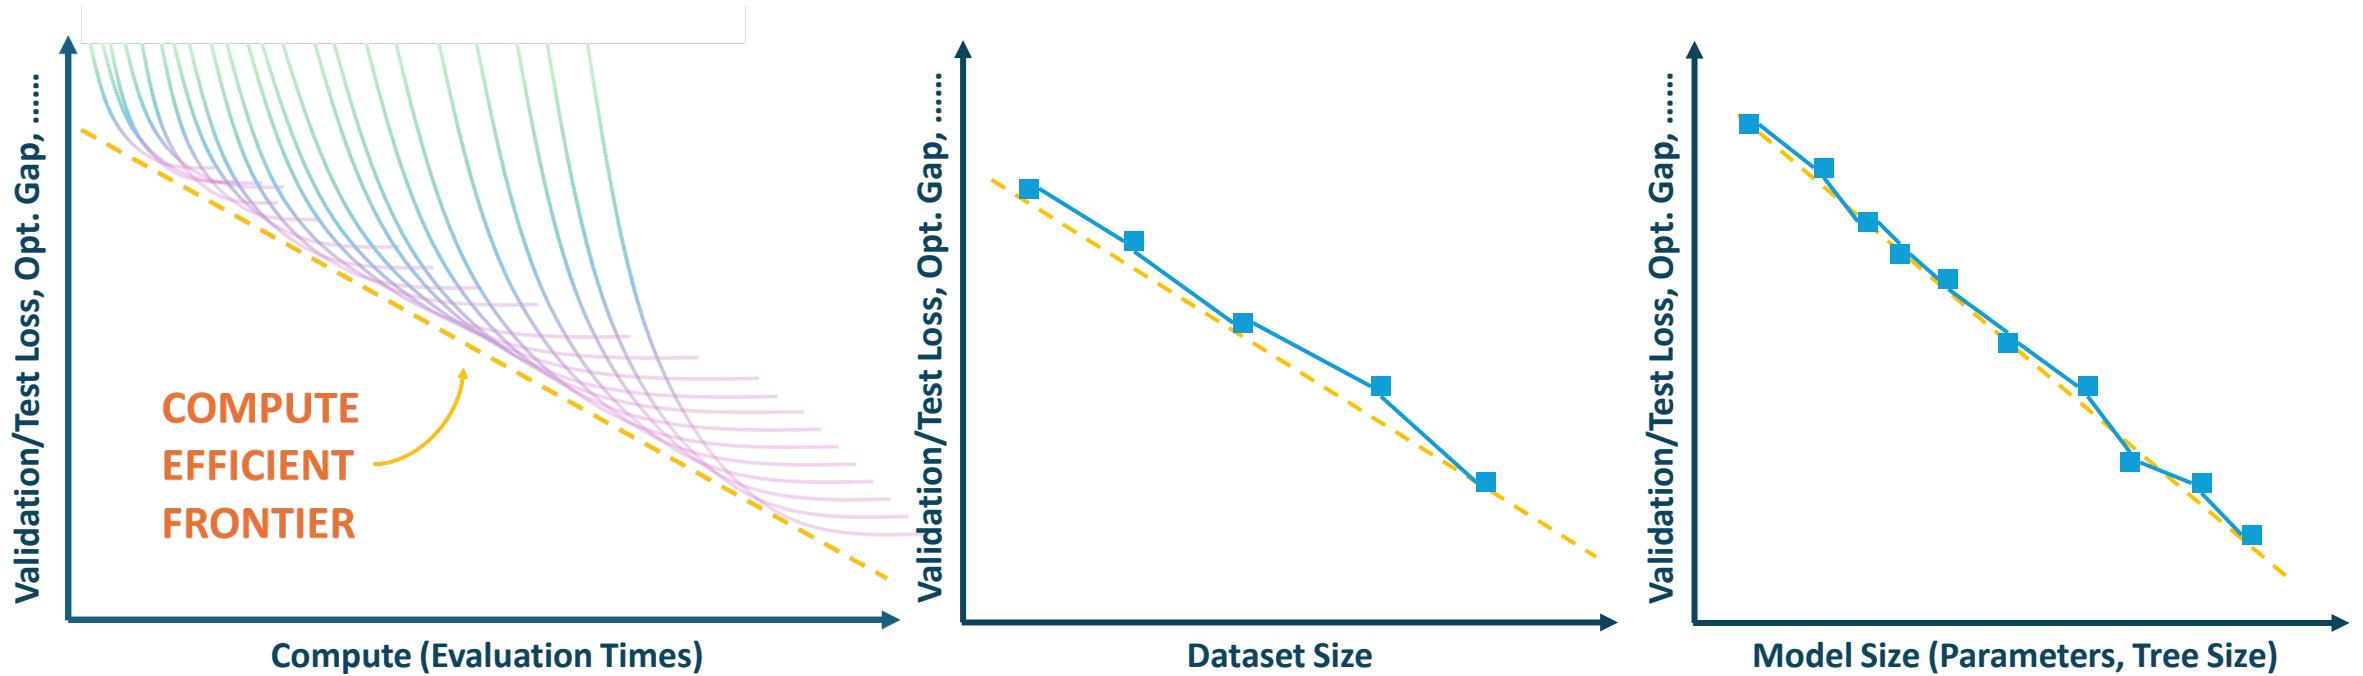
Objective 1

1.1 & 1.2 Analysis the Scaling Behavior of Heuristic & NCO Methods



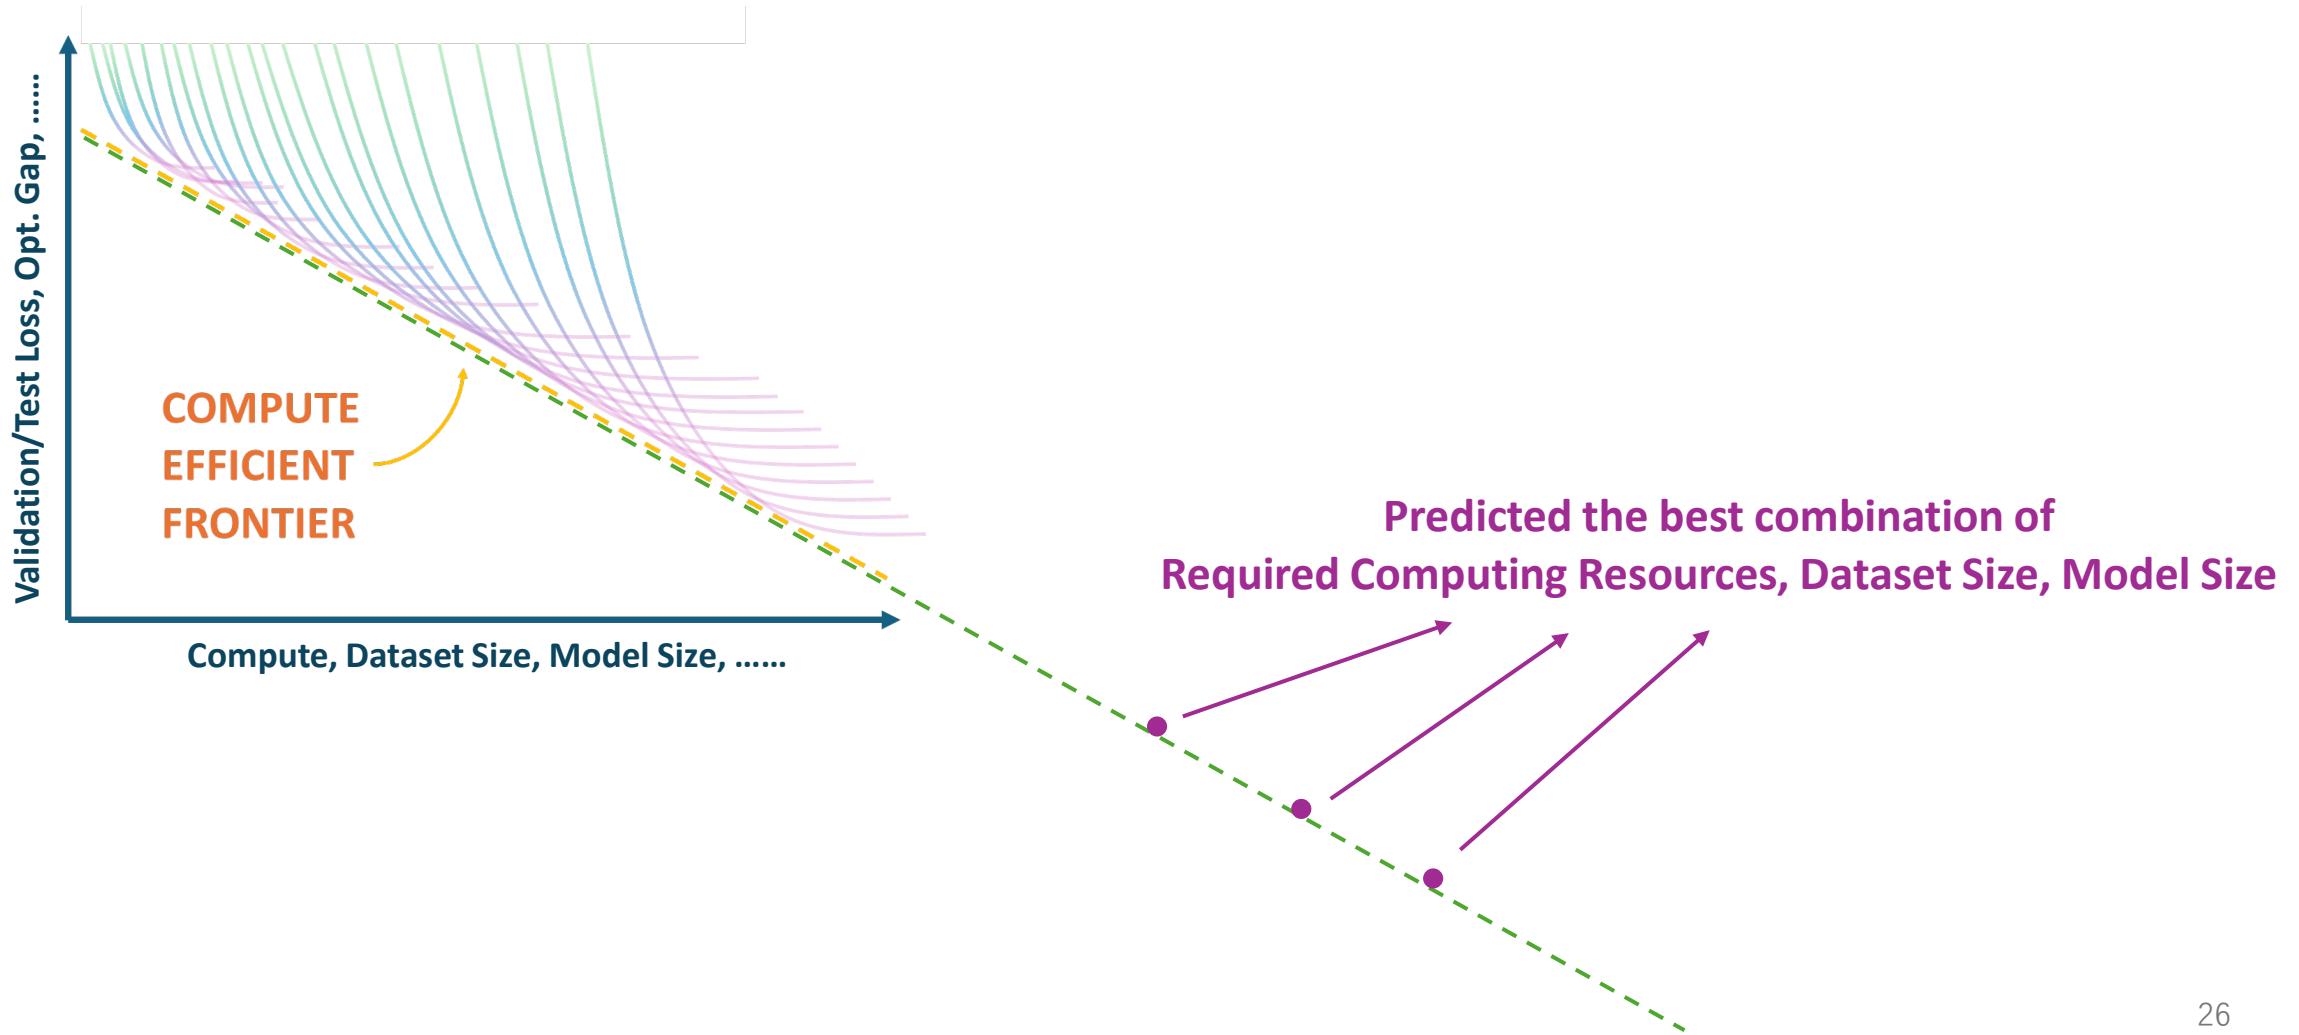
Objective 1

1.1 & 1.2 Analysis the Scaling Behavior of Heuristic & NCO Methods



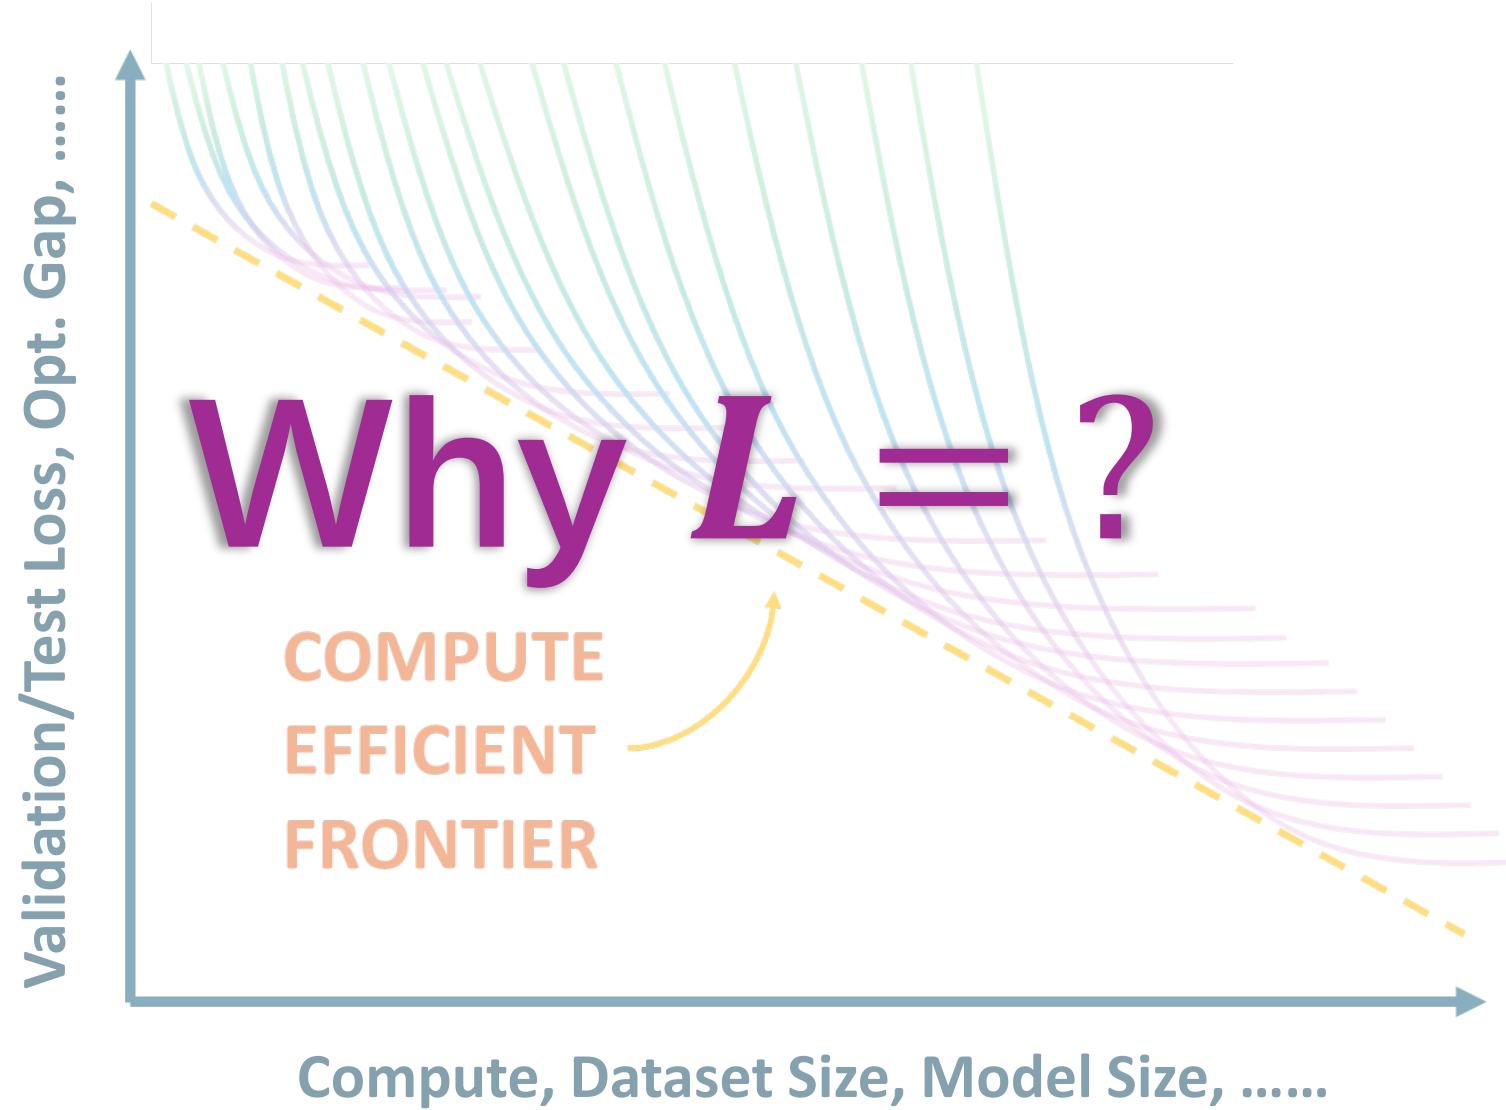
Objective 1

1.3 Formulate Scaling Laws to Guide Resource Allocation



Objective 1

1.4 Develop Theoretical Models to Explain Observed Scaling Behaviors



Objective 2

Integrating **NCO** with **Hyper-Heuristic** Approaches



Current Limitations

2 NCOs' Generalization Ability and Heuristics' Global Information

Heuristic Methods
and (Hyper-Heuristics)
Evolutionary Machine Learning

**Neural Combinatorial
Optimization (NCO)**

High quality solution in reasonable time
Good Generalization ability

High potential to capture global
information with GNNs
High speed Inference with GPU

Hard to capture global information
Trade off between quality and efficiency
Low speed (improvement heuristics)

Low Generalization ability
High Computation resources



Current Limitations

2 NCOs' Generalization Ability and Heuristics' Global Information

Heuristic Methods
and (Hyper-Heuristics)
Evolutionary Machine Learning

**Neural Combinatorial
Optimization (NCO)**

High quality solution in reasonable time

Good Generalization ability

High potential to capture global
information with GNNs
High speed Inference with GPU

Hard to capture global information
Trade off between quality and efficiency
Low speed (improvement heuristics)

Low Generalization ability
High Computation resources



Current Limitations

2 NCOs' Generalization Ability and Heuristics' Global Information

Heuristic Methods
and (Hyper-Heuristics)
Evolutionary Machine Learning

**Neural Combinatorial
Optimization (NCO)**

High quality solution in reasonable time
Good Generalization ability

High potential to capture global
information with GNNs
High speed Inference with GPU

Hard to capture global information
Trade off between quality and efficiency
Low speed (improvement heuristics)

Low Generalization ability
High Computation resources

Objective 2

NCO & Hyper-Heuristics

2.1 Enhance the **Function Set** of GP

2.2 Optimize **Sampling Strategy** in Heuristic Search

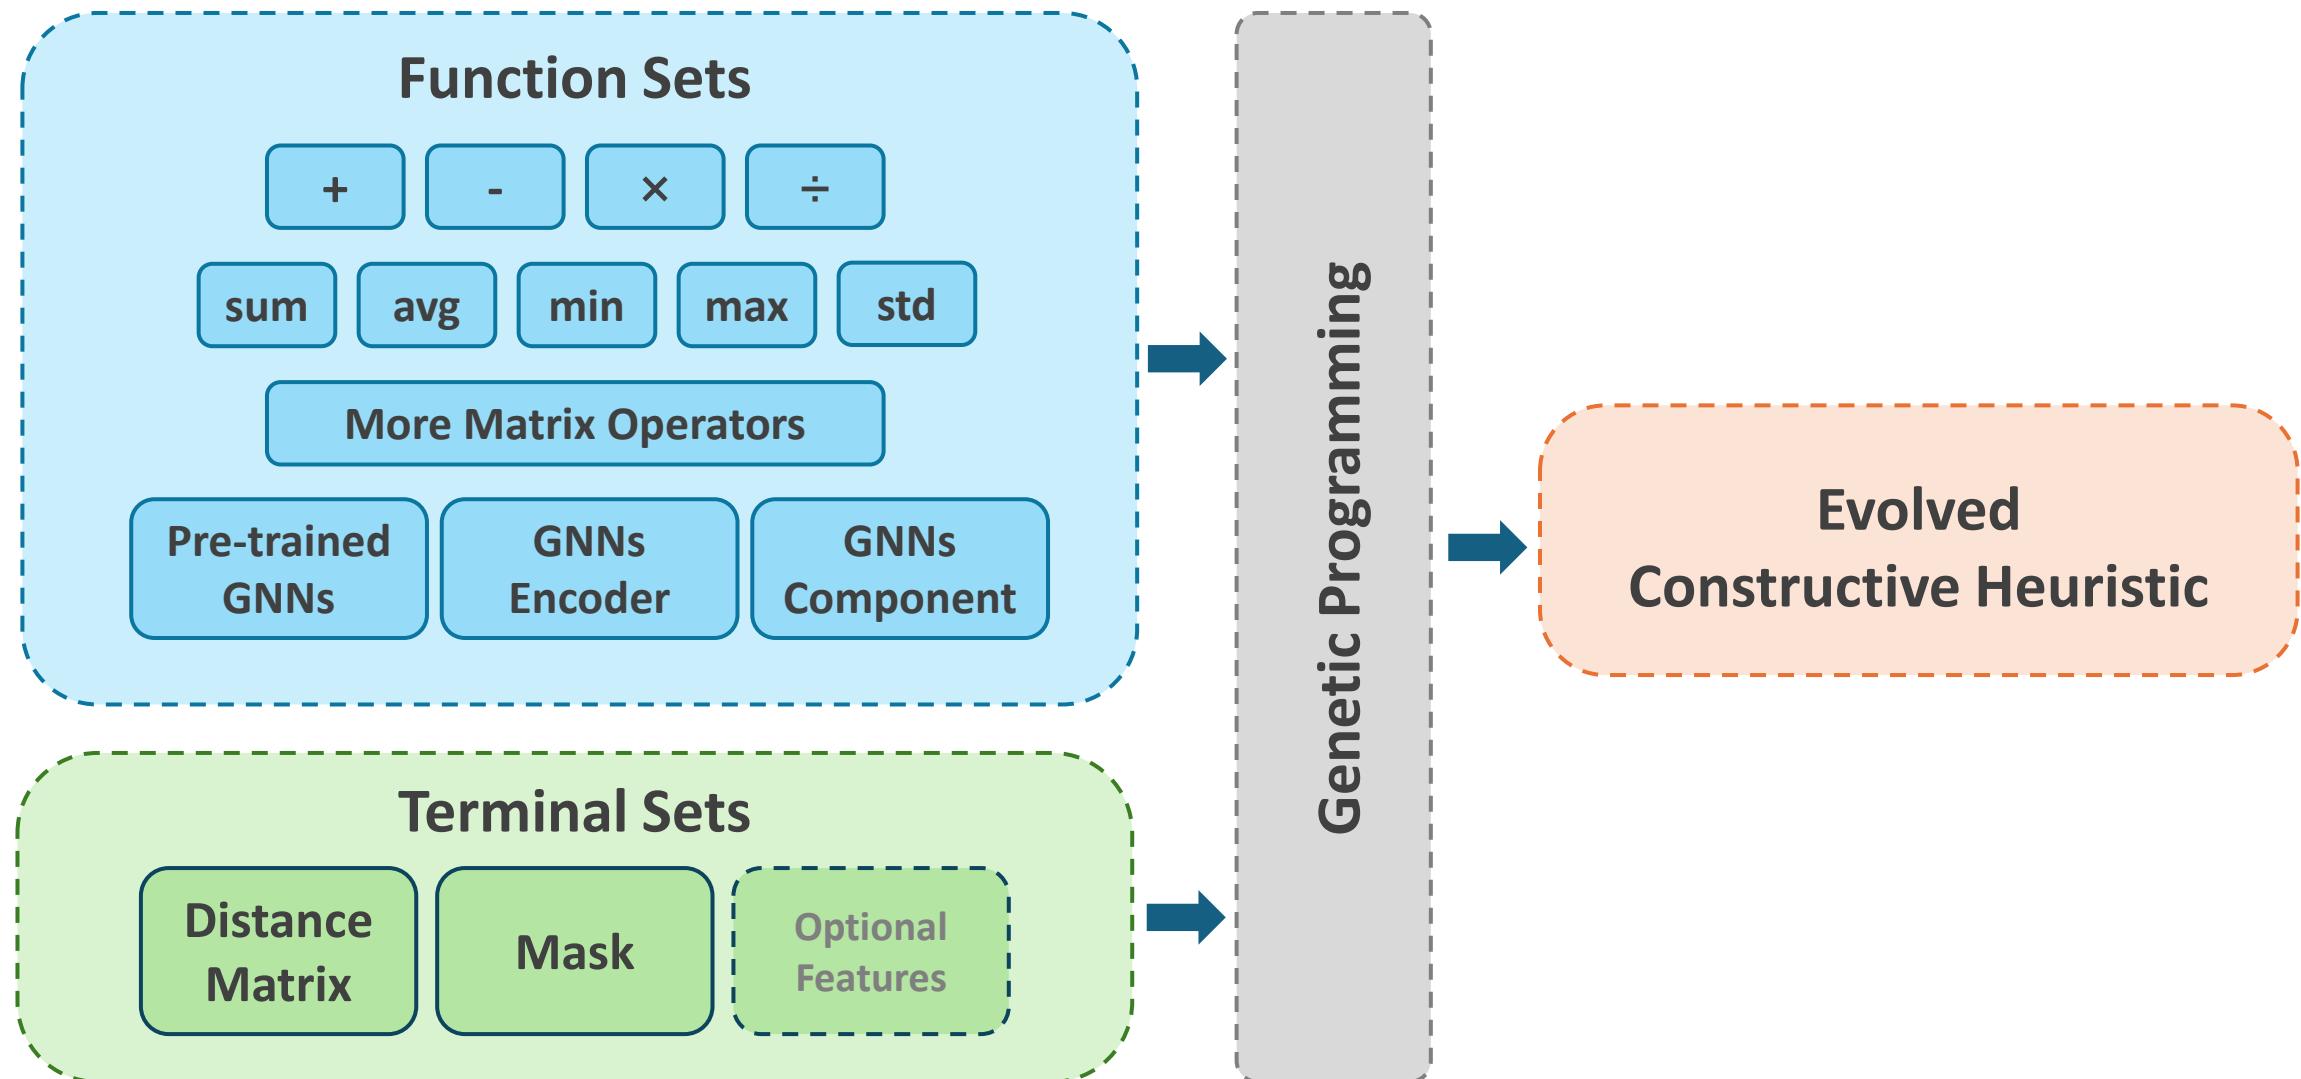
2.3 Leveraging **Historical Search Data** to Improve Performance

for Objective 3

2.4 Prepare Structure Data to Support **Generative AI** Integration

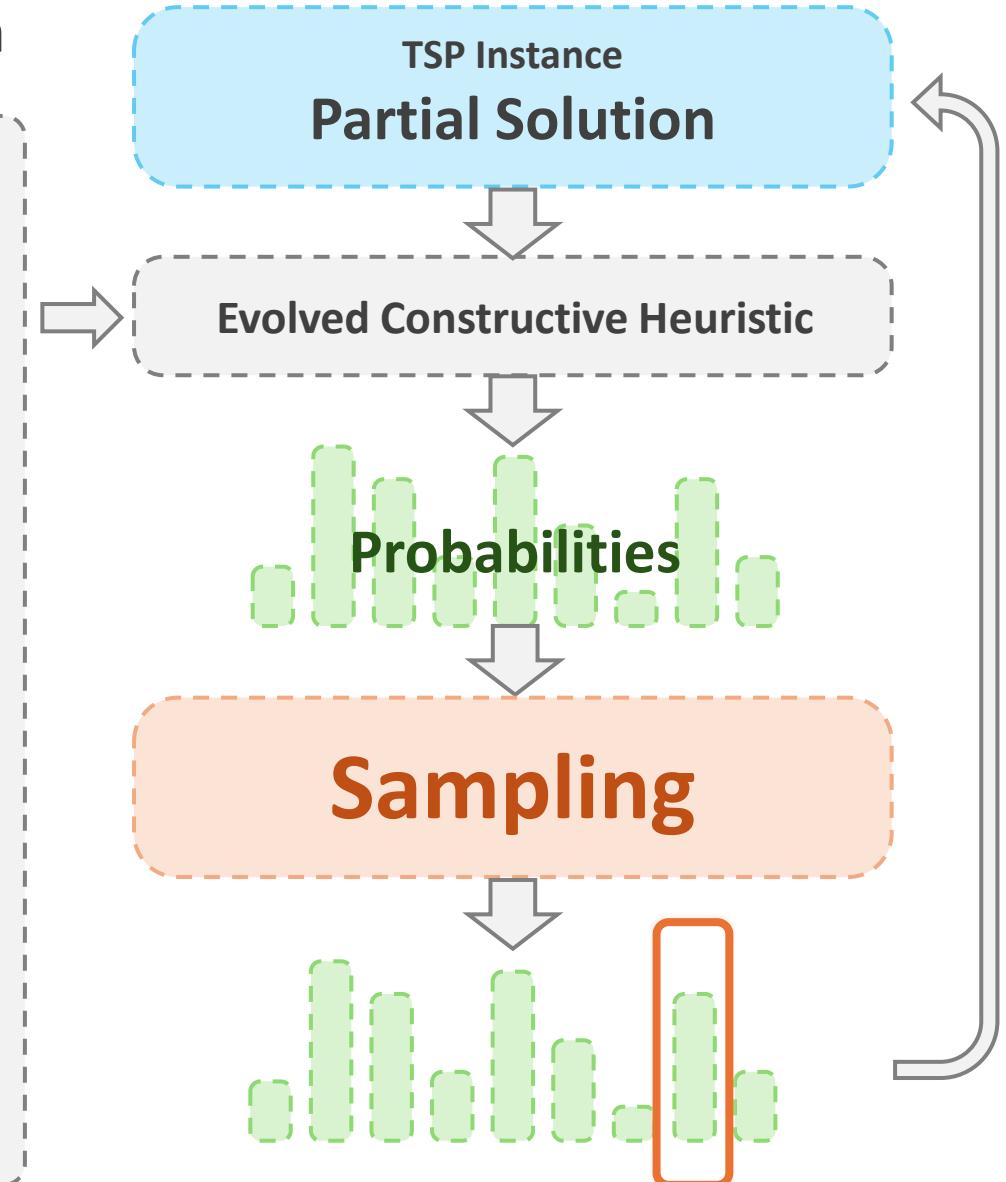
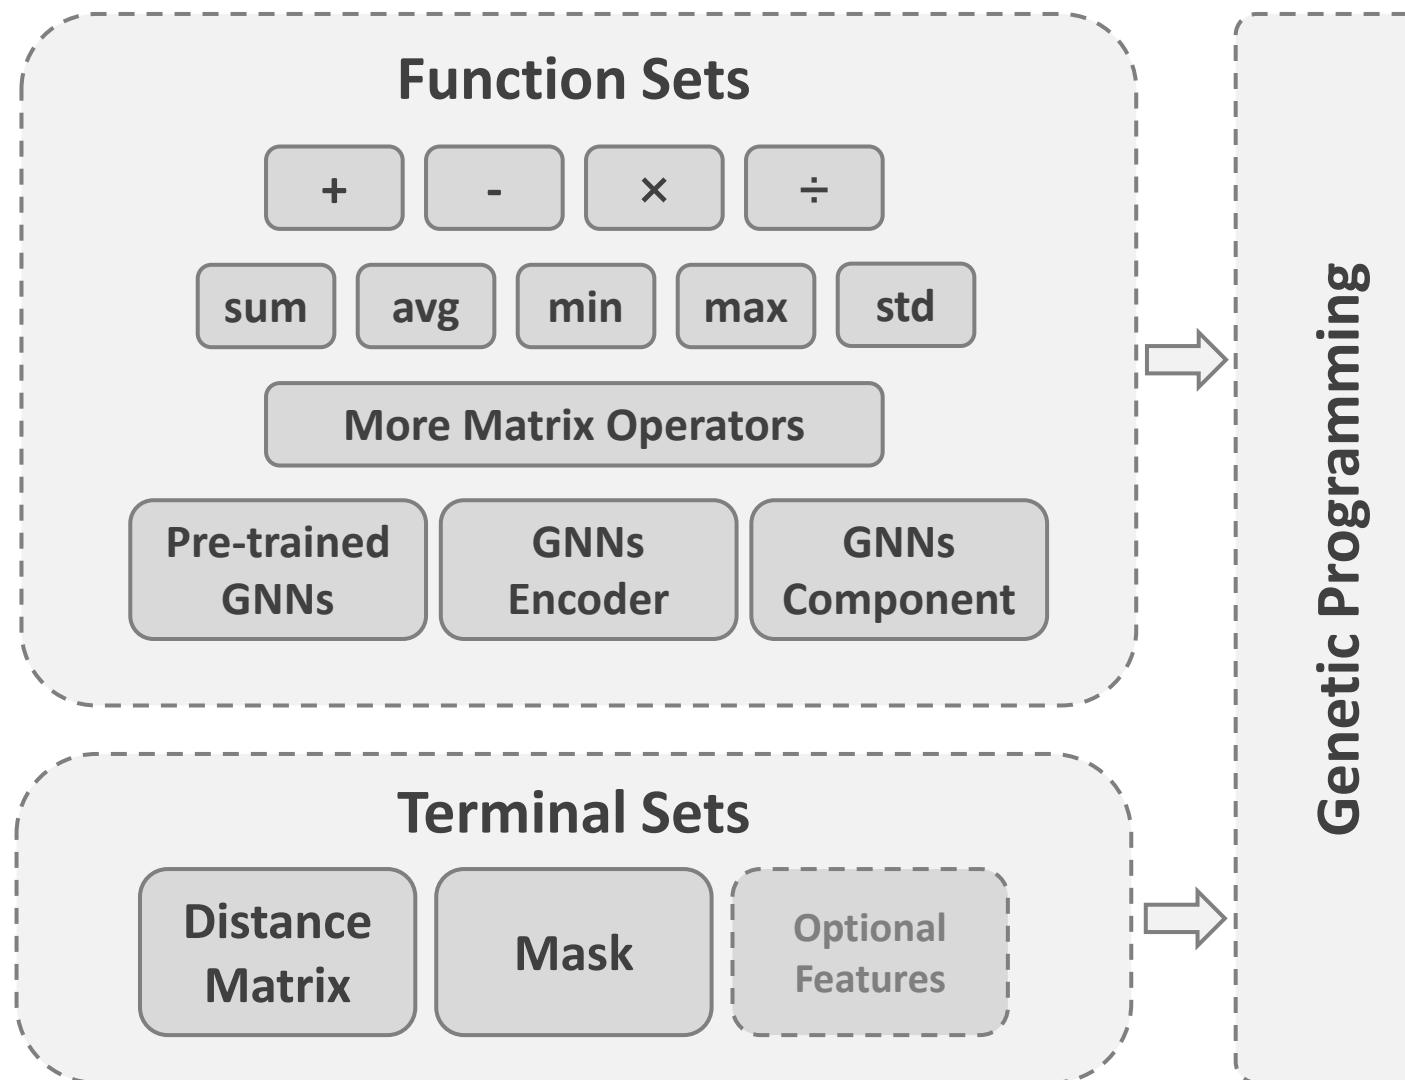
Objective 2

2.1 Enhance the Function Set of GP



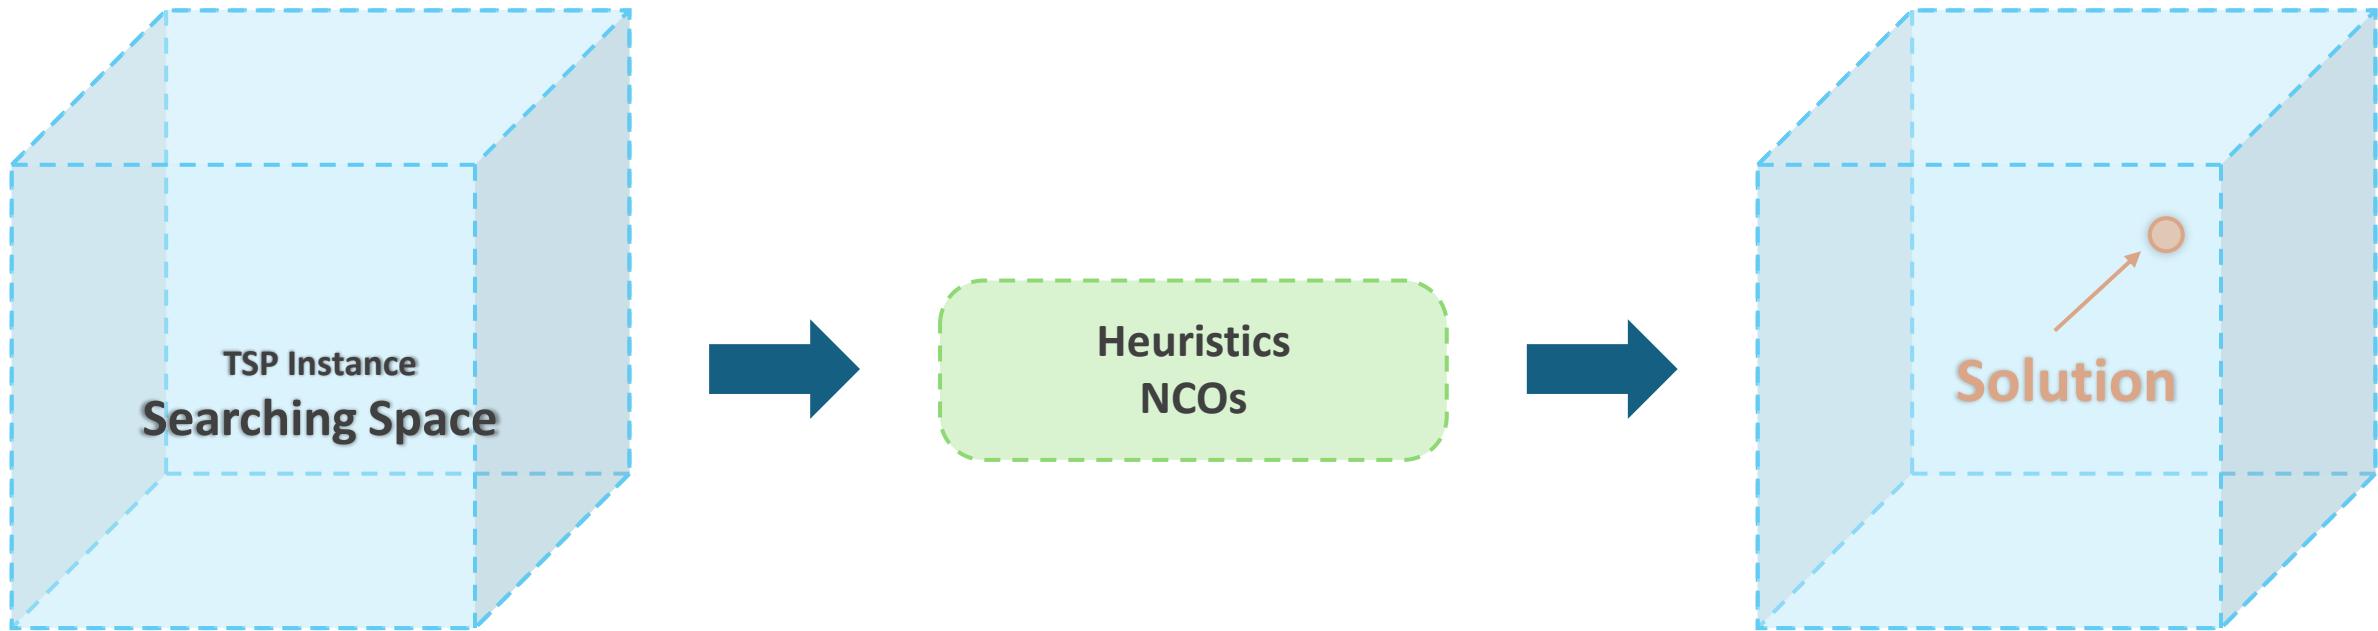
Objective 2

2.2 Optimize Sampling Strategy in Heuristic Search



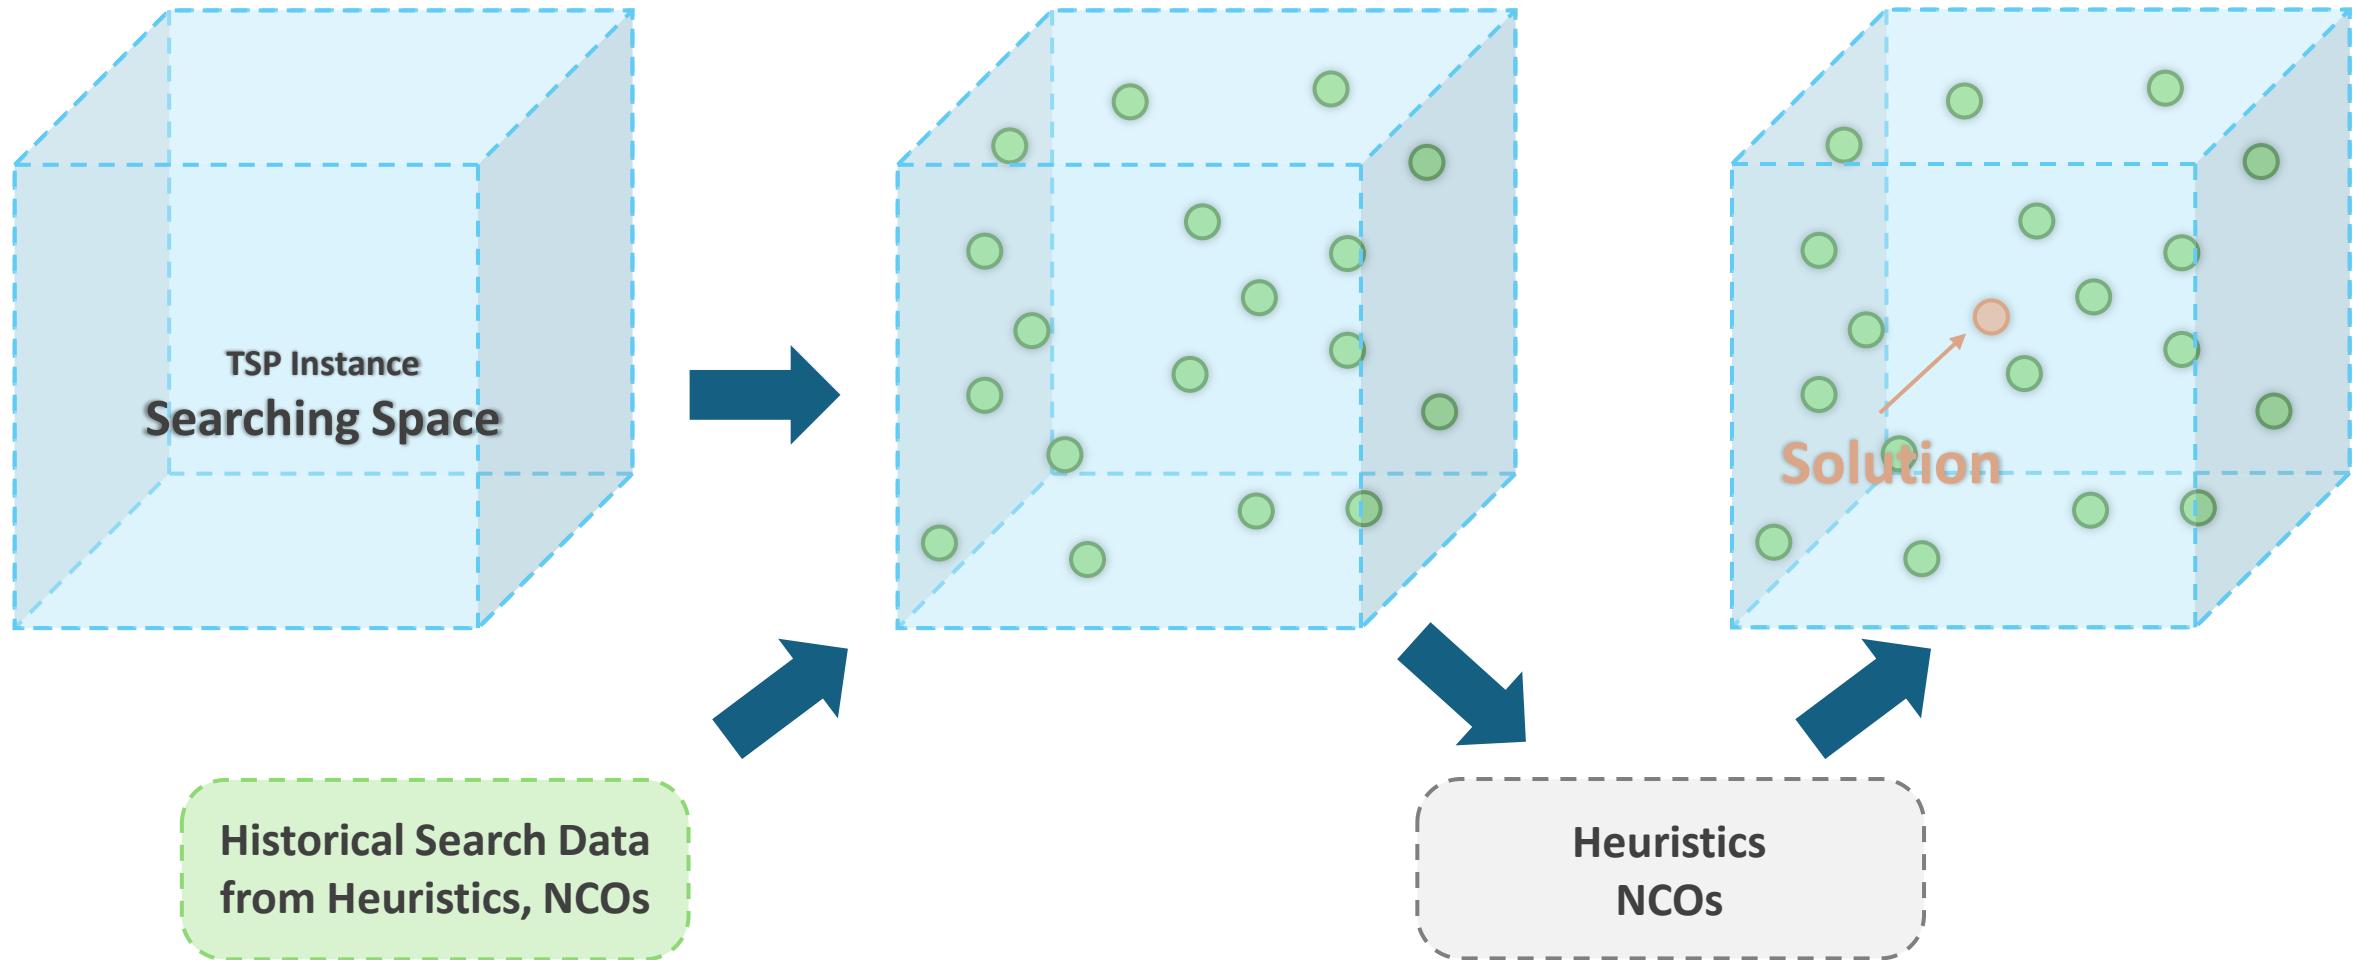
Objective 2

2.3 Leveraging Historical Search Data to Improve Performance



Objective 2

2.3 Leveraging Historical Search Data to Improve Performance



Objective 2

2.4 Prepare Structure Data to Support Generative AI Integration

Paired Data [0]

TSP Instance [0]

Compiled from GP individual
Heuristic [0]

Solution [0]

Gap [0] %

...

Paired Data [1]

TSP Instance [1]

Compiled from GP individual
Heuristic [1]

Solution [1]

Gap [1] %

...

...

Paired Data [i]

TSP Instance [i]

Compiled from GP individual
Heuristic [i]

Solution [i]

Gap [i] %

...

...

Paired Data [n]

TSP Instance [n]

Compiled from GP individual
Heuristic [n]

Solution [n]

Gap [n] %

...

Objective 3

Develop **Generative AI** Models for Heuristic Generation



Current Limitations

3 Insufficient Customization in Generative AI and Heuristics

Generative AI (Diffusion Models) *end-to-end*
cannot adjust model according to specific instance
high computation cost (intermediate)

Generate Heuristics

Heuristics *fixed strategy*
cannot adjust strategy according to specific instance
low computation cost (constructive heuristics)

Objective 3

Generative AI for Heuristic Generation

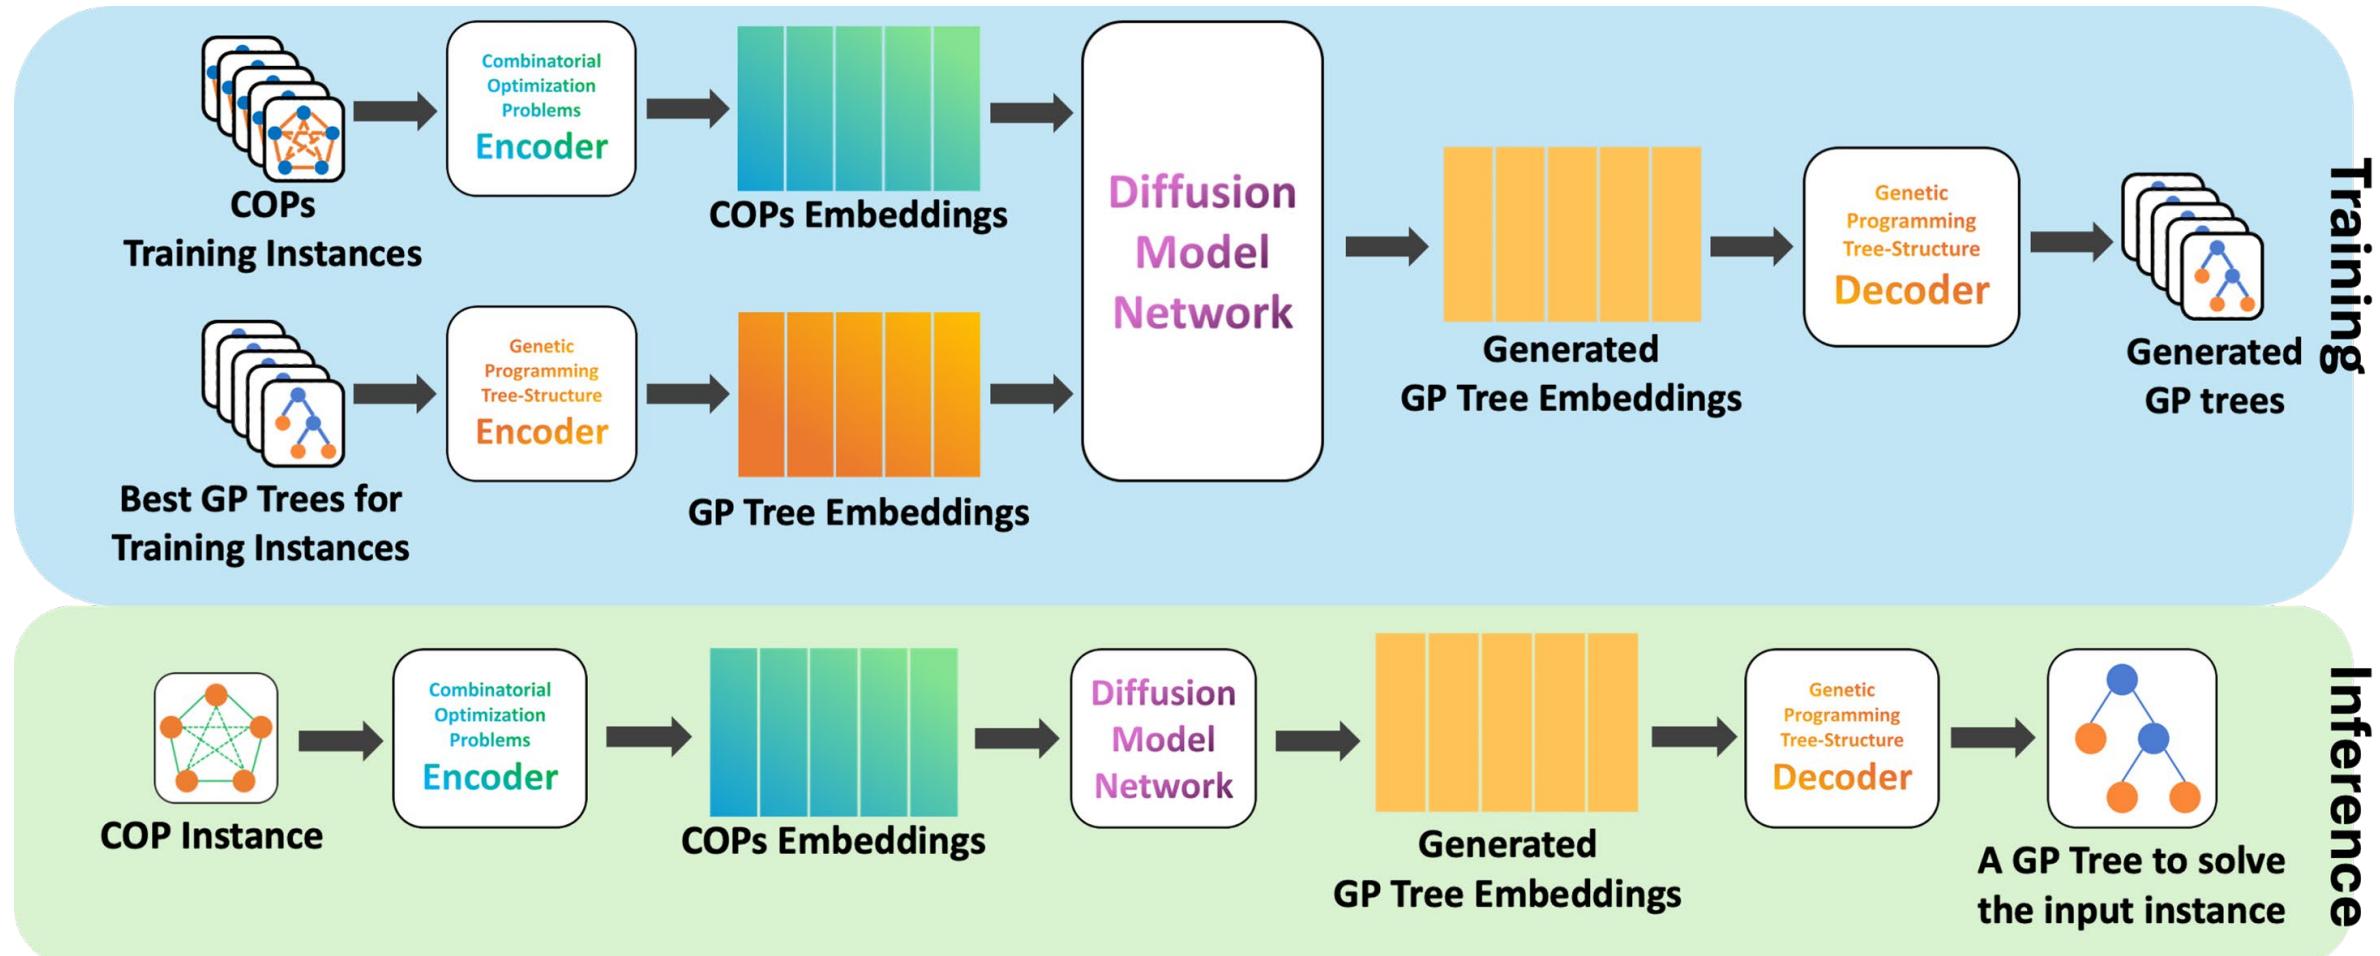
3.1 Diffusion Models for Instance-level Heuristic Generation

3.2 Explore Self-Supervised Training Methods

3.3 Analysis Limitations and Explore Improvement Methods

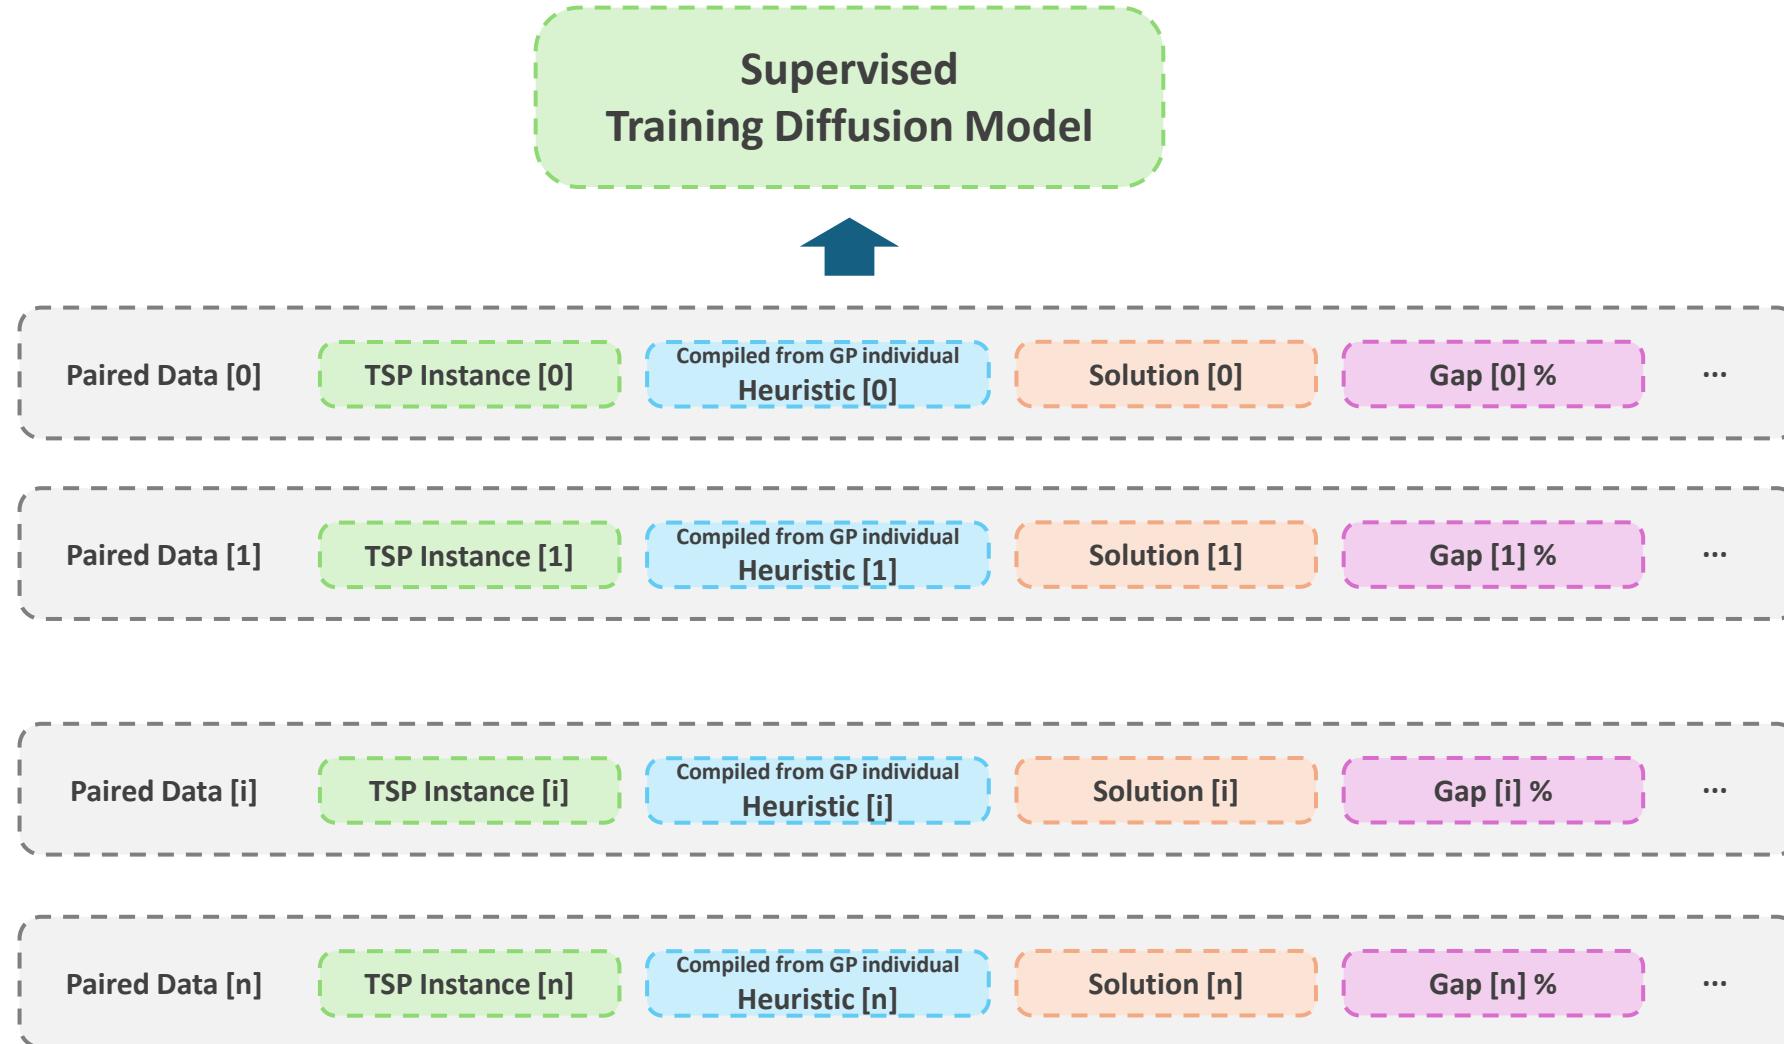
Objective 3

3.1 Diffusion Models for Instance-level Heuristic Generation



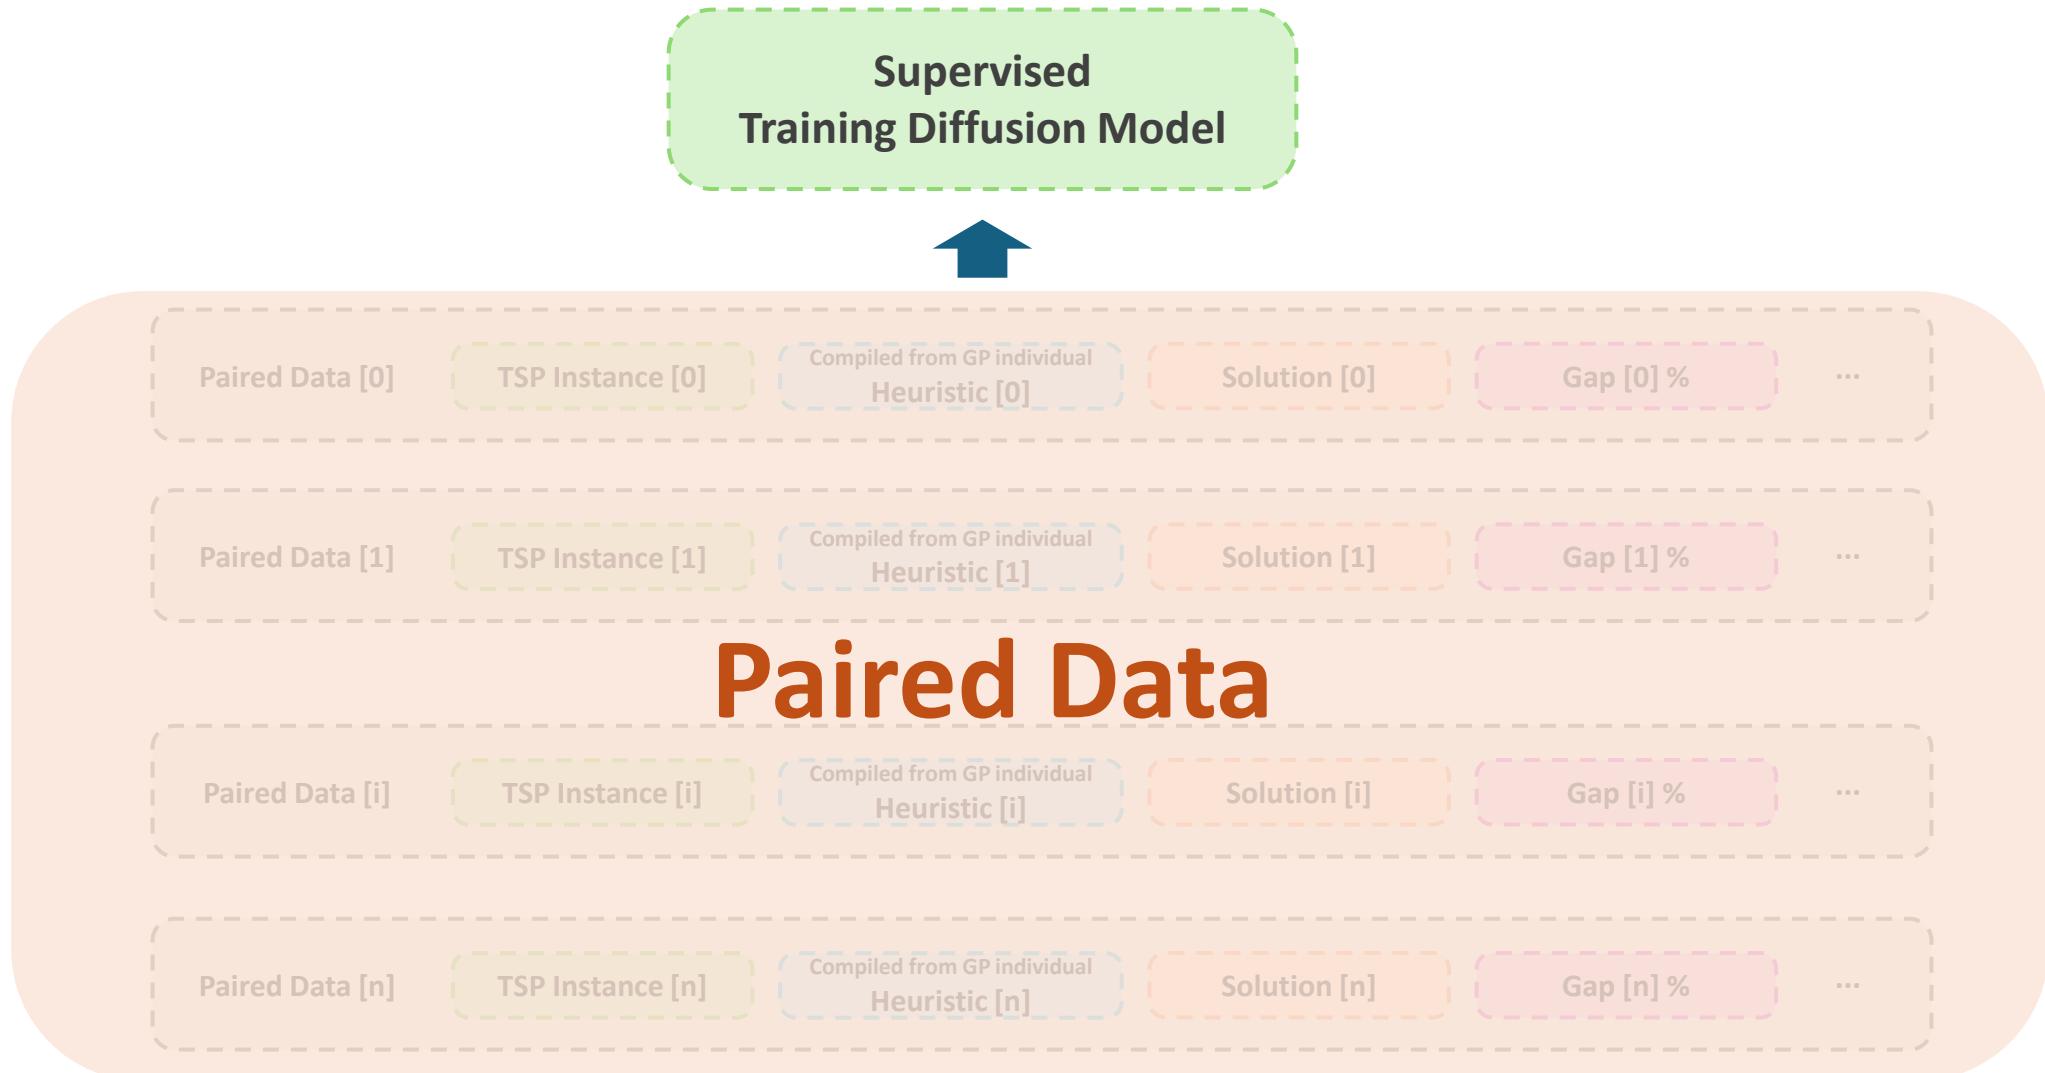
Objective 3

3.2 Explore Self-Supervised Training Methods



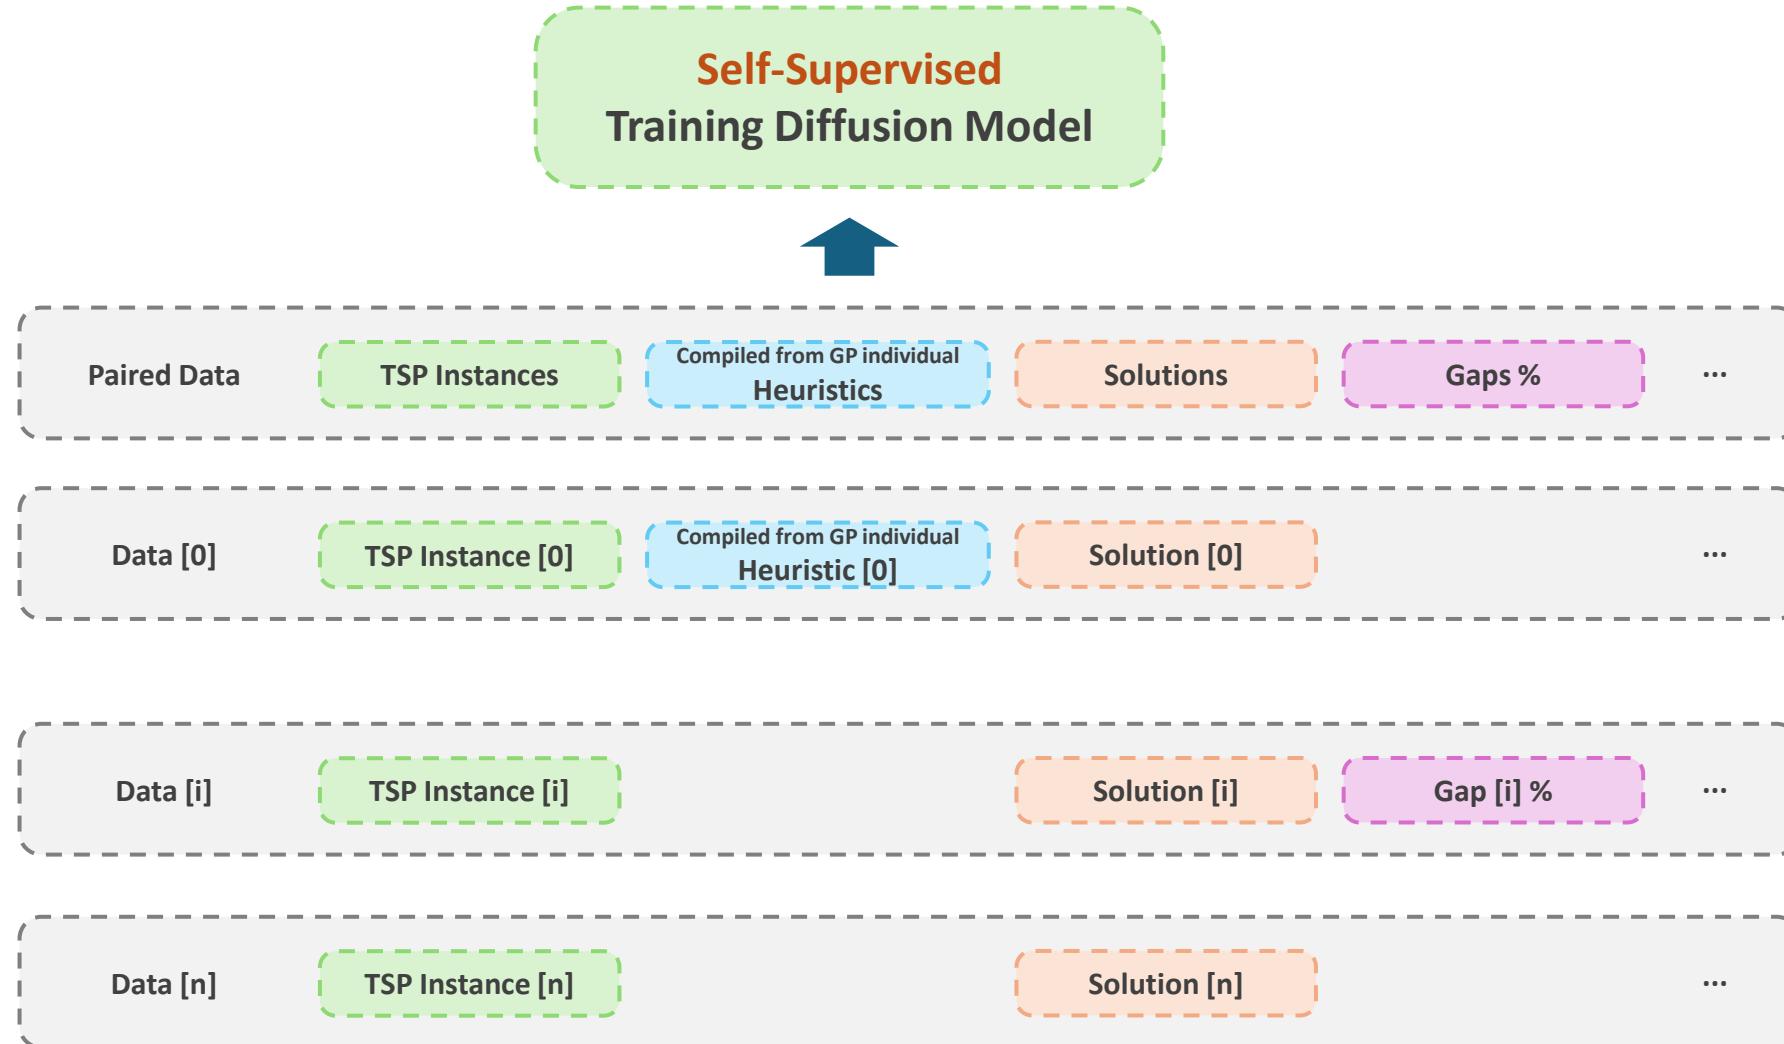
Objective 3

3.2 Explore Self-Supervised Training Methods



Objective 3

3.2 Explore Self-Supervised Training Methods



Optional
Objective 4

Enhancing GP and NCO Methods on Large-Scale and Variant TSP Problems
through **Transfer Learning** and **Domain Adaptation**



Current Limitations

4 Insufficient Research on Transfer Learning

When facing new variant of TSPs
NCOs **Focusing on differences
between variants**
Re-collect training data, re-train

Heuristics

Re-design heuristics manually

Optional Objective 4

Transfer Learning and Domain Adaptation

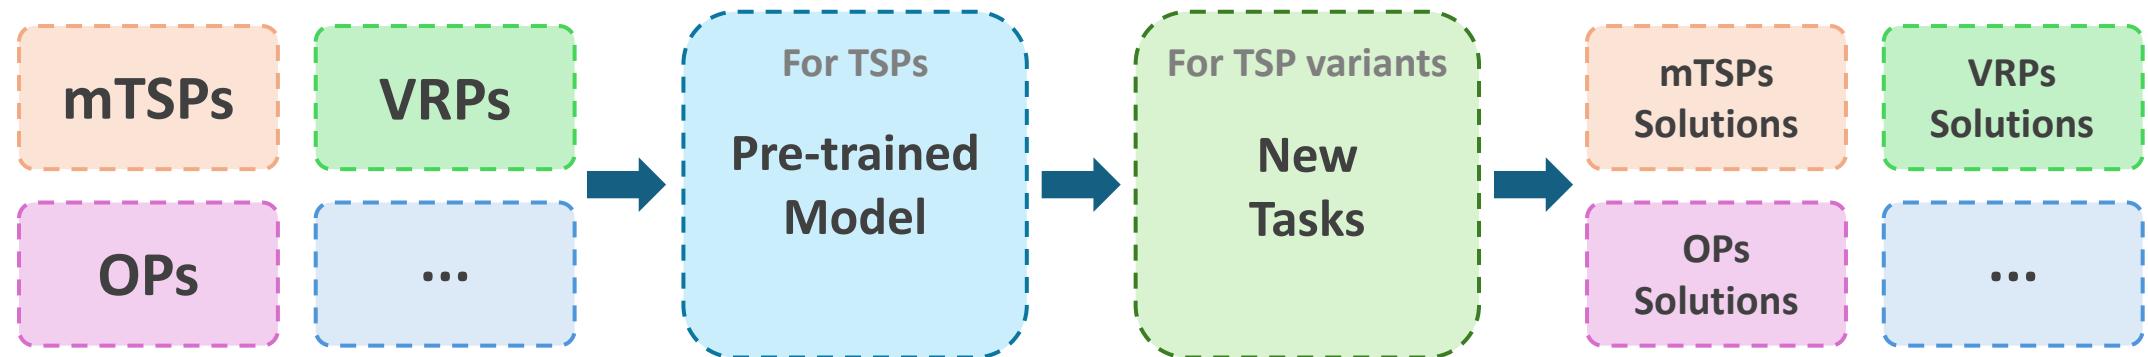
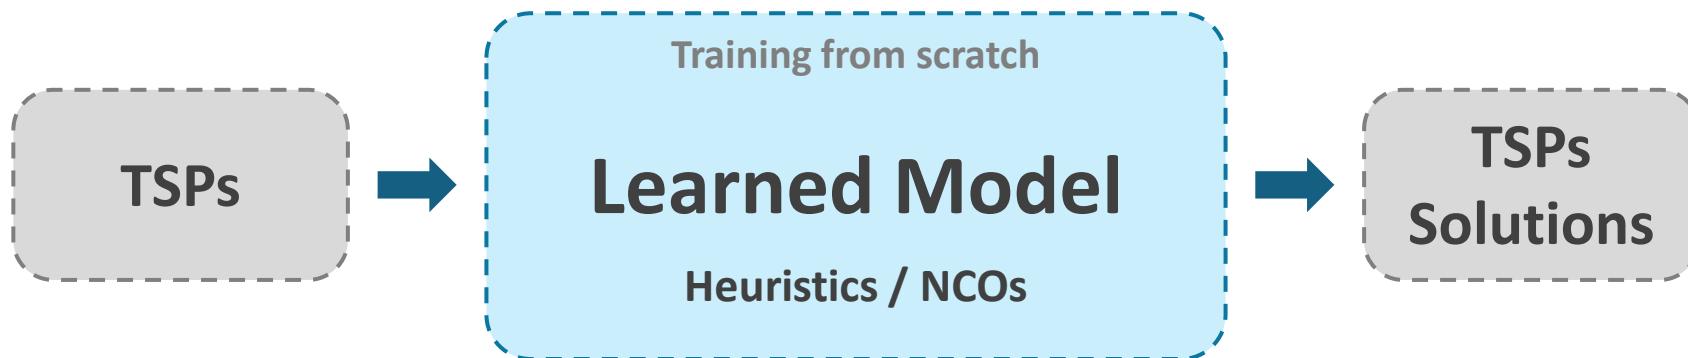
4.1 Develop a **Transfer Learning** Framework for GP and NCO

4.2 **Domain Adaptation** for Variant Problems

4.3 Transfer Learning for **Large-Scale & Variant** Problems

Optional Objective 4

Enhancing GP and Neural NCO Methods on Large-Scale and Variant TSP Problems through Transfer Learning and Domain Adaptation

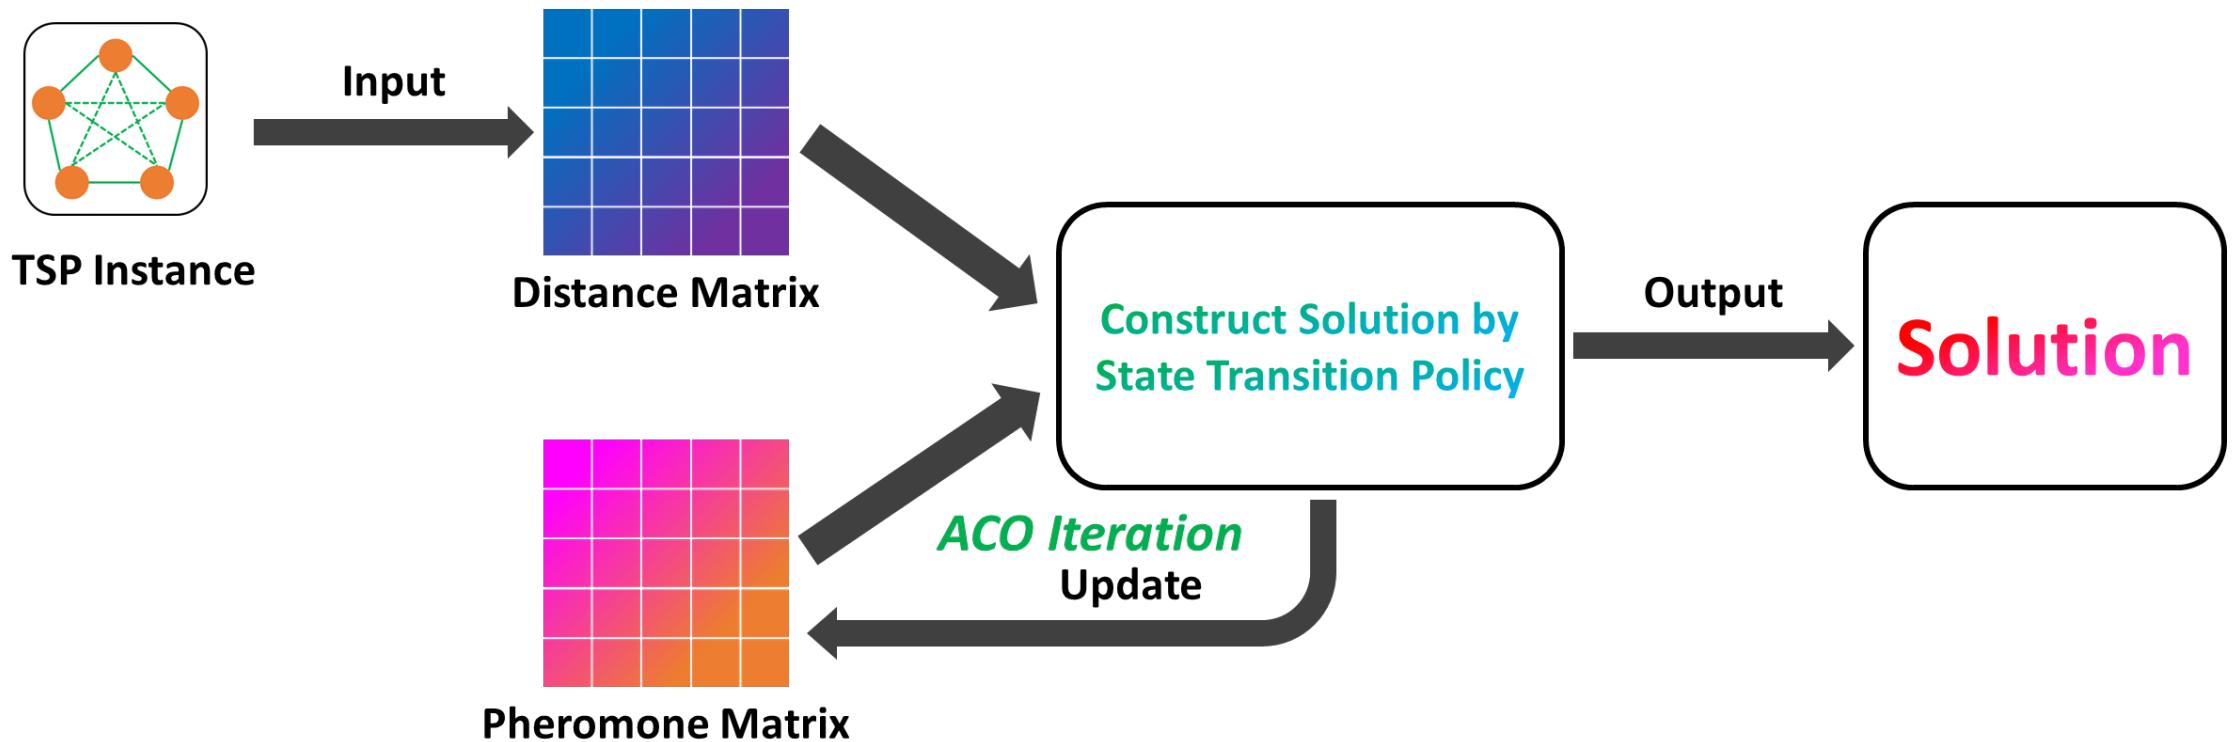


1. Introduction
2. Motivations
3. Research Objectives
4. Preliminary Works
5. Proposed Contributions and Research Plan



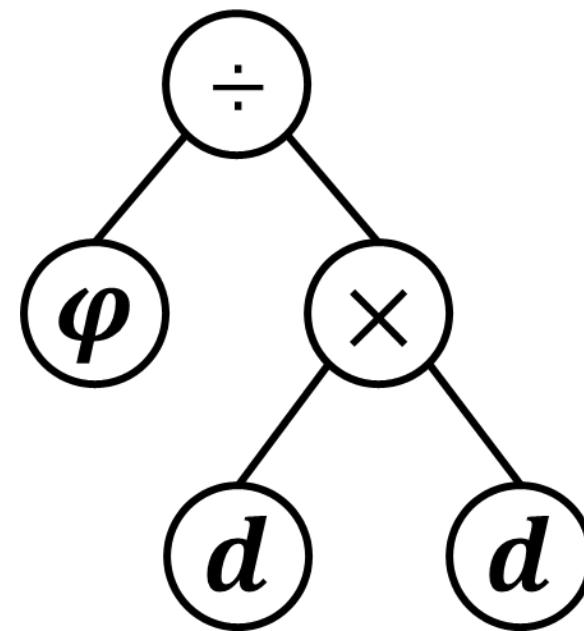
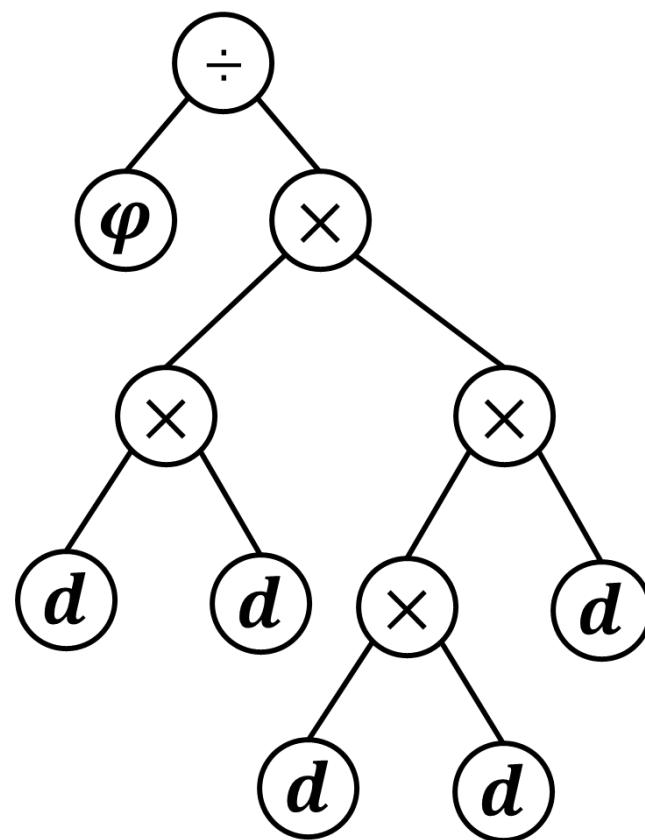
Automated Design of State Transition Rules in Ant Colony Optimization by Genetic Programming: A Comprehensive Investigation

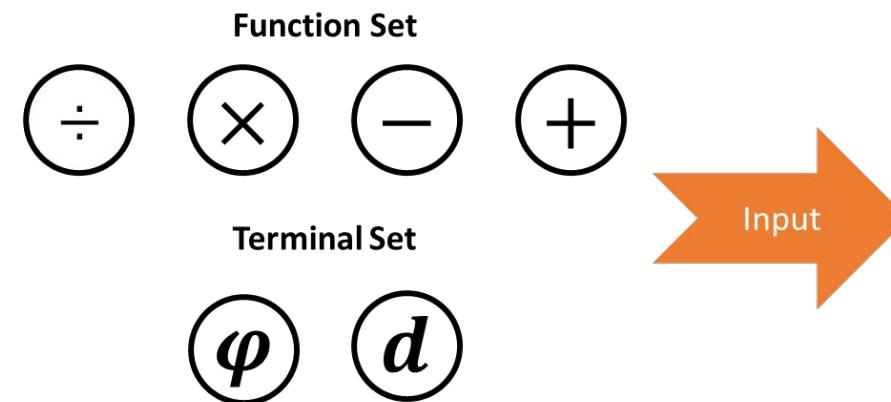
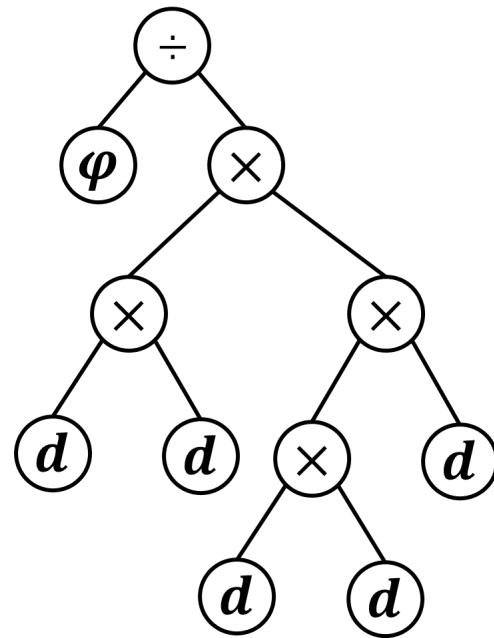
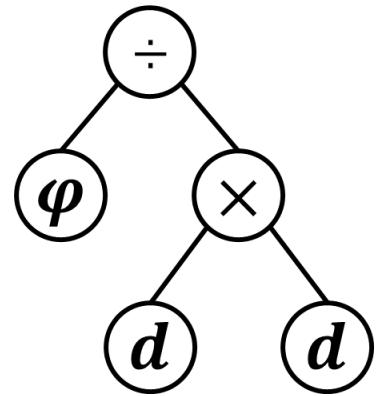
Bocheng Lin, Yi Mei*, Mengjie Zhang



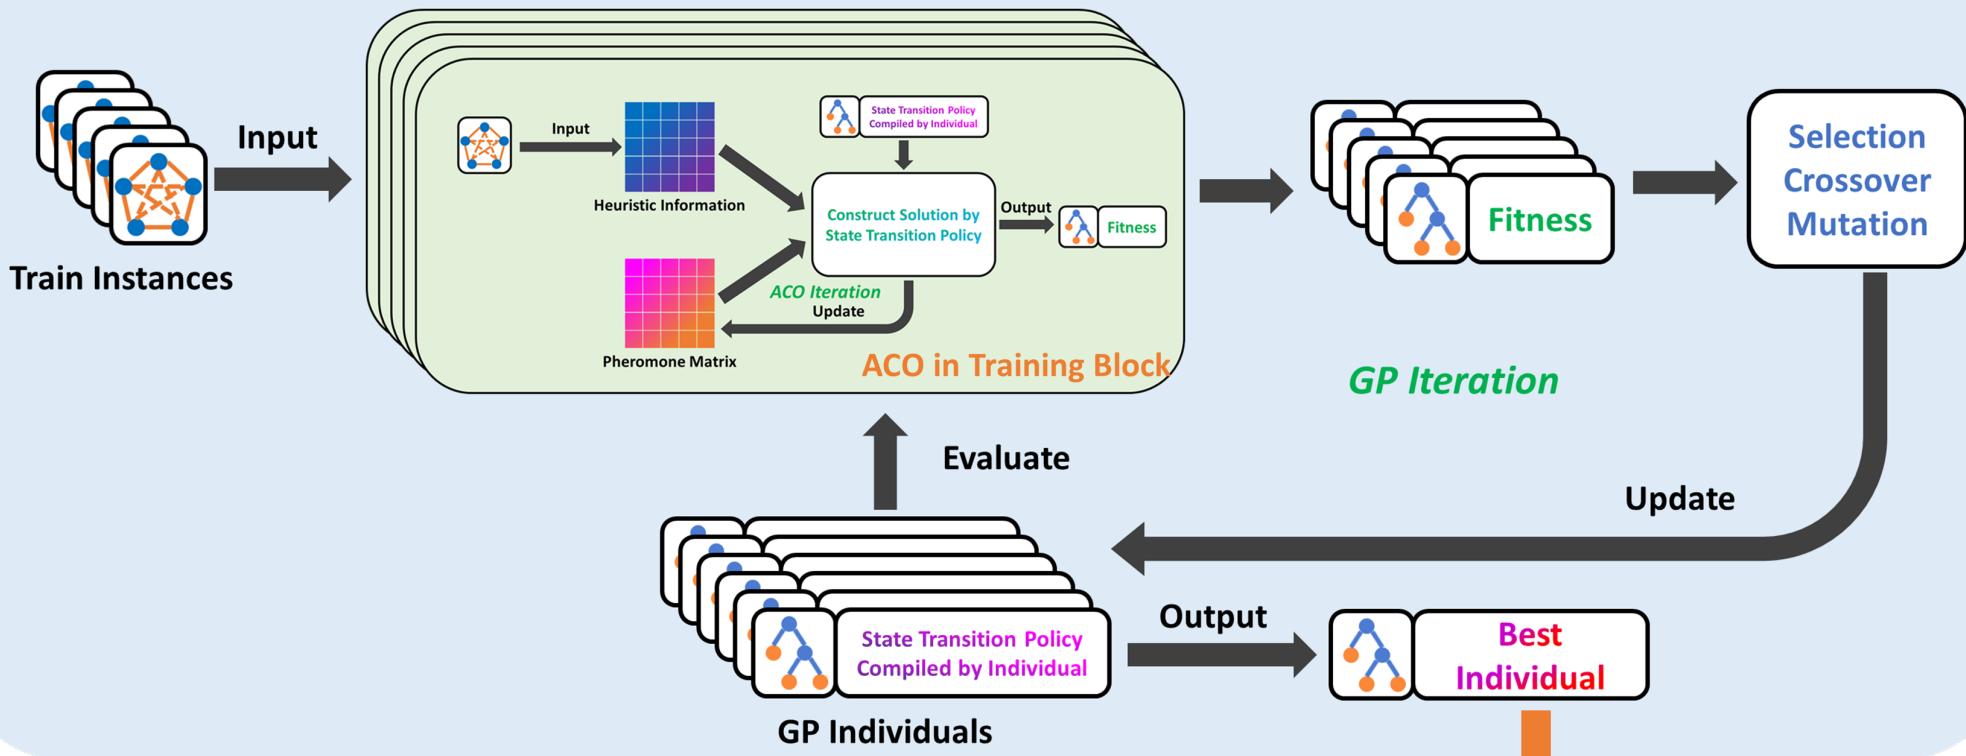


$$p_{ij} = \frac{\varphi_{ij}^{\alpha} / d_{ij}^{\beta}}{\sum_{l \in I_s} \varphi_{il}^{\alpha} / d_{il}^{\beta}}$$

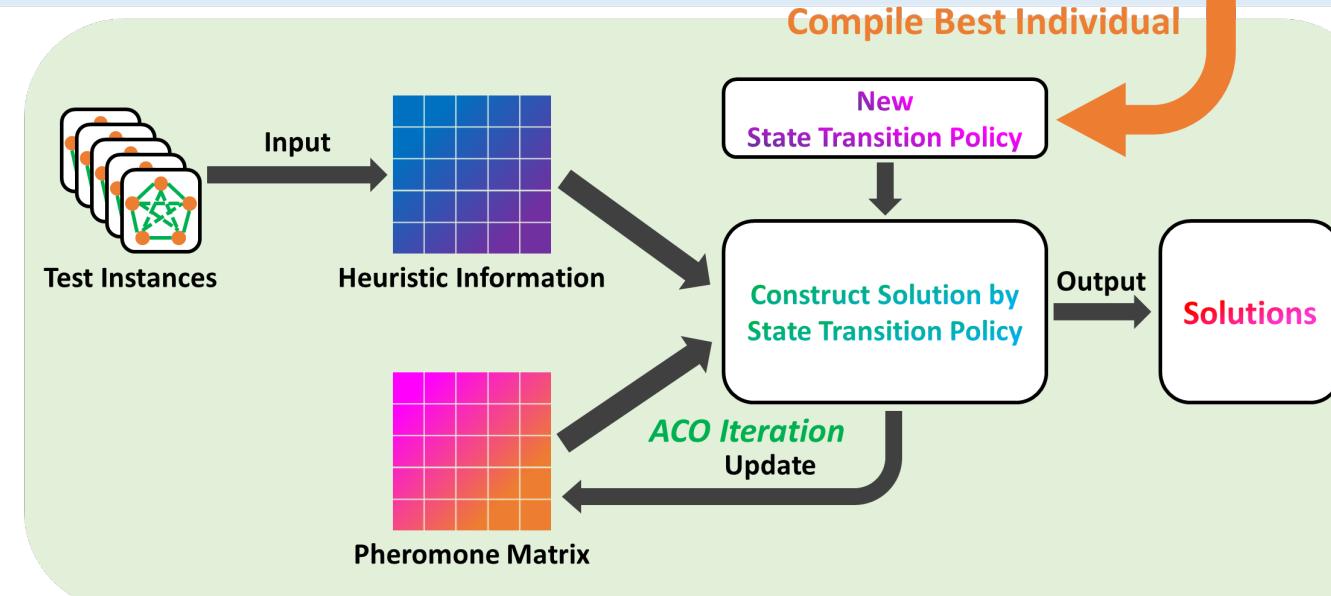




Training Process



ACO Testing Block





Runka's Work in 2009

- Only trained on 1 instance and test on 2 instance
- Only Ant System
- No local search
- No global information



Questions:

1. Does GP-ACO exhibit **generality** when the distributions of training datasets and testing datasets are the same?
2. How do **ACO variants** influence the learning capabilities of GP-ACO?
3. Does the incorporation of **local search** weaken GP-ACO's learning effectiveness?
4. Would incorporating **additional global information** enhance GP-ACO's learning capability?
5. Does GP-ACO remain effective when the **distributions** of training datasets and testing datasets are different?
6. How **interpretable** is GP-ACO?



Experiments to answer Questions:

Experiment 1: To answer Question 1 and Question 2

1. Does GP-ACO exhibit **generality** when the distributions of training datasets and testing datasets are the same?
2. How do **ACO variants** influence the learning capabilities of GP-ACO?

We apply the GP-ACO framework to **AS** and **ACS** and **MMAS**, without considering local search.

Train and test on **three scales of TSP problems (TSP20, TSP20-100, TSP100)**

Experiments to answer Questions:

Experiment 2: To answer Question 3

3. Does the incorporation of **local search** weaken GP-ACO's learning effectiveness?

We introduce **local search** to all the ACO variants in Experiment 1 while keeping the training and testing conditions unchanged.

Experiments to answer Questions:

Experiment 3: To answer Question 4

4. Would incorporating **additional global information** enhance GP-ACO's learning capability?

Building on Experiments 1 and 2, add four terminals containing global information to GP-ACO, which we refer to as xGP-ACO, while maintaining the same training and testing conditions.

Experiments to answer Questions:

Experiment 4: To answer Question 5

5. Does GP-ACO remain effective when the **distributions** of training datasets and testing datasets are different?

We directly use the state transition rules trained by xGP-ACO on the TSP100 dataset from Experiment 3, **test on the publicly available TSPLIB dataset.**



Datasets & Algorithm Settings

- Datasets

Training: TSP20, TSP20-100, TSP100 *50 instances*

Testing: TSP20, TSP20-100, TSP100 *10 instances*

TSPLIB *15 instances*

- Algorithm Settings

GP

50 generations, 100 individuals,
0.8 crossover rate, 0.15 mutation rate,
0.05 reproduction rate,
10 elite individuals, 4 tournament size
5 maximum tree depth
30 independent training

ACO

20 (only TSP20) or 50 ants,
100 generations,
 $\alpha = 1, \beta = 2$ or 5 (only AS)
 $\rho = 0.9$ or 0.98 (only MMAS)
 $q_0 = 0.9$ (only ACS)
2-opt steps = cities/4



Runka's Work in 2009

- Only trained on 1 instance and test on 2 instance
- Only Ant System
- No local search
- No global information

Before experiment, thinking:

Why Runka did not do these?

Why only Runka did this work without anyone following this?

Runka's Work in 2009

- Only trained on 1 instance and test on 2 instance
- Only Ant System
- No local search
- No global information

Execution Time
More than 50 hours

Table 1: GP parameters settings

Parameter	Value
Max Tree Depth	15
Generations	50
Population Size	200
Elitism	1 individual
Initialization	Ramped Half-and-half [2,6]
Selection	Tournament (size 4)
Crossover Type	Subtree Crossover
Crossover Percent	80%
Mutation Type	Subtree Mutation
Mutation Percent	20%
Reproduction	0%
Terminal Node Selection	10%

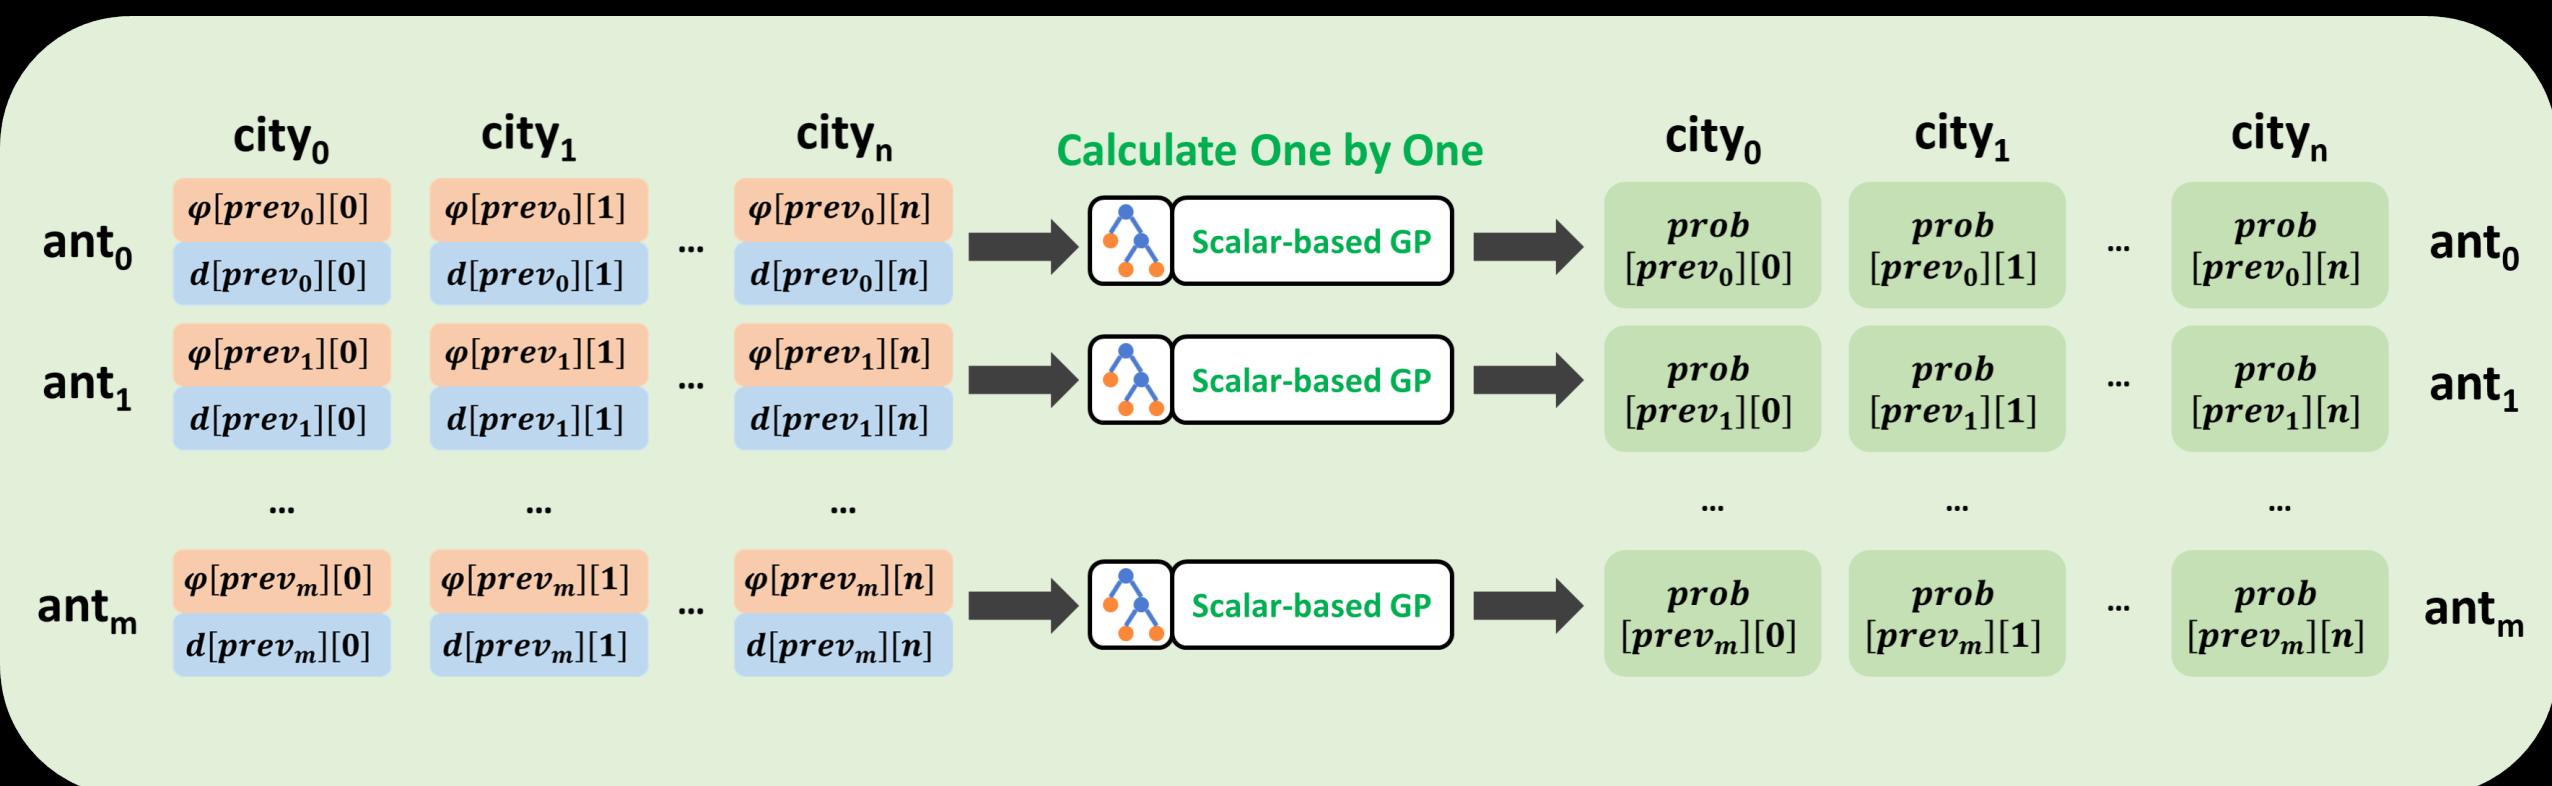
Table 2: AS parameters settings

Parameter	Value
Duration	100
Evaporation Rate (ρ)	0.1
Number of ants	n (problem size)
α (for default formula)	2
β (for default formula)	1
Tournament Size (when needed)	7

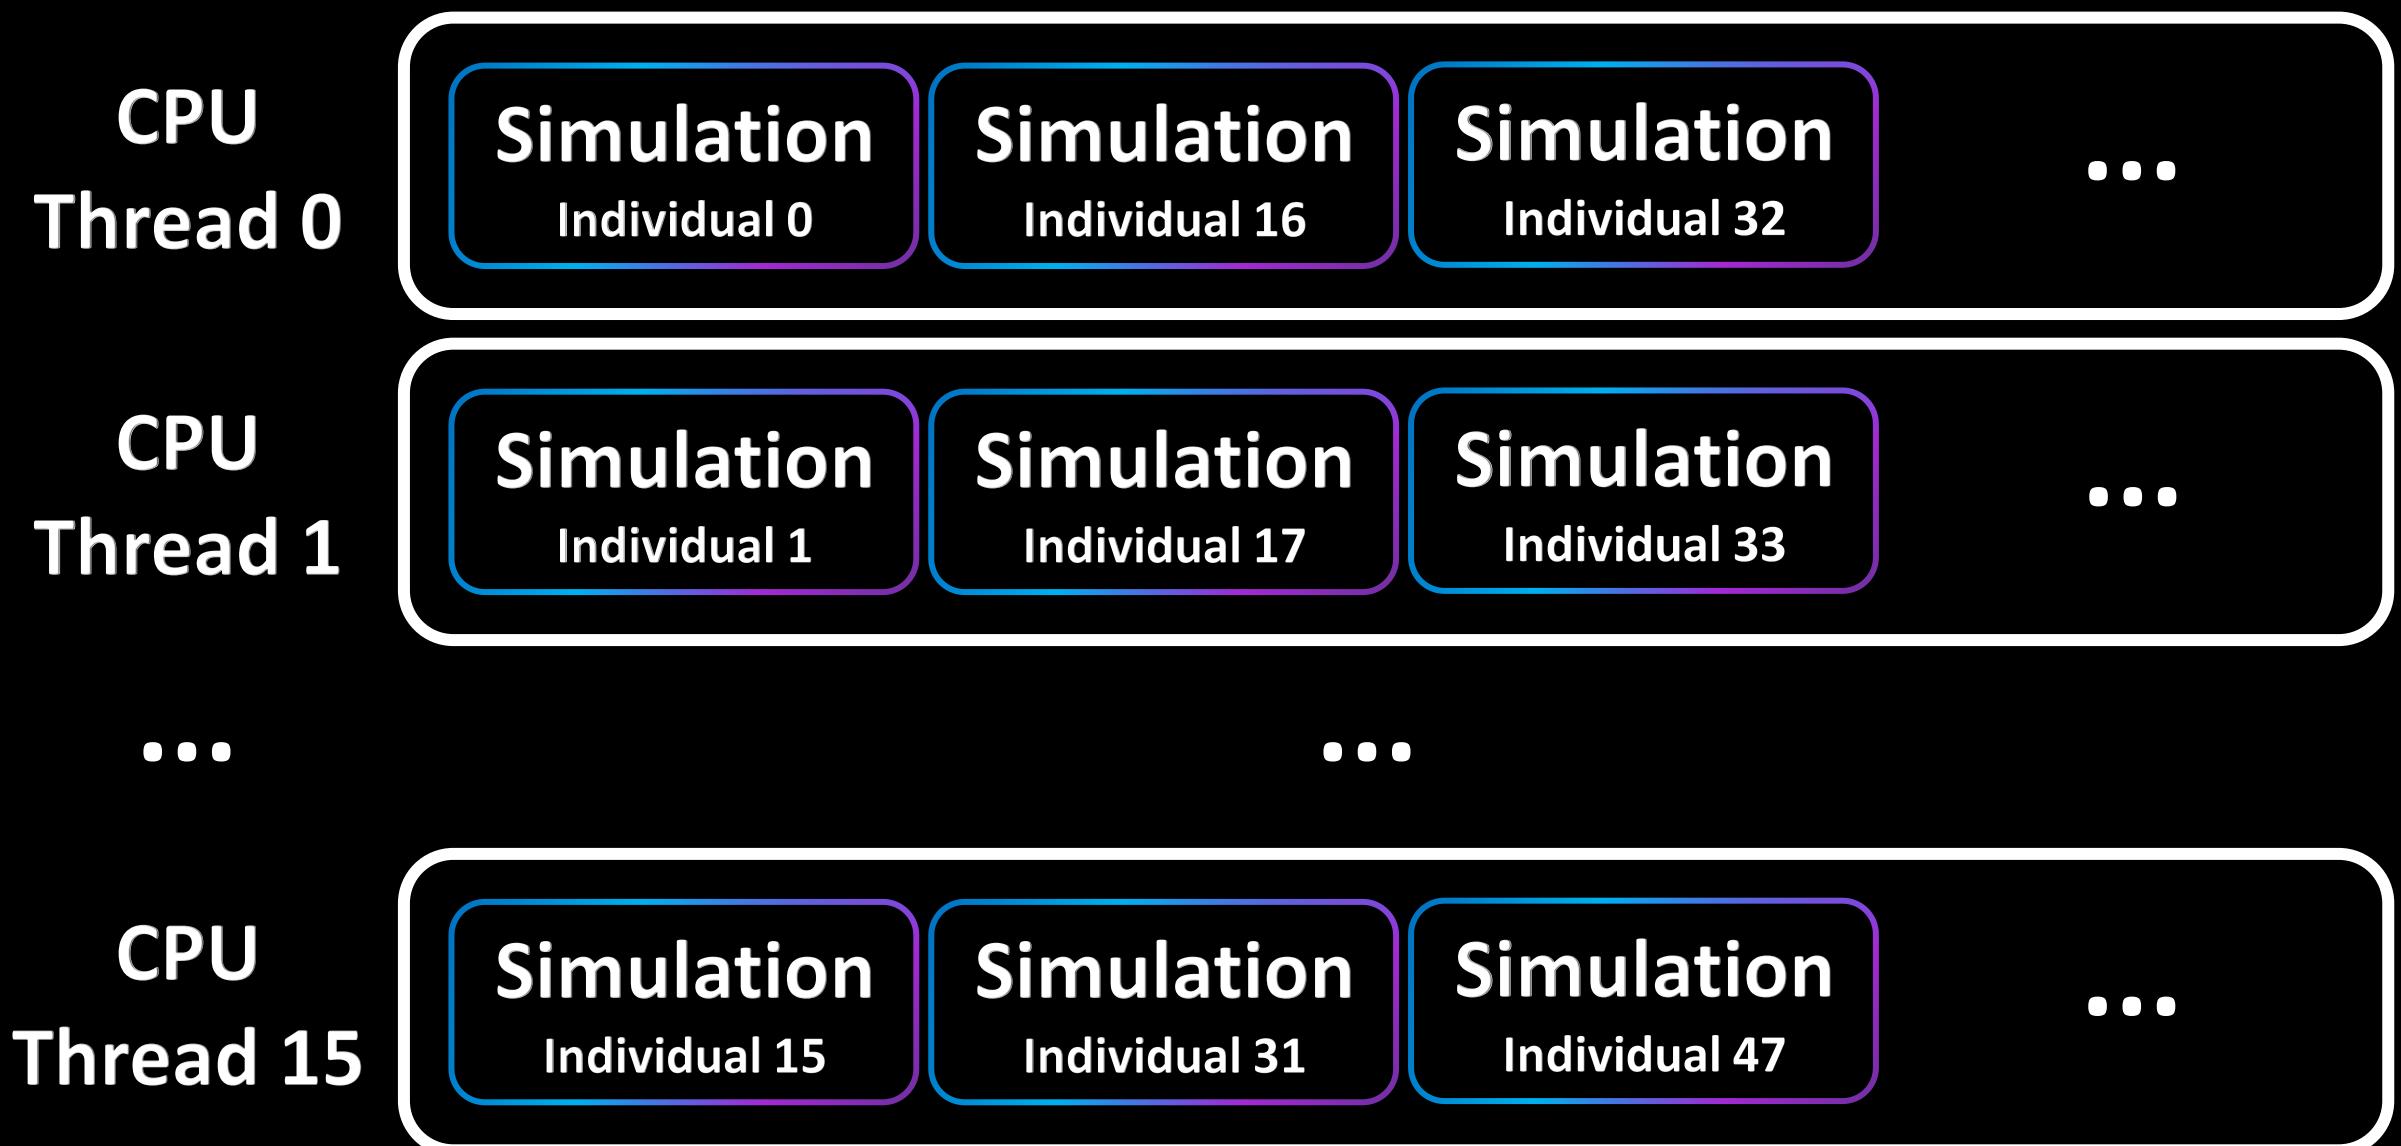


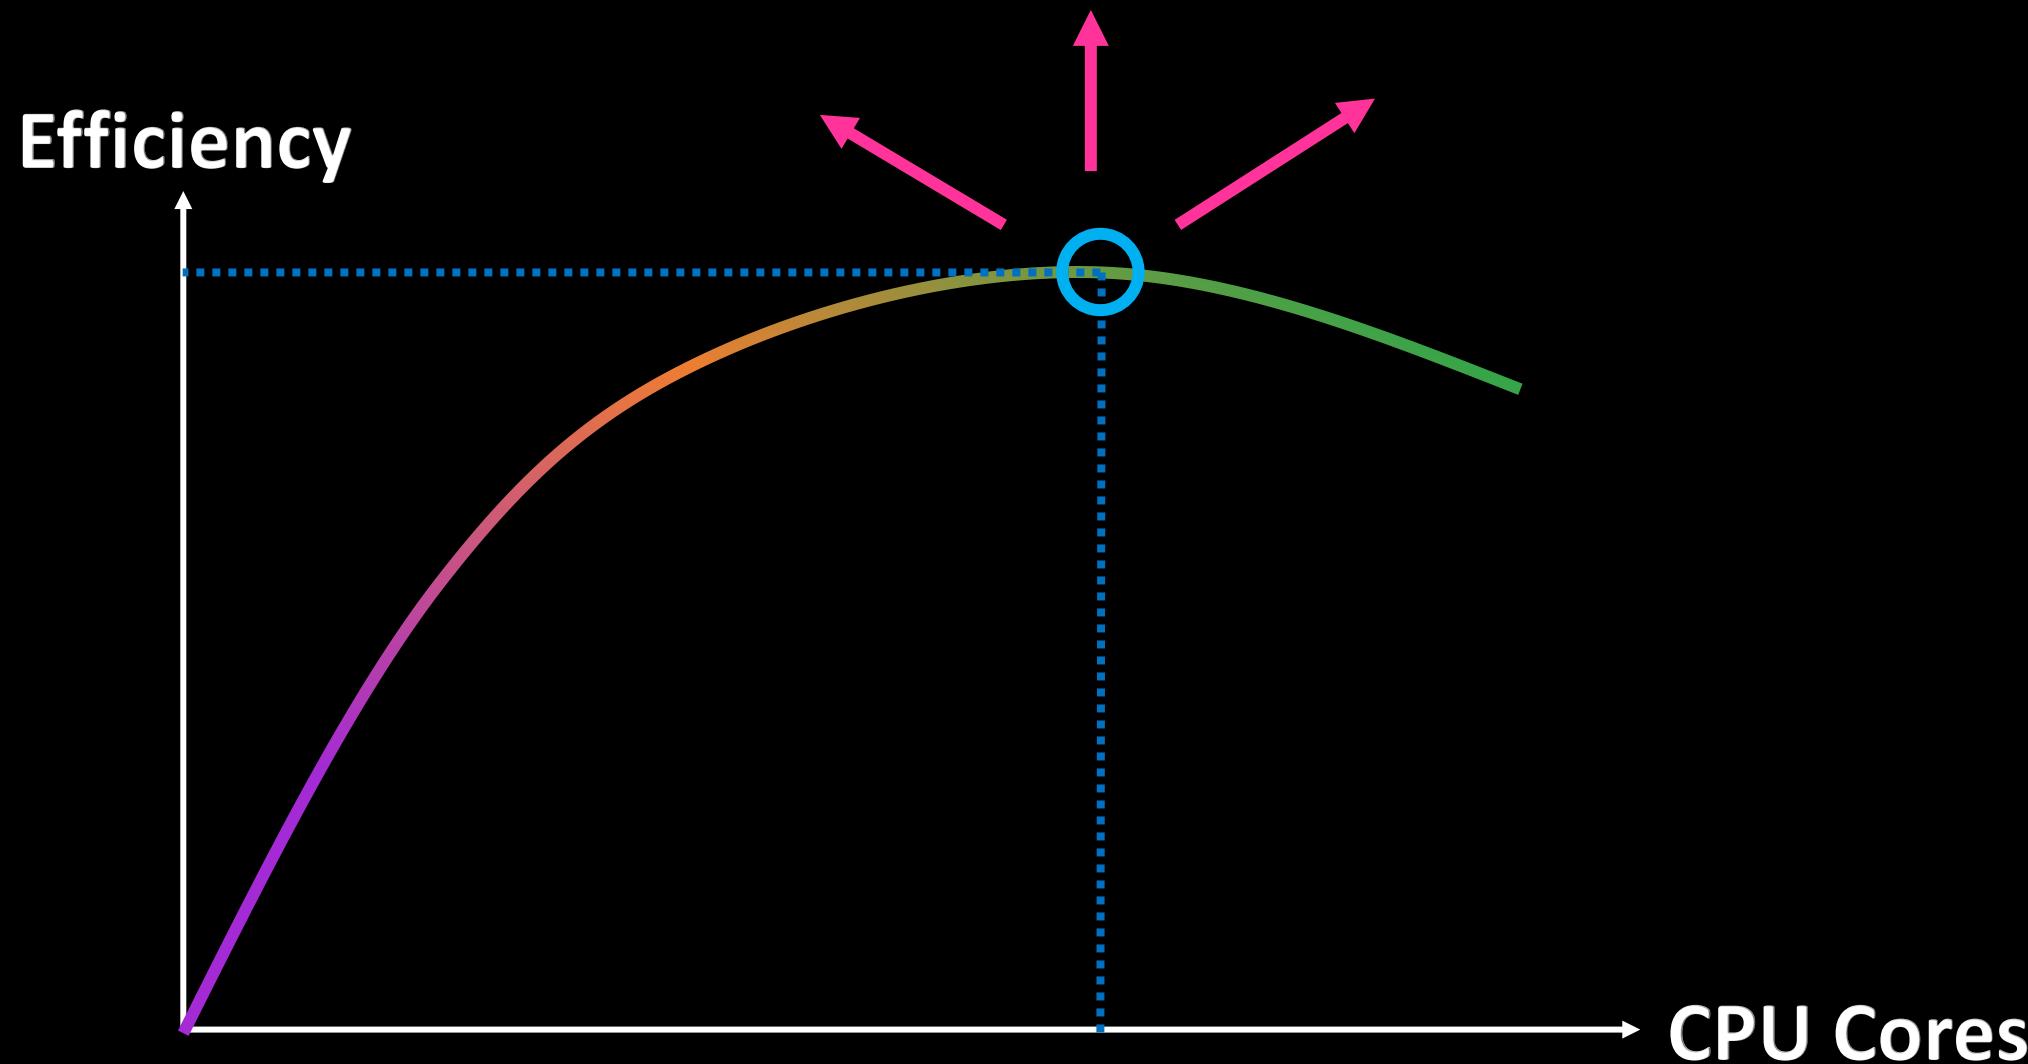
```

for ant in all_ants:
    while unvisited_cities is not empty:
        for cities in unvisted_cities:
    
```



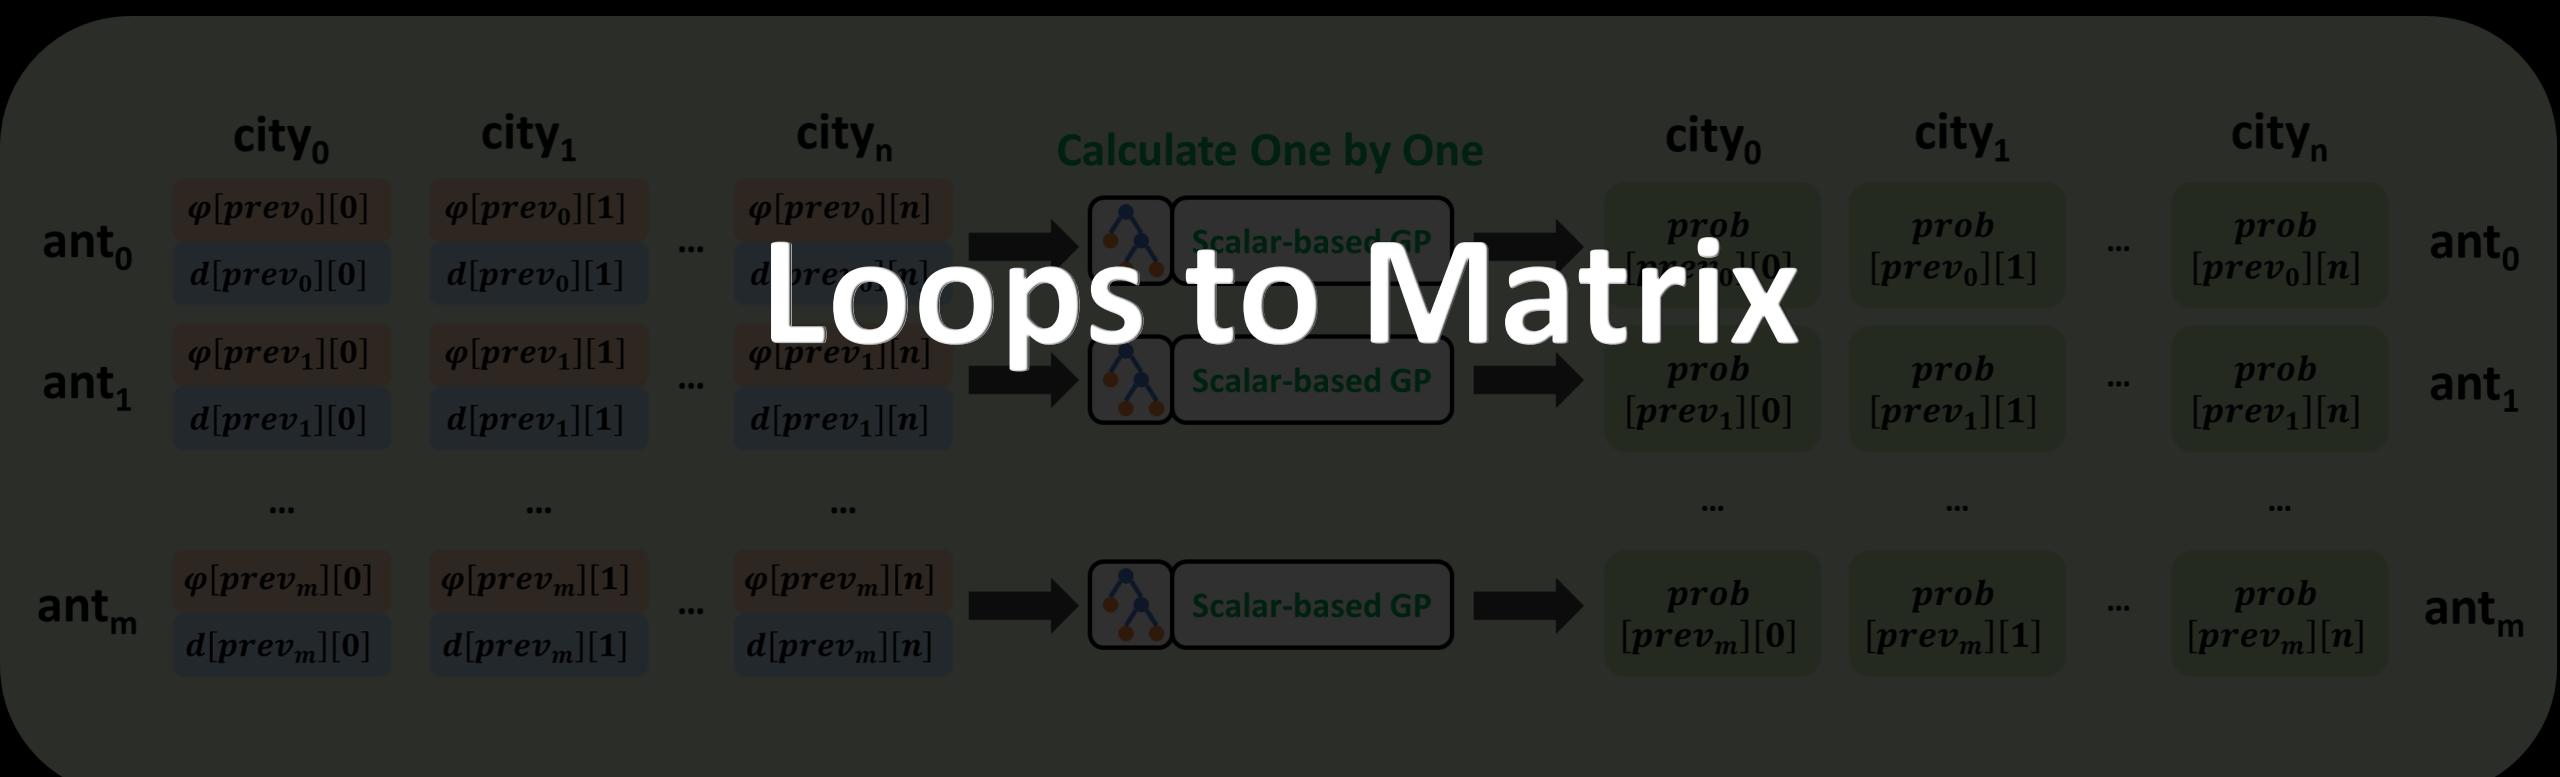
Parallel Simulation/Evaluation

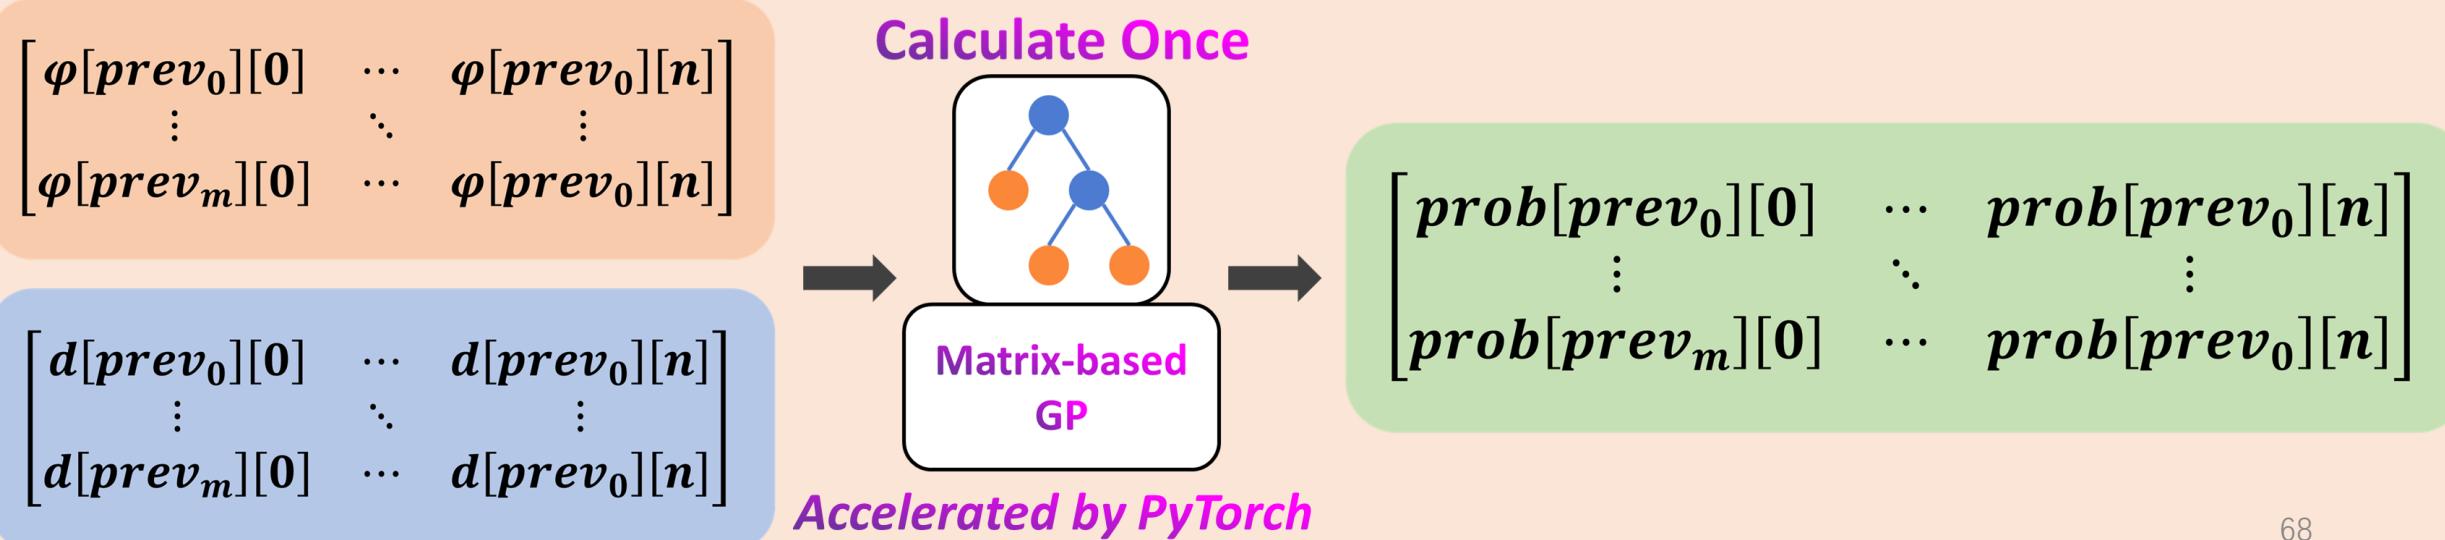
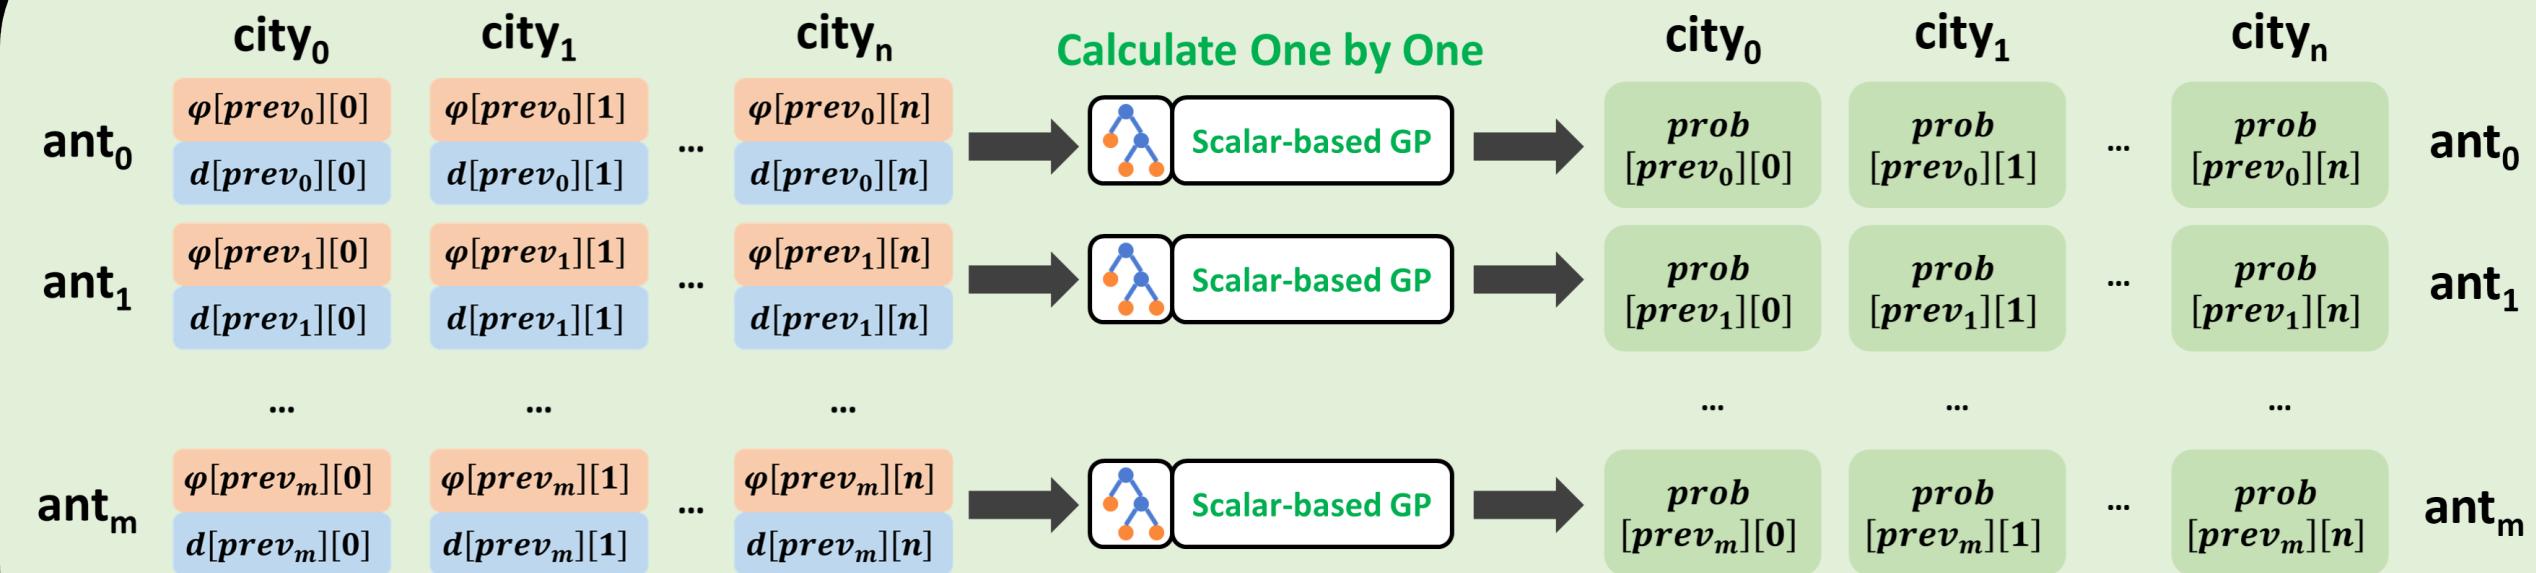




```

for ant in all_ants:
    while unvisited_cities is not empty:
        for cities in unvisted_cities:
    
```





Parallel inside Simulation/Evaluation

CPU
Main
Thread 0

CPU Sub-Thread 0-3
for PyTorch
to accelerate
matrix calculation

Simulation
Individual 0

Simulation
Individual 4

Simulation
Individual 8

...

CPU
Main
Thread 1

CPU Sub-Thread 0-3
for PyTorch
to accelerate
matrix calculation

Simulation
Individual 1

Simulation
Individual 5

Simulation
Individual 9

...

CPU
Main
Thread 2

CPU Sub-Thread 0-3
for PyTorch
to accelerate
matrix calculation

Simulation
Individual 2

Simulation
Individual 6

Simulation
Individual 10

...

CPU
Main
Thread 3

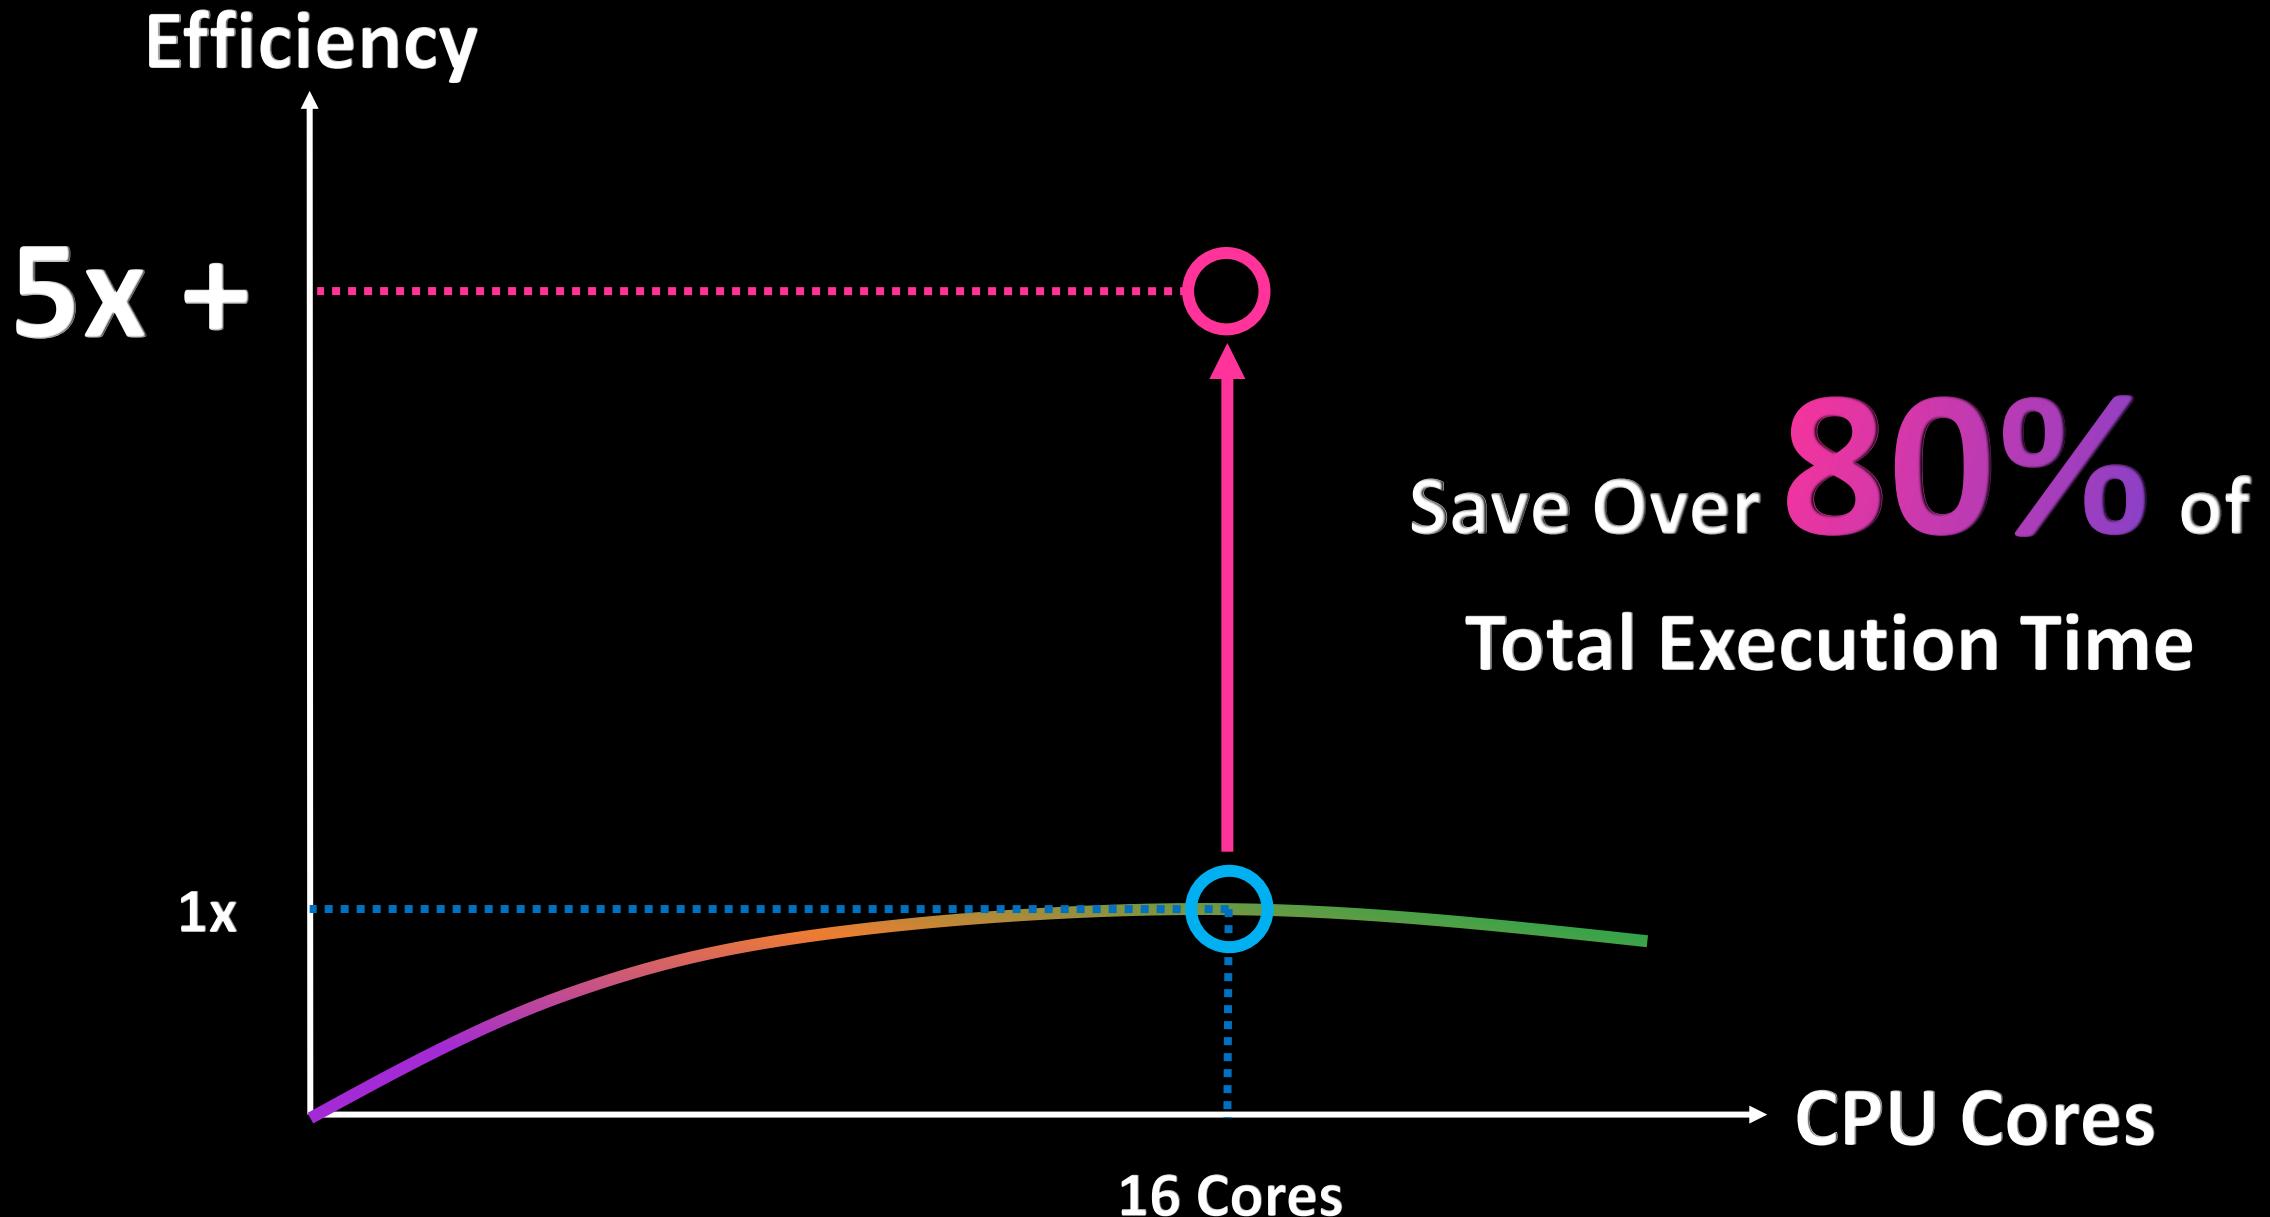
CPU Sub-Thread 0-3
for PyTorch
to accelerate
matrix calculation

Simulation
Individual 3

Simulation
Individual 7

Simulation
Individual 11

...





Experiment 1: To answer Question 1 and Question 2

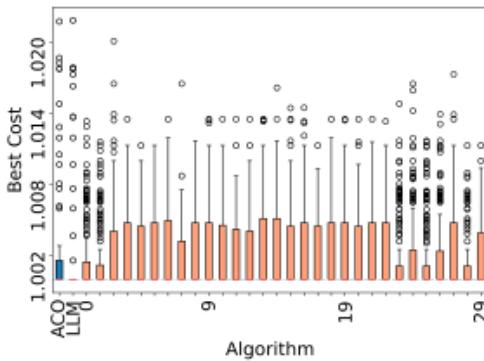
1. Does GP-ACO exhibit **generality** when the distributions of training datasets and testing datasets are the same?
2. How do **ACO variants** influence the learning capabilities of GP-ACO?

We apply the GP-ACO framework to **AS** and **ACS** and **MMAS**, without considering local search.

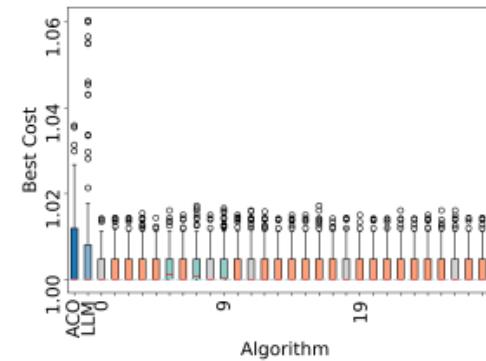
Train and test on **three scales of TSP problems (TSP20, TSP20-100, TSP100)**

Table 1: The normalized costs of GP-ACO, ACO baselines, and LLM-ACO on the test datasets.

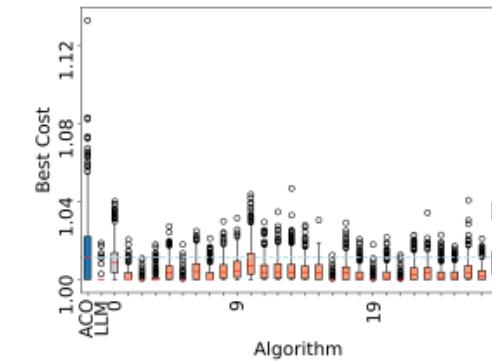
TSP	AS			ACS			MMAS		
	baseline	LLM	GP	baseline	LLM	GP	baseline	LLM	GP
20	mean	1.0028	1.0020	1.0019	1.0063	1.0074	1.0025	1.0164	1.0031
	std	0	0	0.0001	0	0	0.0002	0	0.0002
	min	1	1	1	1	1	1	1	1
	max	1.0217	1.0218	1.0201	1.0359	1.0604	1.0172	1.1327	1.0186
20-100	mean	1.0364	1.0271	1.0265	1.0499	1.0418	1.0306	1.0632	1.0076
	std	0	0	0.0004	0	0	0.0004	0	0.0001
	min	1	1	1	1.0006	1.0016	1	1	1
	max	1.1042	1.0747	1.3813	1.1445	1.0831	1.1660	1.2101	1.0440
100	mean	1.0784	1.0533	1.0105	1.0867	1.0721	1.0446	1.1248	1.0144
	std	0	0	0.0001	0	0	0.0005	0	0.0001
	min	1.0359	1.0127	1.0005	1.0310	1.0148	1.0162	1.0481	1.0071
	max	1.1346	1.0962	1.0274	1.1451	1.1154	1.0799	1.2604	1.0680



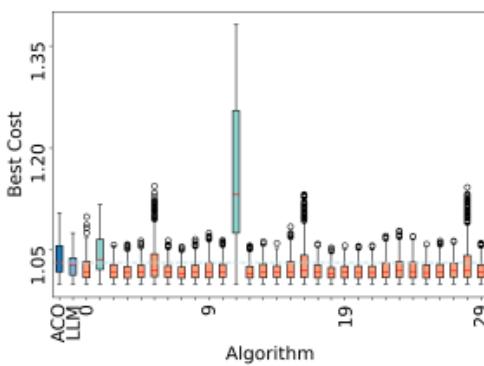
(a) TSP20 - GP-AS



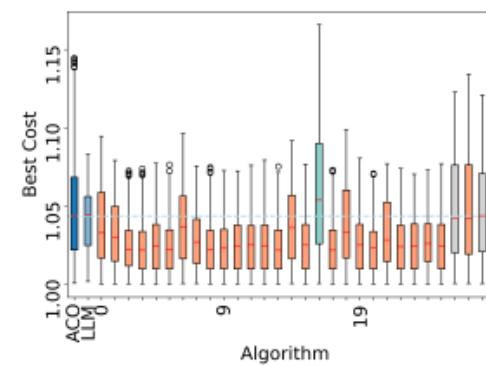
(b) TSP20 - GP-ACS



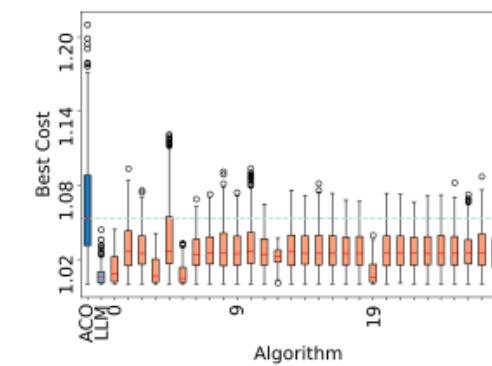
(c) TSP20 - GP-MMAS



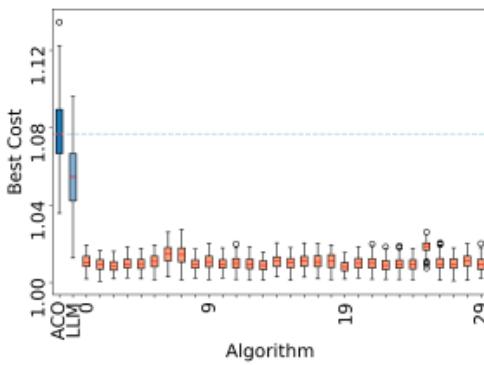
(d) TSP20-100 - GP-AS



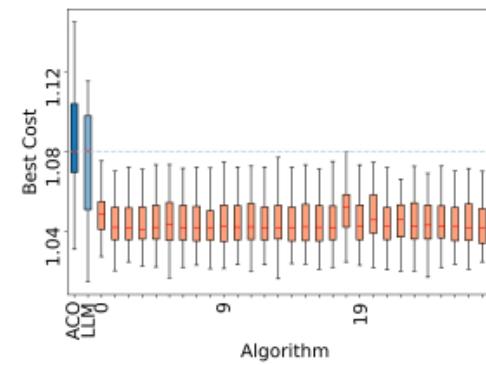
(e) TSP20-100 - GP-ACS



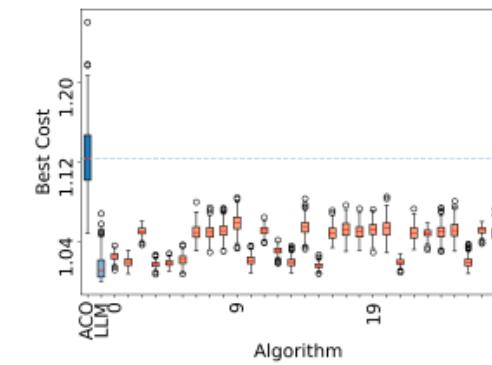
(f) TSP20-100 - GP-MMAS



(g) TSP100 - GP-AS



(h) TSP100 - GP-ACS



(i) TSP100 - GP-MMAS



Experiment 1: To answer Question 1 and Question 2

1. Does GP-ACO exhibit **generality** when the distributions of training datasets and testing datasets are the same?
2. How do **ACO variants** influence the learning capabilities of GP-ACO?

Answer 1: The GP-ACO demonstrates generality.

Across three scales of the TSP problem, the trained GP-ACO consistently outperforms the ACO baseline, and this advantage becomes more pronounced as the problem size increases.

Answer 2: Different ACO frameworks do not affect the learning capability of the GP-ACO.

In three variants of the ACO, the GP-ACO significantly surpasses the ACO baseline in performance. Additionally, we discovered that the state transition rules designed in GP-AS and GP-ACS notably outperform those in LLM-ACO, whereas the reverse is true for GP-MMAS.

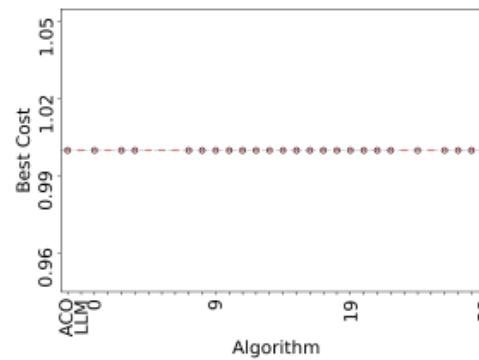
Experiment 2: To answer Question 3

3. Does the incorporation of **local search** weaken GP-ACO's learning effectiveness?

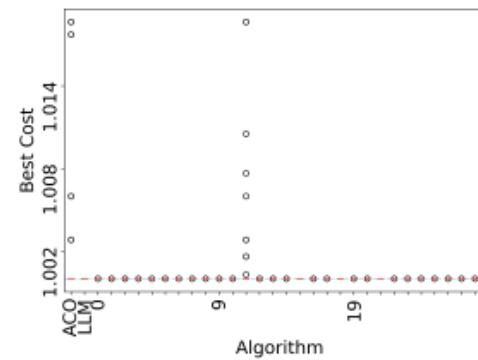
We introduce **local search** to all the ACO variants in Experiment 1 while keeping the training and testing conditions unchanged.

Table 2: The normalized costs of GP-ACO+2-opt, ACO+2-opt baselines, and LLM-ACO+2-opt on the test datasets.

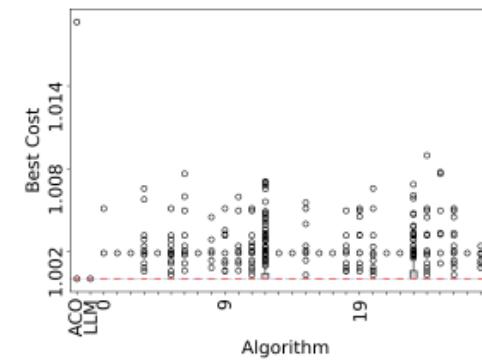
TSP	AS+2-opt			ACS+2-opt			MMAS+2-opt			
	baseline	LLM	GP	baseline	LLM	GP	baseline	LLM	GP	
20	mean	1	1	1	1.0003	1	1	1.0002	1	1.0001
		=	=		+	-		+	-	
	std	0	0	0	0	0	0	0	0	
	min	1	1	1	1	1	1	1	1	
20-100	max	1	1	1	1.0186	1.0000	1.0186	1.0186	1	1.0090
	mean	1.0018	1.0037	1.0008	1.0021	1.0046	1.0008	1.0007	1.0008	1.0008
		+	+		+	+		-	+	
	std	0	0	0	0	0	0.0001	0	0	0.0001
100	min	1	1	1	1	1	1	1	1	1
	max	1.0101	1.0162	1.0452	1.0097	1.0173	1.0123	1.0102	1.0084	1.0495
	mean	1.0045	1.0091	1.0021	1.0045	1.0096	1.0017	1.0017	1.0011	1.0013
		+	+		+	+	=	-	-	
	std	0	0	0.0001	0	0	0.0001	0	0	0.0001
	min	1	1	1	1	1	1	1	1	1
	max	1.0184	1.0269	1.0136	1.0211	1.0281	1.0117	1.0112	1.0185	1.0276



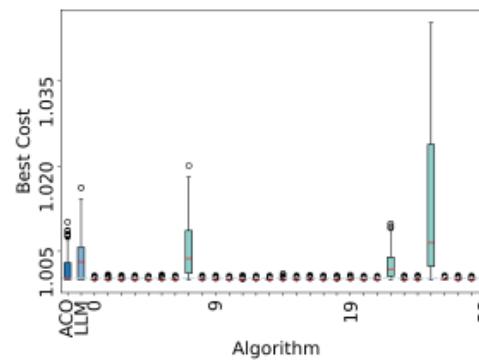
(a) TSP20 - GP-AS+2-opt



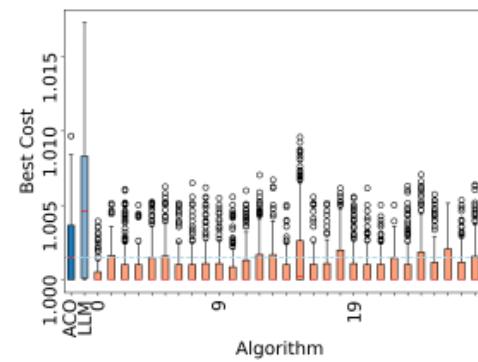
(b) TSP20 - GP-ACS+2-opt



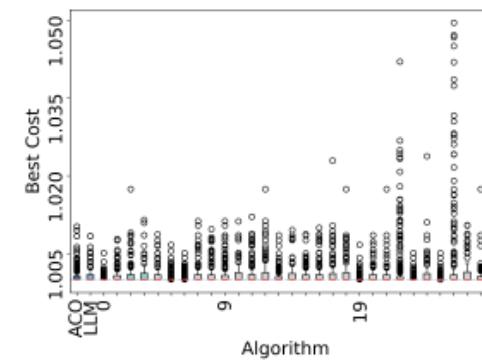
(c) TSP20 - GP-MMAS+2-opt



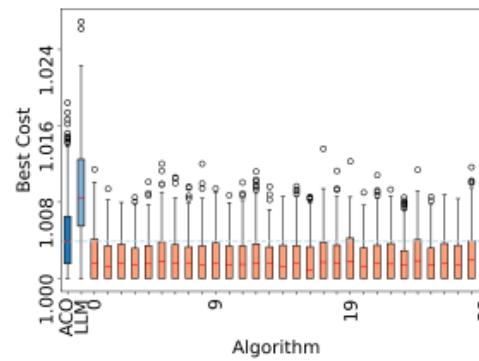
(d) TSP20-100 - GP-AS+2-opt



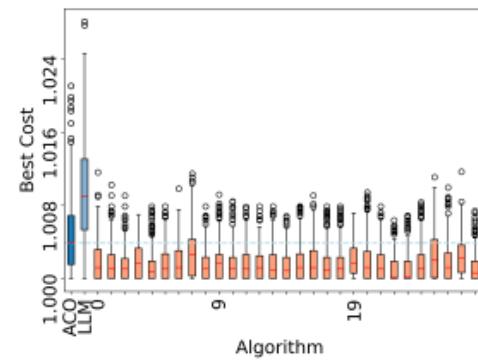
(e) TSP20-100 - GP-ACS+2-opt



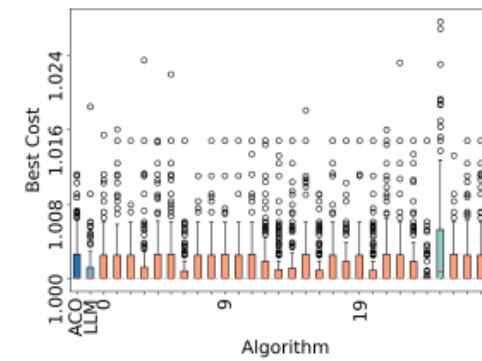
(f) TSP20-100 - GP-MMAS+2-opt



(g) TSP100 - GP-AS+2-opt



(h) TSP100 - GP-ACS+2-opt



(i) TSP100 - GP-MMAS+2-opt



Experiment 2: To answer Question 3

3. Does the incorporation of **local search** weaken GP-ACO's learning effectiveness?

Answer 3: The introduction of local search does not impact the learning effectiveness of GP-ACO in the GP-AS+2-opt and GP-ACS+2-opt scenarios, but it diminishes the learning capabilities of GP-ACO in the GP-MMAS+2-opt scenario.

Enhancing Answer 1: GP-ACO+2-opt continues to exhibit generality.

However, in the case of GP-MMAS+2-opt, its performance in generality did not meet expectations.

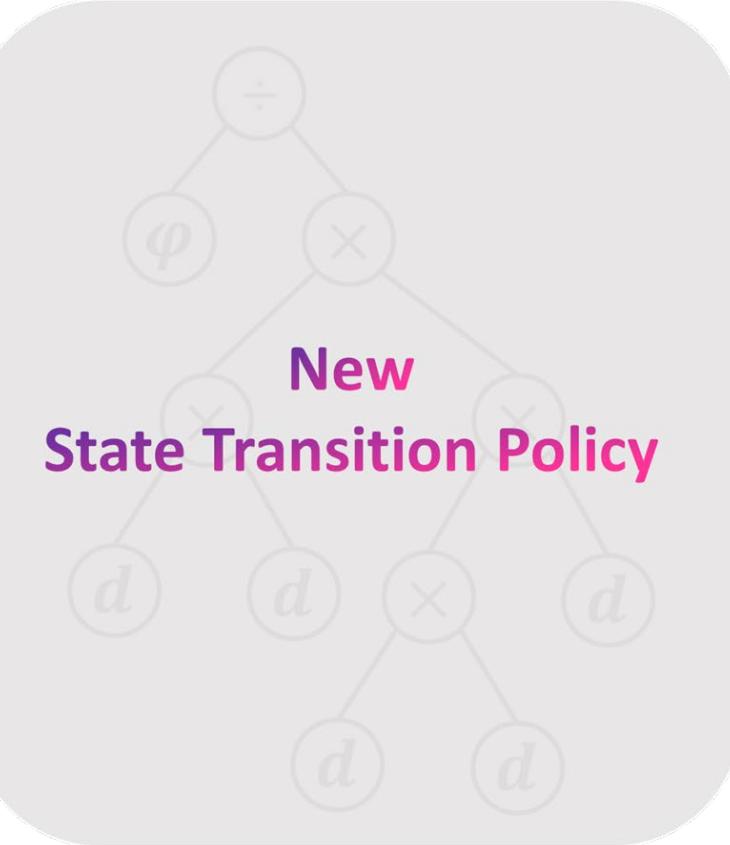
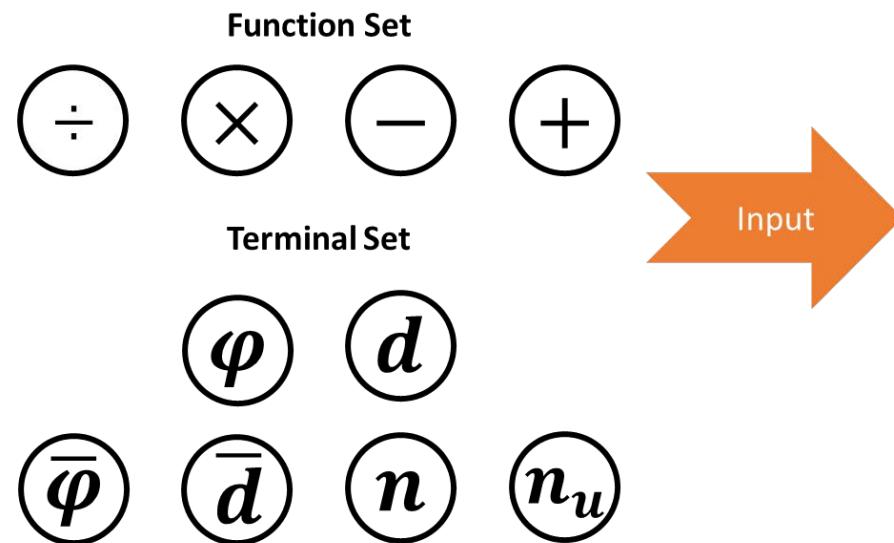
Enhancing Answer 2: ACO variants that rely more on state transition rules are almost unaffected by the introduction of 2-opt local search regarding the learning capabilities of GP-ACO.

Conversely, ACO variants that depend more heavily on pheromone updates show a reduction in GP-ACO's learning abilities when local search is implemented.

Experiment 3: To answer Question 4

4. Would incorporating **additional global information** enhance GP-ACO's learning capability?

Building on Experiments 1 and 2, add four terminals containing global information to GP-ACO, which we refer to as xGP-ACO, while maintaining the same training and testing conditions.



average pheromone to the candidate node from the current node

average distance to the candidate node from the current node

total number of nodes

number of candidate nodes



TSP	AS				ACS				MMAS				
	baseline	LLM	GP	xGP	baseline	LLM	GP	xGP	baseline	LLM	GP	xGP	
20	mean	1.0028	+ 1.0020	+ 1.0019	+ 1.0014	1.0063	+ 1.0074	+ 1.0025	+ 1.0027	1.0164	+ 1.0031	+ 1.0039	+ 1.0015
	std	0	0	0.0001	0.0001	0	0	0.0002	0.0002	0	0	0.0002	0.0002
	min	1	1	1	1	1	1	1	1	1	1	1	1
	max	1.0217	1.0218	1.0201	1.0159	1.0359	1.0604	1.0172	1.0172	1.1327	1.0186	1.0467	1.0403
20-100	mean	1.0364	+ 1.0271	+ 1.0265	+ 1.0235	1.0499	+ 1.0418	+ 1.0306	+ 1.0293	1.0632	+ 1.0076	+ 1.0254	+ 1.0029
	std	0	0	0.0004	0.0004	0	0	0.0004	0.0004	0	0	0.0001	0.0001
	min	1	1	1	1	1.0006	1.0016	1	1	1	1	1	0.9940
	max	1.1042	1.0747	1.3813	1.1882	1.1445	1.0831	1.1660	1.1177	1.2101	1.0440	1.1212	1.0579
100	mean	1.0784	+ 1.0533	+ 1.0105	+ 1.0079	1.0867	+ 1.0721	+ 1.0446	+ 1.0440	1.1248	+ 1.0144	+ 1.0400	+ 1.0039
	std	0	0	0.0001	0.0001	0	0	0.0005	0.0005	0	0	0.0001	0.0001
	min	1.0359	1.0127	1.0005	1.0008	1.0310	1.0148	1.0162	1.0156	1.0481	1	1.0071	0.9881
	max	1.1346	1.0962	1.0274	1.0194	1.1451	1.1154	1.0799	1.0797	1.2604	1.0680	1.0852	1.0400



TSP	AS+2-opt					ACS+2-opt					MMAS+2-opt					
	baseline	LLM	GP	xGP		baseline	LLM	GP	xGP		baseline	LLM	GP	xGP		
20	mean	1	-	1	=	1	=	1		1.0003	+	1	=	1	1.0002	
	std	0		0		0		0		0		0		0	0	
	min	1		1		1		1		1		1		1	1	
	max	1		1		1		1.0063		1.0186		1.0000		1.0186	1.0186	
20-100	mean	1.0018	+	1.0037	+	1.0008	+	1.0002		1.0021	+	1.0046	+	1.0008	+	1.0008
	std	0		0		0		0		0		0.0001		0.0001		0.0001
	min	1		1		1		1		1		1		1	1	1
	max	1.0101		1.0162		1.0452		1.0053		1.0097		1.0173		1.0123		1.0080
100	mean	1.0045	+	1.0091	+	1.0021	+	1.0017		1.0045	+	1.0096	+	1.0017	-	1.0018
	std	0		0		0.0001		0.0001		0		0		0.0001		0.0001
	min	1		1		1		1		1		1		1	1	1
	max	1.0184		1.0269		1.0136		1.0142		1.0211		1.0281		1.0117		1.0141

Experiment 3: To answer Question 4

4. Would incorporating **additional global information** enhance GP-ACO's learning capability?

Answer 4: The introduction of additional global information in xGP-ACO without local search proves highly effective,
significantly enhancing learning capabilities compared to GP-ACO.

With local search, incorporating more global information significantly improves the learning capabilities of GP-AS+2-opt

but weakens those of xGP-MMAS+2-opt, which has less dependency on state transition rules.

Enhancing Answer 2 & 3: After introducing more global information, and excluding the consideration of local search, **variations in ACO frameworks do not reduce the learning capabilities of xGP-ACO.**

However, when the local search is introduced, **the lesser dependence on state transition rules in MMAS variants diminishes the learning capabilities of xGP-ACO+2-opt.**



Experiment 4: To answer Question 5

5. Does GP-ACO remain effective when the **distributions** of training datasets and testing datasets are different?

We directly use the state transition rules trained by xGP-ACO on the TSP100 dataset from Experiment 3, test on the publicly available TSPLIB dataset.



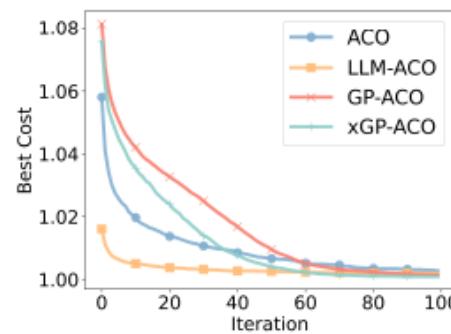
TSP instances	AS				ACS				MMAS			
	baseline	LLM	xGP	xGP-best	baseline	LLM	xGP	xGP-best	baseline	LLM	xGP	xGP-best
bier127	10.75	5.52	6.62	8.83	17.84	6.60	7.83	6.56	6.80	2.06	5.45	0.54
ch130	9.84	5.57	0.88	0.75	20.59	6.19	10.11	9.63	7.98	2.70	0.71	1.53
d493	21.23	11.31	11.20	3.15	32.53	10.90	16.25	13.97	11.22	8.42	4.16	4.98
eil51	7.57	4.01	1.11	0.78	13.20	5.04	10.47	9.29	5.21	1.71	1.11	0.31
fl417	16.90	14.46	4.26	1.81	26.12	15.11	12.68	12.18	9.98	9.29	2.53	4.83
kroA150	13.76	9.26	3.47	2.23	23.55	10.02	13.67	11.62	8.83	4.51	1.92	1.05
kroB100	7.03	5.46	1.35	1.36	16.36	7.41	5.76	6.30	7.43	1.42	1.46	0.26
kroC100	6.92	4.68	1.25	1.08	17.05	5.49	7.75	7.22	7.96	1.14	1.39	0.18
lin318	17.34	10.21	11.25	6.40	26.70	9.65	16.51	15.74	10.73	6.50	3.87	3.93
pr226	15.00	7.28	10.33	9.30	17.58	9.89	12.01	10.27	8.74	2.20	5.62	2.24
pr264	13.62	8.92	8.62	7.52	25.30	9.30	12.75	11.49	9.82	5.99	3.81	3.30
pr299	20.97	10.36	13.36	7.92	33.72	11.21	19.04	17.28	12.23	8.13	4.23	3.69
pr439	23.52	10.94	21.67	16.86	36.07	13.35	19.71	16.19	12.27	7.31	7.53	4.45
rat99	9.00	6.50	1.18	1.01	15.58	8.60	10.54	10.76	8.49	3.61	1.85	1.97
ts225	9.25	3.70	14.82	19.71	23.17	3.74	8.91	5.56	12.08	3.10	8.71	0.70
Avg. opt. gap	13.51	7.88	7.43	5.91	23.02	8.83	12.27	10.93	9.32	4.54	3.62	2.26

TSP instances	AS+2opt				ACS+2opt				MMAS+2opt			
	baseline	LLM	xGP	xGP-best	baseline	LLM	xGP	xGP-best	baseline	LLM	xGP	xGP-best
bier127	0.85	0.74	0.99	0.51	1.26	0.71	0.38	0.26	3.60	0.07	2.05	0.03
ch130	0.94	1.25	0.03	0.25	1.89	1.40	0.94	1.01	4.00	0.22	0.20	0.14
d493	2.30	2.60	1.82	1.42	3.30	2.25	2.07	1.82	7.96	1.12	4.41	0.58
eil51	0.95	0.90	0.39	0.68	1.50	0.94	0.72	0.71	4.38	0.70	0.41	0.69
fl417	1.13	1.67	0.35	0.83	1.74	1.62	1.07	1.05	2.60	0.59	5.49	0.52
kroA150	1.30	1.81	0.52	0.69	2.04	2.11	0.67	0.56	4.60	0.21	1.41	0.04
kroB100	0.52	0.65	0.27	0.26	0.81	0.40	0.18	0.21	3.06	0.09	0.25	0.08
kroC100	0.28	0.67	0.22	0.008	0.65	0.81	0.04	0.04	2.56	0.05	0.26	0.03
lin318	2.38	2.41	1.26	1.72	3.29	2.06	1.65	1.03	5.74	0.58	5.74	0.45
pr226	0.40	0.86	0.55	0.26	0.89	1.03	0.51	0.52	1.14	0.03	8.63	0.007
pr264	1.04	0.73	0.89	0.21	1.43	2.49	0.27	0.16	5.45	0.007	5.49	0.07
pr299	1.62	1.85	1.66	1.22	2.43	1.96	1.18	1.02	7.60	0.33	6.35	0.18
pr439	1.49	1.84	3.50	1.26	2.20	1.63	1.34	1.00	7.17	0.55	9.36	0.22
rat99	1.04	1.44	0.38	0.71	2.09	1.16	0.74	0.75	7.35	0.72	0.46	0.69
ts225	0.54	0.74	3.29	0.48	0.72	1.37	0.32	0.29	3.34	0.33	7.61	0.002
Avg. opt. gap	1.12	1.34	1.08	0.70	1.75	1.46	0.80	0.70	4.70	0.37	3.87	0.25

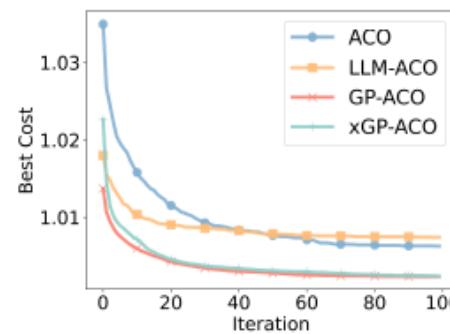
Experiment 4: To answer Question 5

5. Does GP-ACO remain effective when the **distributions** of training datasets and testing datasets are different?

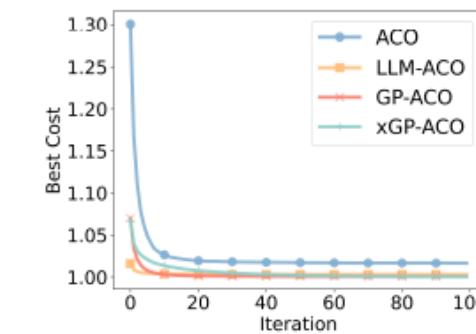
Answer 5: The state transition rules learned by xGP-ACO and xGP-ACO+2-opt demonstrate strong performance across maps of varying scales and distributions. **This indicates that the generality capabilities of xGP-ACO extend beyond training scenarios, adapting effectively to diverse and larger-scale environments.**



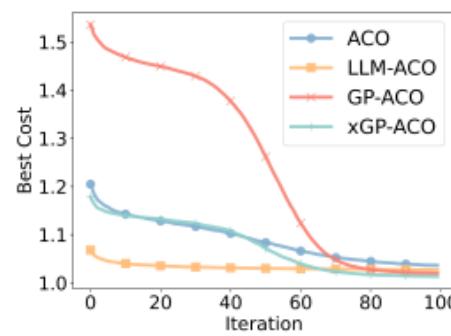
(a) TSP20-AS



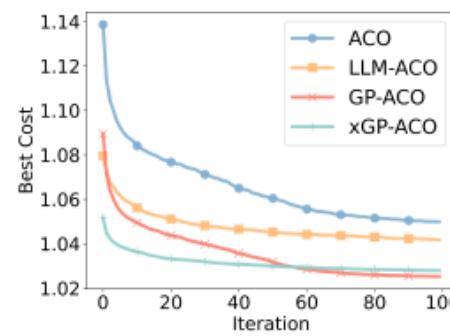
(b) TSP20-ACS



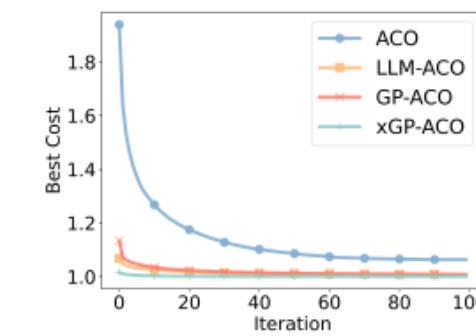
(c) TSP20-MMAS



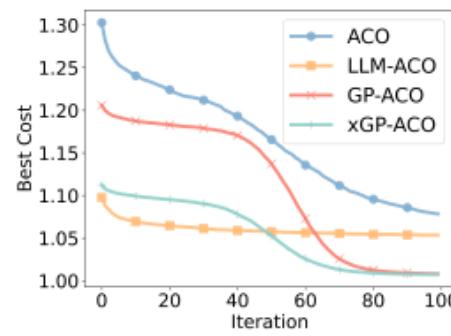
(d) TSP20-100-AS



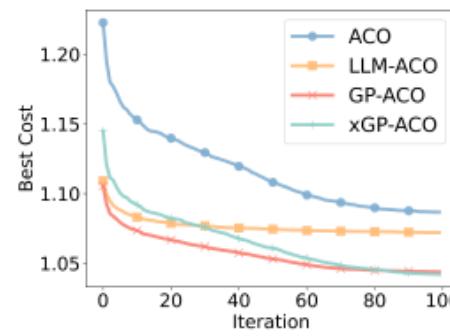
(e) TSP20-100-ACS



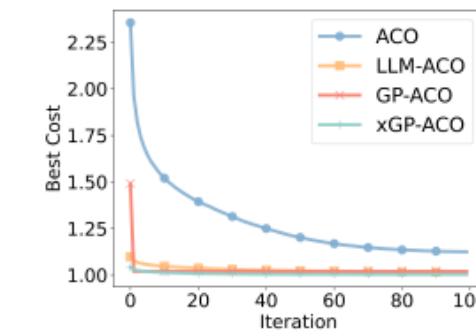
(f) TSP20-100-MMAS



(g) TSP100-AS

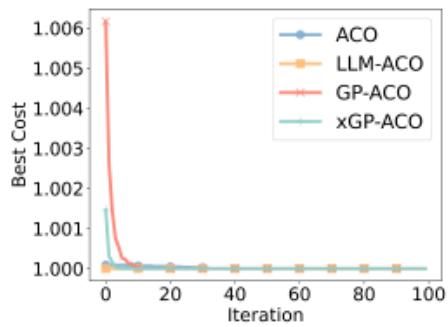


(h) TSP100-ACS

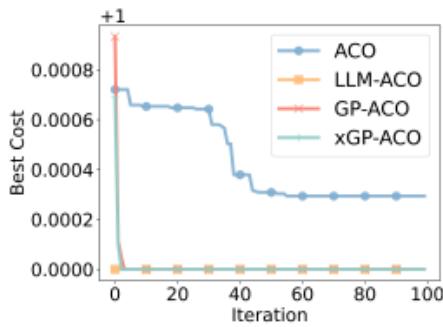


(i) TSP100-MMAS

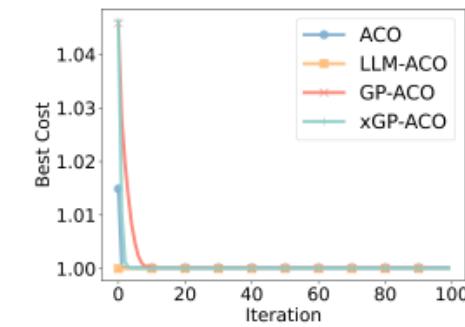
Fig. 7: Convergence of ACO baseline, LLM-ACO, GP-ACO and xGP-ACO



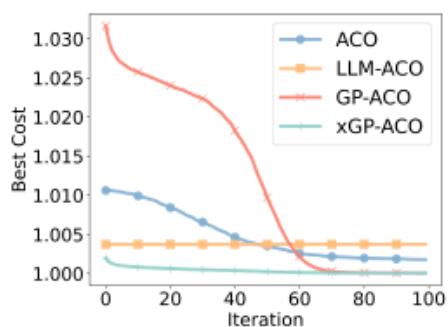
(a) TSP20-AS+2-opt



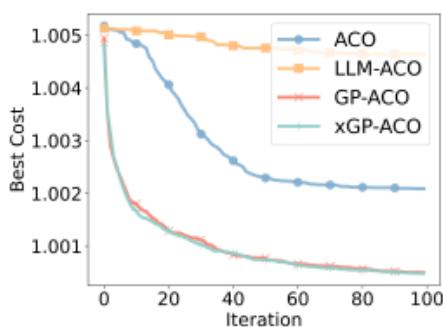
(b) TSP20-ACS+2-opt



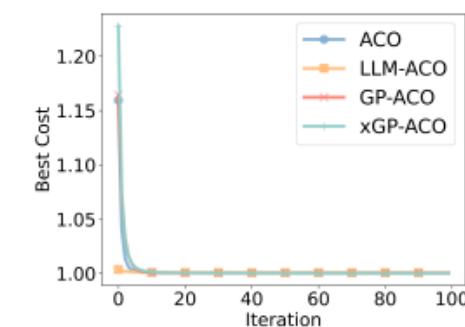
(c) TSP20-MMAS+2-opt



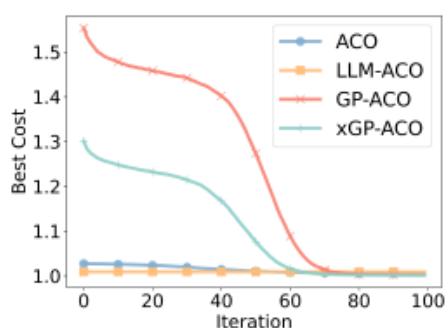
(d) TSP20-100-AS+2-opt



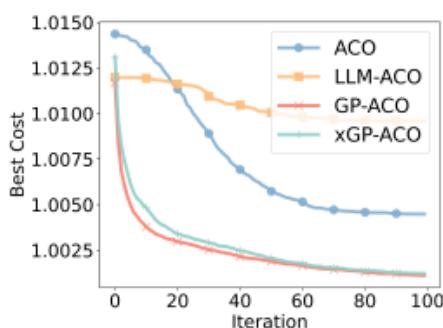
(e) TSP20-100-ACS+2-opt



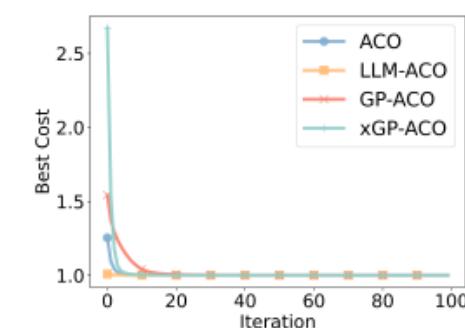
(f) TSP20-100-MMAS+2-opt



(g) TSP100-AS+2-opt



(h) TSP100-ACS+2-opt



(i) TSP100-MMAS+2-opt

Fig. 8: Convergence of ACO+2-opt baseline, LLM-ACO+2-opt, GP-ACO+2-opt and xGP-ACO+2-opt

Table 7: New State Transition rules learned by GP-ACO.

Algorithm	TSP20	TSP20-100	TSP100
GP-AS	$\frac{\varphi}{(d-2\varphi)d^2} - \varphi$	$\frac{\varphi}{(\varphi+d^4)d^5}$	$\frac{2\varphi}{(2\varphi+d^2)(2\varphi+d^4)d^3}$
GP-ACS	$\varphi d^3 - \frac{d^3}{\varphi} + 1$	$\frac{(d^3-1)\varphi}{(\varphi-d)d^4} - d$	$\frac{2\varphi}{(\varphi+d)^2 d} - 2d^2 + 1$
GP-MMAS	$-\varphi d - \frac{2d^2}{\varphi+d}$	$-\frac{2\varphi}{(\varphi+\varphi d+2d^2)d}$	$\frac{\varphi+\varphi(\varphi+d)d^2}{(\varphi+d)^2 d^4}$

Table 8: New State Transition rules learned by xGP-ACO.

Local Search	xGP-AS	xGP-ACS	xGP-MMAS
None	$\frac{n^2 \varphi^2 \bar{d}^2}{\varphi d \bar{d} + d^2}$	$\left(\frac{\bar{d}}{d}\right)^3 - \frac{\varphi}{d^2}$	$\frac{2\bar{\varphi} + \varphi d}{\bar{\varphi}^4 + d^8}$
2-opt	$\frac{n^3 \varphi^3 \bar{d}^2}{(n\varphi+1)(n+\varphi)d}$	$\frac{\varphi^2 \bar{d}}{d^2}$	$-\varphi^2 + (n - n_u - 2) \varphi - n_u - \bar{\varphi}$



To answer Question 6

6. How **interpretable** is GP-ACO?

Answer 6: GP-ACO and xGP-ACO exhibit excellent interpretability.

The primary goals of this paper are to answer four research questions that were formulated concerning

1. generality,
2. the impact of different ACO variants on GP-ACO,
3. the influence of the 2-opt local search on GP-ACO,
4. and the effects of incorporating more global information into GP-ACO.

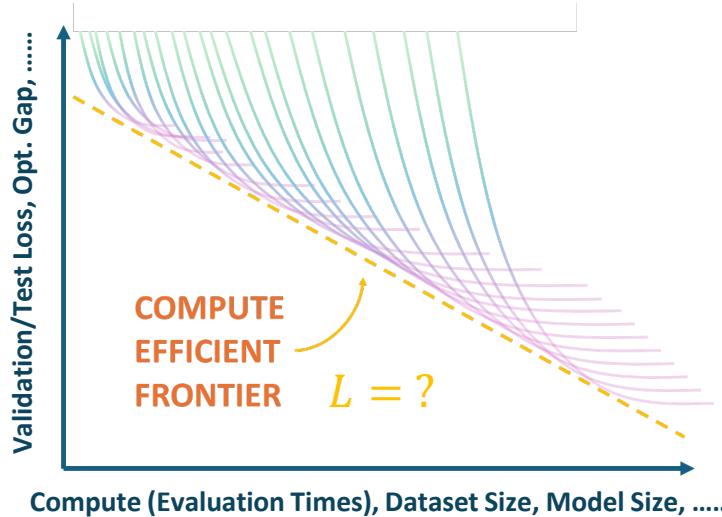
From the experimental results, we observe that:

1. GP-ACO exhibits **robust generality** and performs well across problems of **different scales and distributions**;
2. **Different ACO variants do not affect** the learning capability of GP-ACO;
3. The introduction of 2-opt local search **diminishes the learning capabilities of GP-ACO in the MMAS+2-opt variant**, which has a reduced dependency on state transition rules;
4. Incorporating **additional global information (xGP-ACO)** significantly **enhances the learning capabilities** of GP-ACO;
5. GP-ACO and xGP-ACO show **excellent interpretability**.

1. Introduction
2. Motivations
3. Research Objectives
4. Preliminary Works
5. Proposed Contributions and Research Plan

Proposed Contribution 1

Development of Mathematical Scaling Laws for Heuristic and NCO Methods in TSP



Key Focus:

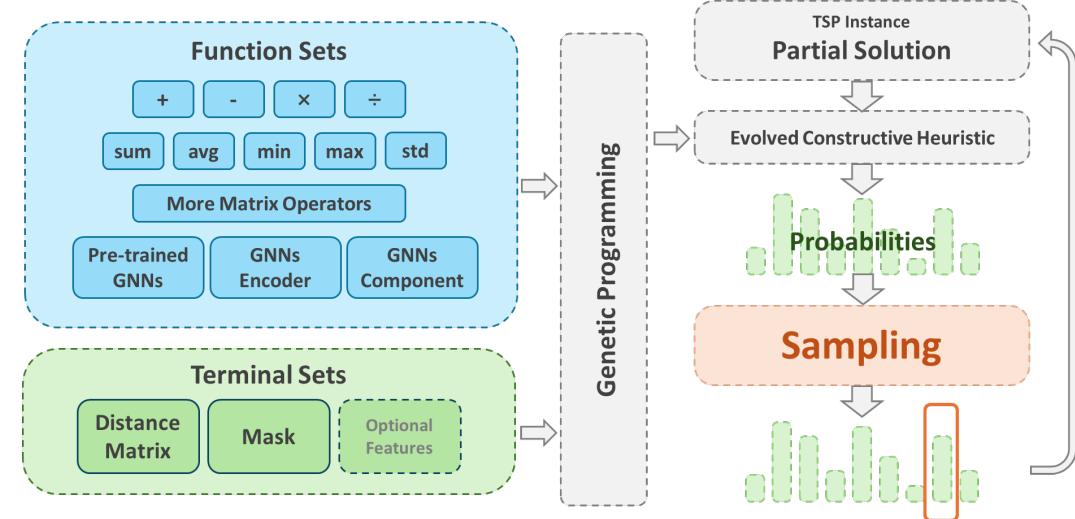
- Characterize performance scaling with problem size, dataset size, model size, and ...
- Predict performance and resource needs for large-scale TSP instances

Expect:

Enable efficient resolution of large TSP problems with improved energy efficiency and predictive performance.

Proposed Contribution 2

Integration of NCO Methods with GP for Enhanced Heuristic Algorithms



Key Focus:

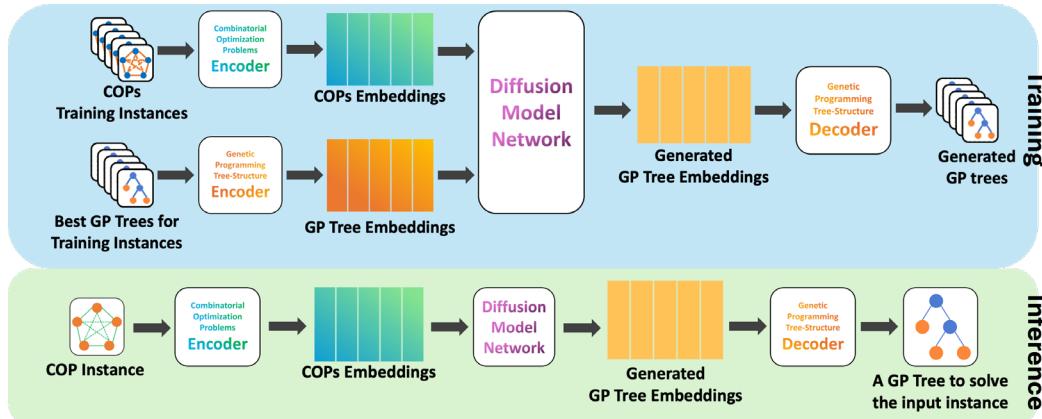
- Augment GP function set with advanced operators and NCO techniques
- Improve generalization and scalability

Expect:

Outperform state-of-the-art heuristic approaches and deliver efficient and adaptable solutions for large-scale TSP.

Proposed Contribution 3

Development of Generative AI Models for Instance-Level Heuristic Generation in TSP



Key Focus:

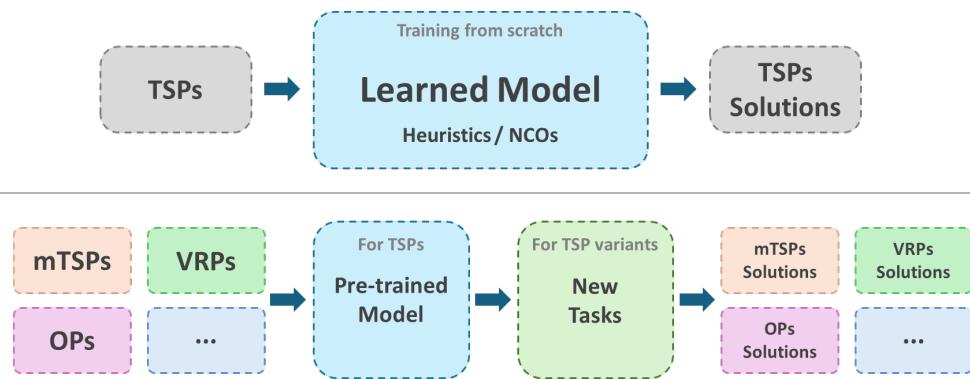
- First to apply generative AI (diffusion models) for heuristic generation in combinatorial optimization
- Adapt heuristics to instance-specific characteristics

Expect:

Enhanced solution quality and computational efficiency; Improved generalization and robustness

Optional Proposed Contribution 4

Enhancing GP and NCO Methods through Transfer Learning and Domain Adaptation for Large-Scale and Variant TSPs



Key Focus:

- Apply transfer learning to adapt models trained on small/simple instances to larger, complex instances, and different TSP variants

Expect:

Improve scalability and adaptability; Enhance generalization and robustness of GP and NCO methods

Research Plan

Phase	Task	Duration
1	Literature Review, overall design, and writing the proposal	12 (completed)
2.1	Develop Scaling Laws for NCO methods	10
2.2	Develop Scaling Laws for population-based heuristics methods	10
3.1	Extending the function set of GP	10
3.2	Develop a new sampling strategy	6
3.3	Prepare and structure data for training generative models	12
4.1	Develop a diffusion model for generating heuristics	12
4.2	Explore self-supervised training methods	6
5	Writing the thesis	4
2.3	Develop theoretical models to explain observed scaling laws	6 <optional>
6.1	Develop a transfer learning framework for GP and NCO methods	6 <optional>
6.2	Develop a domain adaptation method	6 <optional>

Timeline

Phase	Task	Time in Months											
		2	4	6	8	10	12	14	16	18	20	22	24
N/A	Literature Review	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓	✓
2.1	Develop Scaling Laws for NCO methods	✓	✓	✓	✓	✓							
2.2	Develop Scaling Laws for population-based heuristics methods	✓	✓	✓	✓	✓							
3.1	Extending the function set of GP		✓	✓	✓	✓	✓						
3.2	Develop a new sampling strategy			✓	✓	✓							
3.3	Prepare and structure data for training generative models		✓	✓	✓	✓	✓	✓					
4.1	Develop a diffusion model for generating heuristics for specific instances				✓	✓	✓	✓	✓	✓	✓		
4.2	Explore self-supervised training methods						✓	✓	✓	✓			
5.1	Write the first draft of the thesis											✓	✓
5.2	Edit the final draft												✓
2.3	Develop theoretical models to explain observed scaling laws <optional>										✓	✓	✓
6.1	Develop a transfer learning framework for GP and NCO methods <optional>							✓	✓	✓			
6.2	Develop a domain adaptation method <optional>								✓	✓	✓		

Thesis Outline

Chapter 1: Introduction

Chapter 2: Literature Review

Chapter 3: Scaling Laws for Heuristics and NCO in TSP

Chapter 4: Integrating NCO with GP

Chapter 5: Diffusion Models for Generating Heuristics

<optional> Chapter 6: Transfer Learning and Domain Adaptation for GP and NCO

Chapter 7: Conclusion and Future Work



Thank you for your listening!

Bocheng Lin 2024.12

Shot at Tongariro National Park